

OPERATING SYSTEMS

MASTER IN COMPUTER SCIENCE & BUSINESS TECHNOLOGY

Threads(2) & Process Cooperation

Professor: Olivier Perard

Email: operard@faculty.ie.edu

Github:

https://github.com/operard/opsys_parallel/blob/master/mcsbt/README.md

Threads Implementation

Contents

- Multithreading Levels
- Multithreading Models
- Threading Issues



Multithreading vs. Single threading

- Single threading: when the OS does not recognize the concept of thread.
- Multithreading: when the OS supports multiple threads of execution within a single process.
- MS-DOS supports a single user process and a single thread.
- Older UNIXs supports multiple user processes but only support one thread per process.
- Solaris and Windows NT support multiple threads.

Multithreading Levels

- Thread library provides programmer with API for creating and managing threads.
- Three multithreading levels:
 - 1) User-Level Threads (ULT)
 - Library entirely in user space.
 - 2) Kernel-Level Threads (KLT)
 - Kernel-level library supported by the OS.
 - 3) Hybrid ULT/KLT Approach




User(-level) and Kernel(-level) Threads


- User(-level) threads – management by user-level threads library.
- Three primary thread libraries:
 - POSIX Pthreads
 - Windows threads
 - Java threads
- Kernel(-level) threads – supported by the Kernel.
- Examples – virtually all general purpose operating systems, including:
 - Windows
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X

1) User-Level Threads (ULT)

- Thread management done by user-level threads library
- The kernel is not aware of the existence of threads.
- All thread management is done by the application by using a thread library.
- Thread switching does not require kernel mode privileges.
- Scheduling is application specific



Implementing Threads in User Space



ULT Idea

- Thread management done by user-level threads library.
- Threads library contains code for:
 - creating and destroying threads.
 - passing messages and data between threads.
 - scheduling thread execution.
 - saving and restoring thread contexts.
- Three primary thread libraries:
 - POSIX Pthreads
 - Win32 threads
 - Java threads



POSIX Pthreads

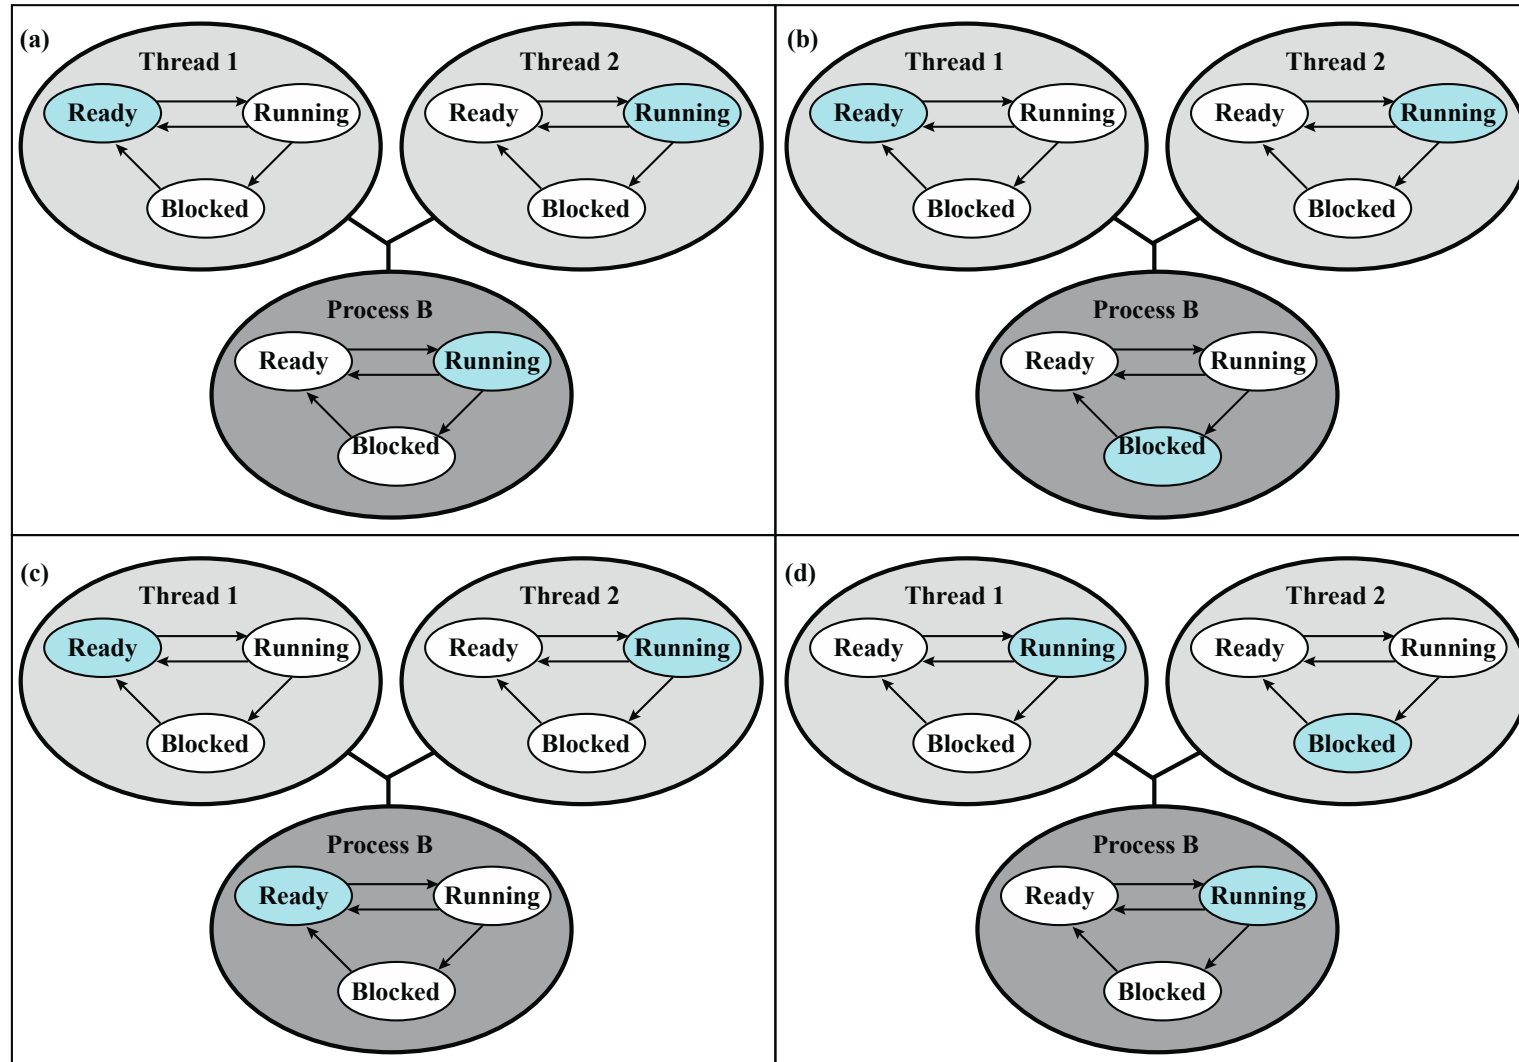
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization.
- May be provided either as ULT or KLT.
- API specifies behavior of the thread library, implementation is up to development of the library.
- Common in UNIX operating systems (Solaris, Linux, Mac OS X).

Some of the Pthreads function calls

Thread call	Description
Pthread_create	Create a new thread
Pthread_exit	Terminate the calling thread
Pthread_join	Wait for a specific thread to exit
Pthread_yield	Release the CPU to let another thread run
Pthread_attr_init	Create and initialize a thread's attribute structure
Pthread_attr_destroy	Remove a thread's attribute structure

Kernel activity for ULTs


- The kernel is not aware of thread activity but it is still managing process activity.
- When a thread makes a system call, the whole task will be blocked.
- But for the thread library that thread is still in the running state.
- So thread states are independent of process states.



Colored state
is current state

Figure 4.6 Examples of the Relationships Between User-Level Thread States and Process States

Advantages of ULTs



Advantages and inconveniences of ULT

- Advantages


- Thread switching does not involve the kernel: no mode switching.
- Scheduling can be application specific: choose the best algorithm.
- ULTs can run on any OS. Only needs a thread library.

- Inconveniences


- Most system calls are blocking and the kernel blocks processes. So all threads within the process will be blocked.
- The kernel can only assign processes to processors. Two threads within the same process cannot run simultaneously on two processors.

2) Kernel-Level Threads (KLT)

- All thread management is done by kernel.
- No thread library but an API to the kernel thread facility.
- Kernel maintains context information for the process and the threads.
- Switching between threads requires the kernel.
- Scheduling on a thread basis.



Implementing Threads in the Kernel



KLT Idea

- Threads supported by the Kernel.
- Examples:
 - Windows 2000/XP
 - OS/2
 - Linux
 - Solaris
 - Tru64 UNIX
 - Mac OS X



Linux Threads

- Linux refers to them as tasks rather than threads.
- Thread creation is done through **clone()** system call.
- **clone()** allows a child task to share the address space of the parent task (process).
- This sharing of the address space allows the cloned child task to behave much like a separate thread.

Advantages and inconveniences of KLT

- Advantages

- the kernel can simultaneously schedule many threads of the same process on many processors.
- blocking is done on a thread level.
- kernel routines can be multithreaded.

- Inconveniences

- thread switching within the same process involves the kernel. We have 2 mode switches per thread switch.
- this results in a significant slow down.


Thread operation latencies () μs

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840


Source: Anderson, T. et al, “Scheduler Activations: Effective Kernel Support for the User-Level Management of Parallelism”, ACM TOCS, February 1992.

3) Hybrid ULT/KLT Approaches

- Thread creation done in the user space.
- Bulk of scheduling and synchronization of threads done in the user space.
- The programmer may adjust the number of KLTs.
- May combine the best of both approaches.
- Example is Solaris prior to version 9.




Hybrid Implementation (1)

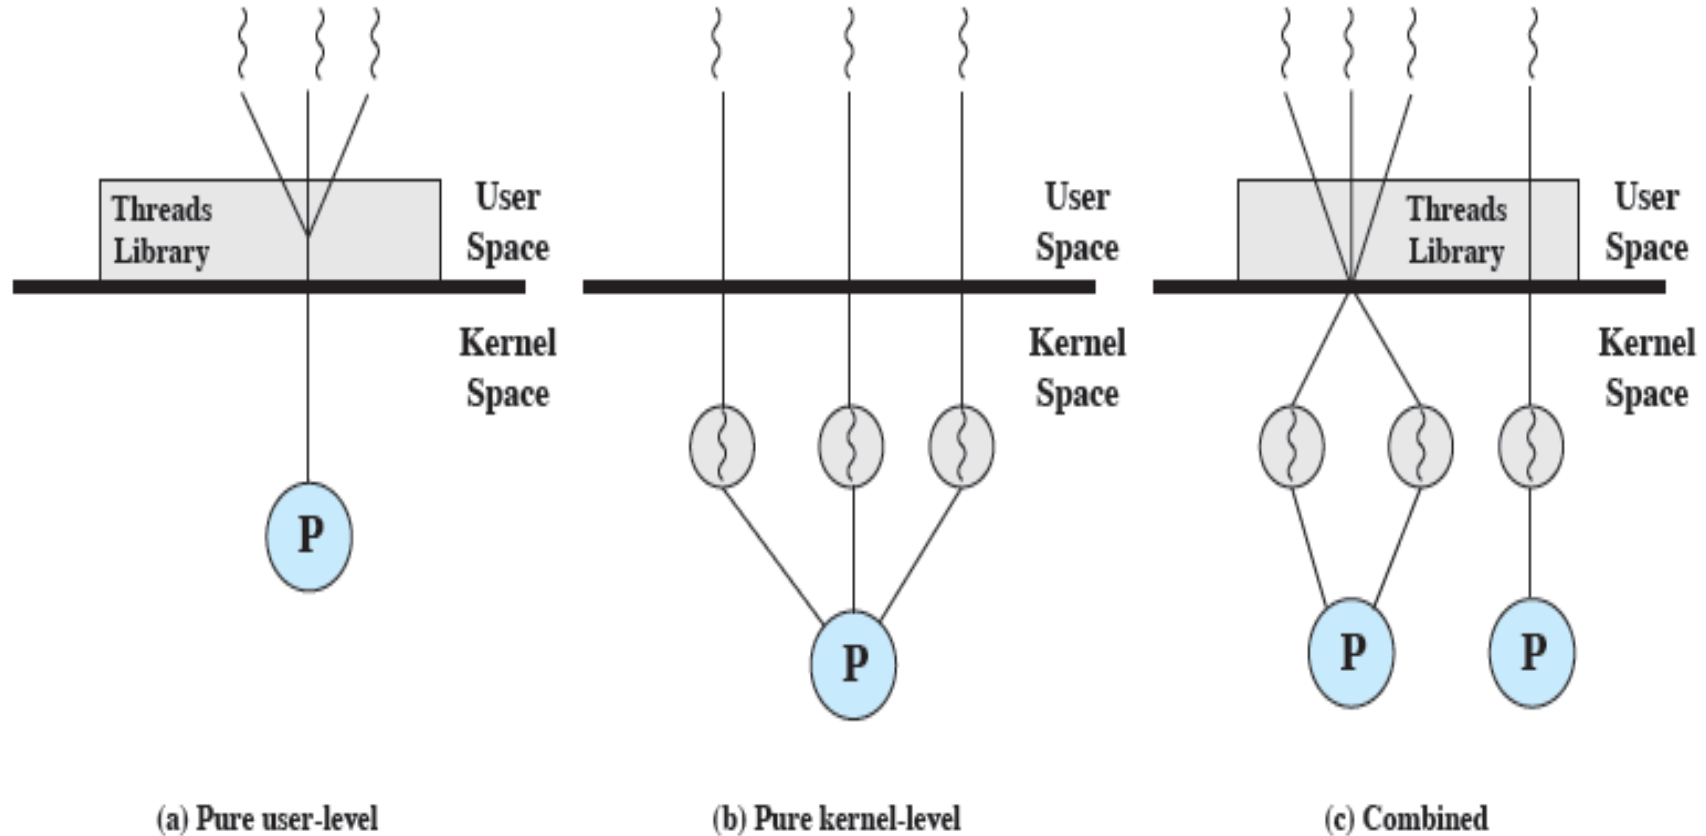


Multiplexing user-level threads onto kernel-level threads.

Hybrid Implementation (2)



ULT, KLT and Combined Approaches



Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many
- Two-level Model




Relationship between Threads and Processes

Threads:Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	Combines attributes of M:1 and 1:M cases.	TRIX

Many-to-One Model (1)

- Many user-level threads mapped to single kernel-level thread.




Many-to-One Model (2)

- One thread blocking causes all to block.
- Multiple threads may not run in parallel on multi-core system because only one may be in kernel at a time.
- Few systems currently use this model.
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads

One-to-One Model (1)

- Each user-level thread maps to kernel-level thread.




One-to-One Model (2)

- Creating a user-level thread creates a kernel thread.
- More concurrency than many-to-one.
- Number of threads per process sometimes restricted due to overhead.
- Examples
 - Windows
 - Linux
 - Solaris 9 and later

Many-to-Many Model (1)

Allows many user level threads to be mapped to many kernel threads.




Many-to-Many Model (2)

- Allows the operating system to create a sufficient number of kernel threads.
- Solaris prior to version 9.
- Windows with the *ThreadFiber* package.

Two-level Model (1)


- Similar to Many-to-Many model, except that it allows a user thread to be **bound** to kernel thread.




Two-level Model (2)

- Examples:
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier

Lightweight Process (LWP)

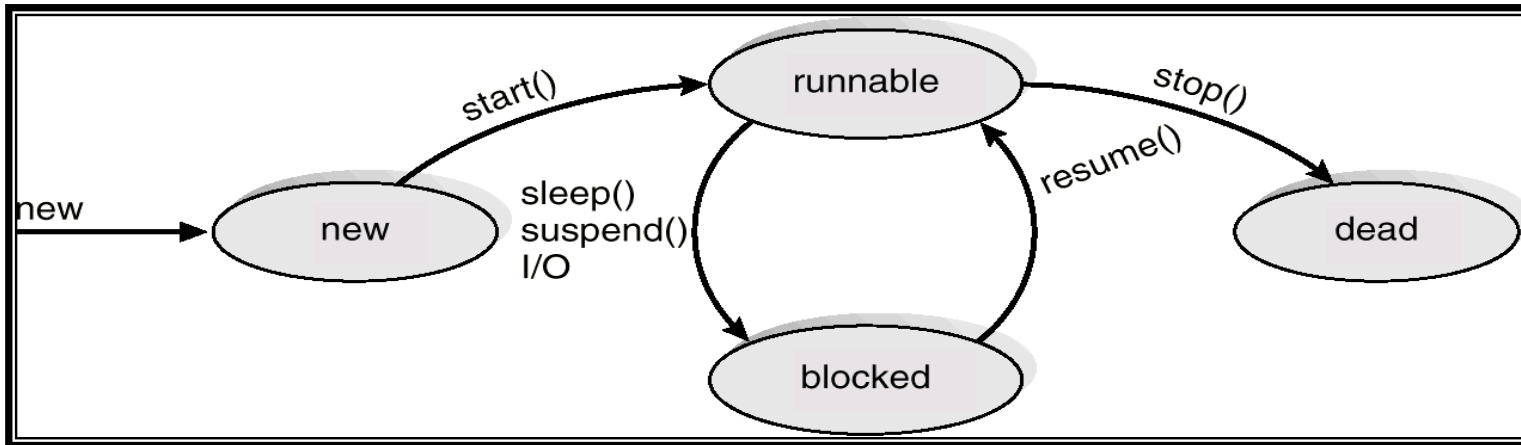


Solaris 2 Threads



Java Threads

- Java threads are managed by the JVM.
- Typically implemented using the threads model provided by underlying OS.
- Java threads may be created by:
 - Extending Thread class (at language-level)
 - Implementing the Runnable interface



Threading Issues

- Semantics of **fork()** and **exec()** system calls
 - Does **fork()** duplicate only the calling thread or all threads?
- Thread cancellation of target thread
 - Asynchronous or deferred
- Signal handling
- Thread pools
- Thread-local storage
- Scheduler activations



Thread Cancellation

- Terminating a thread before it has finished.
- Two general approaches:
 - Asynchronous cancellation terminates the target thread immediately.
 - Deferred cancellation allows the target thread to periodically check if it should be cancelled.

Signal Handling (1)

- Signals are used in UNIX systems to notify a process that a particular event has occurred.
- A signal handler is used to process signals:
 1. Signal is generated by particular event.
 2. Signal is delivered to a process.
 3. Signal is handled by one of two signal handlers:
 1. default
 2. user-defined
- Every signal has default handler that kernel runs when handling signal:
 - | User-defined signal handler can override default.
 - | For single-threaded, signal delivered to process.

Signal Handling (2)

- Where should a signal be delivered for multi-threaded?
 - Deliver the signal to the thread to which the signal applies.
 - Deliver the signal to every thread in the process.
 - Deliver the signal to certain threads in the process.
 - Assign a specific thread to receive all signals for the process.

Thread Pools

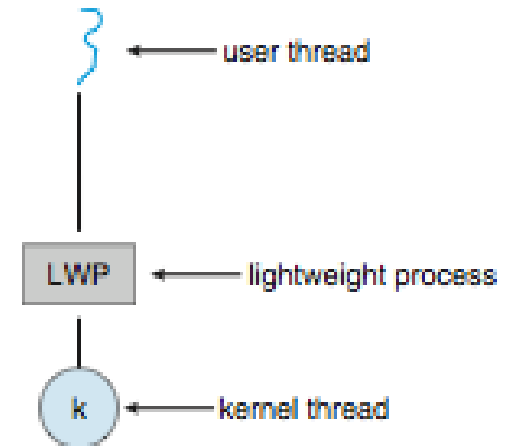
- Create a number of threads in a pool where they await work.
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread.
 - Allows the number of threads in the application(s) to be bound to the size of the pool.

Thread-Local Storage

- Thread-local storage (TLS) allows each thread to have its own copy of data.
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool).
- Different from local variables:
 - Local variables visible only during single function invocation.
 - TLS visible across function invocations.
- Similar to **static** data:
 - TLS is unique to each thread.

Scheduler Activations

- Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application/
- Typically use an intermediate data structure between user and kernel threads – lightweight process (LWP):
 - Appears to be a virtual processor on which process can schedule user thread to run.
 - Each LWP attached to kernel thread.
 - How many LWPs to create?
- Scheduler activations provide upcalls – a communication mechanism from the kernel to the upcall handler in the thread library.
- This communication allows an application to maintain the correct number kernel threads.



Operating System Examples

- Windows Threads
- Linux Threads

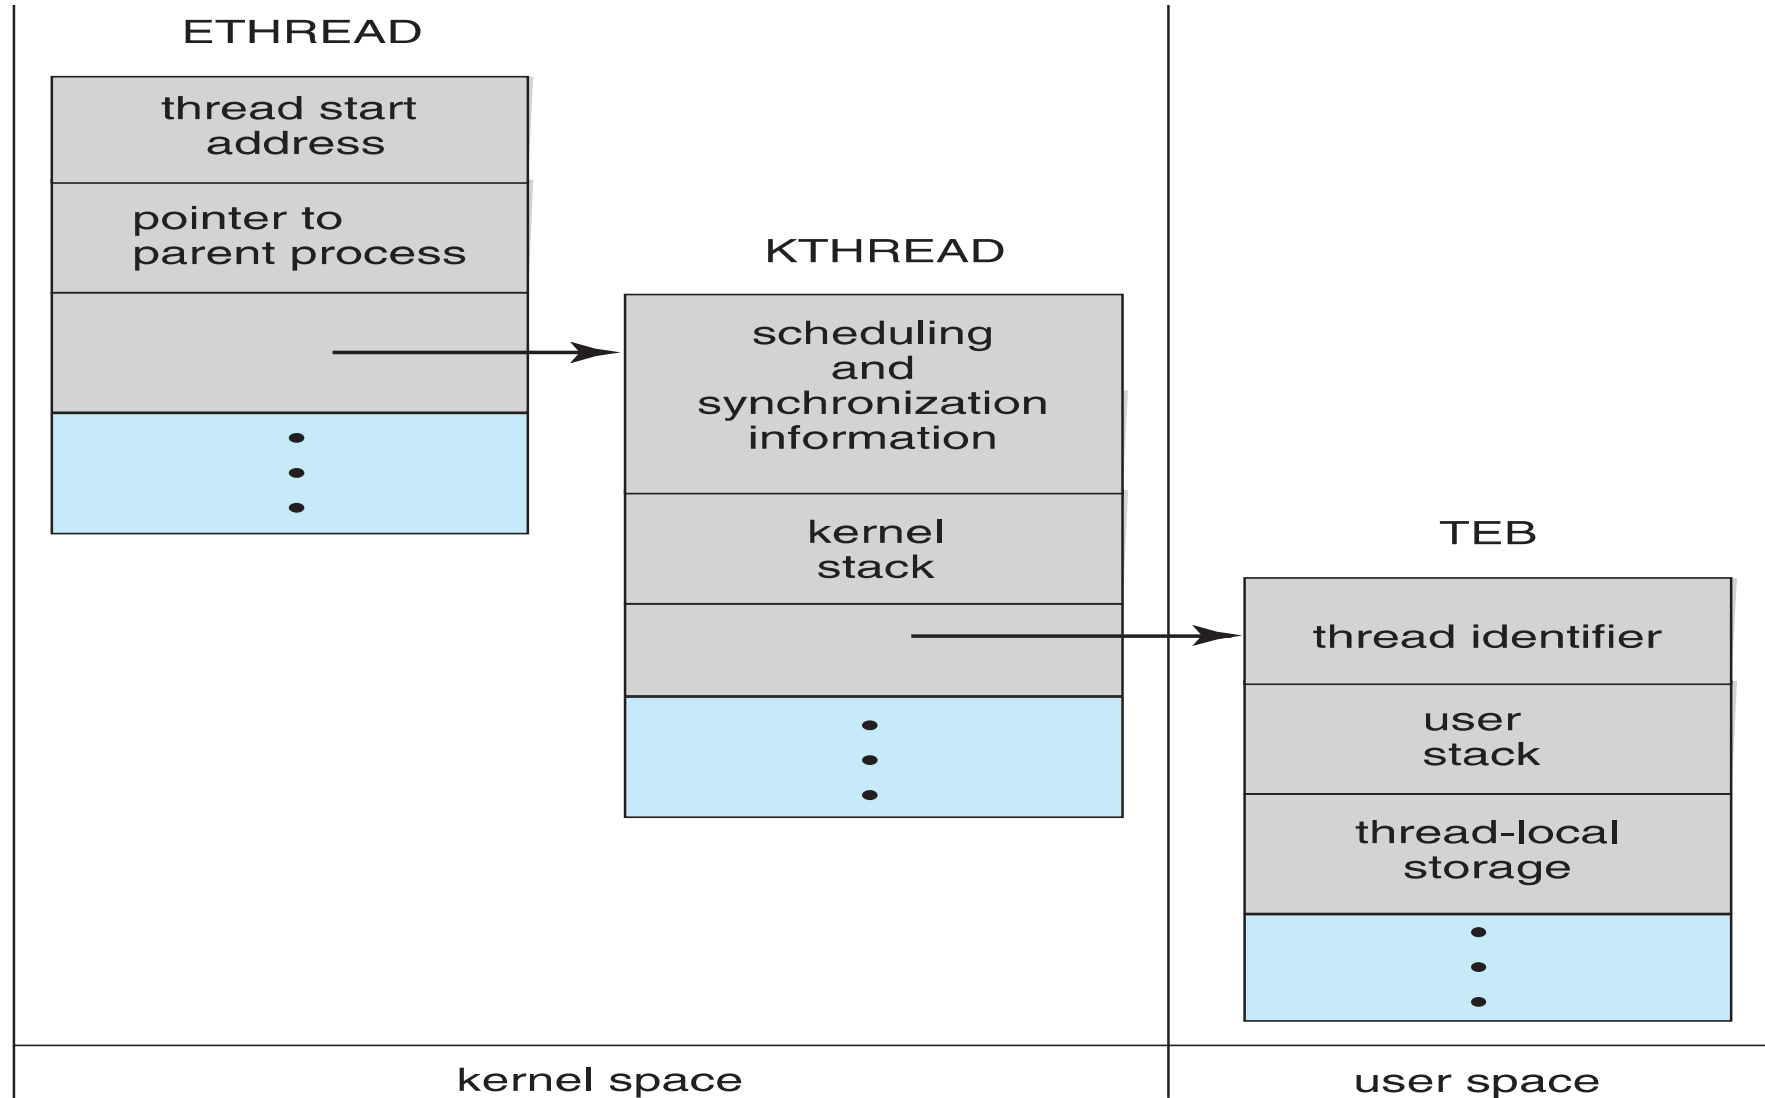
Windows Threads (1)

- Windows implements the Windows API – primary API for Win 98, Win NT, Win 2000, Win XP, and Win 7.
- Implements the one-to-one mapping, kernel-level.
- Each thread contains:
 - A thread id.
 - Register set representing state of processor.
 - Separate user and kernel stacks for when thread runs in user mode or kernel mode.
 - Private data storage area used by run-time libraries and dynamic link libraries (DLLs).
- The register set, stacks, and private storage area are known as the context of the thread.

Windows Threads (2)

- The primary data structures of a thread include:
 - ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD, in kernel space.
 - KTHREAD (kernel thread block) – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space.
 - TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space.

Windows Threads Data Structures



Linux Threads

- Linux refers to them as ***tasks*** rather than ***threads***.
- Thread creation is done through **`clone()`** system call.
- **`clone()`** allows a child task to share the address space of the parent task (process).

- Flags control behavior.

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- **`struct task_struct`** points to process data structures (shared or unique).

Thread Python Implementation

Starting a new Thread

EXAMPLE:

```
#!/usr/bin/python
import thread
import time
# Define a function for the thread
def print_time( threadName, delay):
    count = 0
    while count < 5:
        time.sleep(delay)
        count += 1
        print "%s: %s" % ( threadName, time.ctime(time.time()) )
# Create two threads as follows
try:
    thread.start_new_thread( print_time, ("Thread-1", 2, ) )
    thread.start_new_thread( print_time, ("Thread-2", 4, ) )
except:
    print "Error: unable to start thread"
while 1:
    pass
```

Creating Thread using Threading Module

- To implement a new thread using the threading module, need to do the following:
 - Define a new subclass of the Thread class.
 - Override the `__init__(self [,args])` method to add additional arguments.
 - override the `run(self [,args])` method to implement what the thread should do when started.
- Once you have created the new Thread subclass, you can create an instance of it and then start a new thread by invoking the `start()`, which will in turn call `run()` method.

EXAMPLE:

```
#!/usr/bin/python
import threading
import time
exitFlag = 0
class myThread (threading.Thread):
    def __init__(self, threadID, name, counter):
        threading.Thread.__init__(self)
        self.threadID = threadID
        self.name = name
        self.counter = counter
    def run(self):
        print "Starting " + self.name
        print_time(self.name, self.counter, 5)
        print "Exiting " + self.name
def print_time(threadName, delay, counter):
    while counter:
        if exitFlag:
            thread.exit()
        time.sleep(delay)
        print "%s: %s" % (threadName,
time.ctime(time.time()))
        counter -= 1
```

```
# Create new threads
thread1 = myThread(1, "Thread-1", 1)
thread2 = myThread(2, "Thread-2", 2)
# Start new Threads
thread1.start()
thread2.start()
print "Exiting Main Thread"
```

When the above code is executed, it produces the following result:

```
Starting Thread-1
Starting Thread-2
Exiting Main Thread
Thread-1: Thu Mar 21 09:10:03 2013
Thread-1: Thu Mar 21 09:10:04 2013
Thread-2: Thu Mar 21 09:10:04 2013
Thread-1: Thu Mar 21 09:10:05 2013
Thread-1: Thu Mar 21 09:10:06 2013
Thread-2: Thu Mar 21 09:10:06 2013
Thread-1: Thu Mar 21 09:10:07 2013
Exiting Thread-1
Thread-2: Thu Mar 21 09:10:08 2013
Thread-2: Thu Mar 21 09:10:10 2013
Thread-2: Thu Mar 21 09:10:12 2013
Exiting Thread-2
```

Synchronizing Threads

- The threading module provided with Python includes a simple-to-implement locking mechanism that will allow you to synchronize threads
- • A new lock is created by calling the `Lock()` method, which returns the new lock.
- • The `acquire(blocking)` method of the new lock object would be used to force threads to run synchronously
- • The optional blocking parameter enables you to control whether the thread will wait to acquire the lock.
- • If blocking is set to 0, the thread will return immediately with a 0 value if the lock cannot be acquired and with a 1 if the lock was acquired. If blocking is set to 1, the thread will block and wait for the lock to be released.
- • The `release()` method of the the new lock object would be used to release the lock when it is no longer required.

EXAMPLE:

```
#!/usr/bin/python
import threading
import time
class myThread (threading.Thread):
    def __init__(self, threadID, name, counter):
        threading.Thread.__init__(self)
        self.threadID = threadID
        self.name = name
        self.counter = counter
    def run(self):
        print "Starting " + self.name
        # Get lock to synchronize threads
        threadLock.acquire()
        print_time(self.name, self.counter, 3)
        # Free lock to release next thread
        threadLock.release()
def print_time(threadName, delay, counter):
    while counter:
        time.sleep(delay)
        print "%s: %s" % (threadName,
            time.ctime(time.time()))
        counter -= 1
threadLock = threading.Lock()
threads = []
```

```
# Create new threads
thread1 = myThread(1, "Thread-1", 1)
thread2 = myThread(2, "Thread-2", 2)
# Start new Threads
thread1.start()
thread2.start()
# Add threads to thread list
threads.append(thread1)
threads.append(thread2)
# Wait for all threads to complete
for t in threads:
    t.join()
print "Exiting Main Thread"
```

When the above code is executed, it produces the following result:

```
Starting Thread-1
Starting Thread-2
Thread-1: Thu Mar 21 09:11:28 2013
Thread-1: Thu Mar 21 09:11:29 2013
Thread-1: Thu Mar 21 09:11:30 2013
Thread-2: Thu Mar 21 09:11:32 2013
Thread-2: Thu Mar 21 09:11:34 2013
Thread-2: Thu Mar 21 09:11:36 2013
Exiting Main Thread
```


MultiProcessing

```
1 import multiprocessing
2 def spawn():
3     print('test!')
4
5 if __name__ == '__main__':
6     for i in range(5):
7         p = multiprocessing.Process(target=spawn)
8         p.start()
```

If you have a shared database, you want to make sure that you're waiting for relevant processes to finish before starting new ones.

```
1 for i in range(5):
2     p = multiprocessing.Process(target=spawn)
3     p.start()
4     p.join() # this line allows you to wait for processes
```

MultiProcessing

```
1 import multiprocessing
2 def spawn(num):
3     print(num)
4
5 if __name__ == '__main__':
6     for i in range(25):
7         ## right here
8         p = multiprocessing.Process(target=spawn, args=(i,))
9         p.start()
```

R Implementation

Methods of Parallelization

- There are two main ways in which code can be parallelized, via *sockets* or via *forking*. These function slightly differently:
 - The *socket* approach launches a new version of R on each core. Technically this connection is done via networking (e.g. the same as if you connected to a remote server), but the connection is happening all on your own computer³ I mention this because you may get a warning from your computer asking whether to allow R to accept incoming connections, you should allow it.
 - The *forking* approach copies the entire current version of R and moves it to a new core.

Methods of Parallelization

- Socket:
 - Pro: Works on any system (including Windows).
 - Pro: Each process on each node is unique so it can't cross-contaminate.
 - Con: Each process is unique so it will be slower
 - Con: Things such as package loading need to be done in each process separately. Variables defined on your main version of R don't exist on each core unless explicitly placed there.
 - Con: More complicated to implement.
- Forking:
 - Con: Only works on POSIX systems (Mac, Linux, Unix, BSD) and not Windows.
 - Con: Because processes are duplicates, it can cause issues specifically with random number generation (which should usually be handled by parallel in the background) or when running in a GUI (such as RStudio). This doesn't come up often, but if you get odd behavior, this may be the case.
 - Pro: Faster than sockets.
 - Pro: Because it copies the existing version of R, your entire workspace exists in each process.
 - Pro: Trivially easy to implement.

Forking with `mclapply`

The most straightforward way to enable parallel processing is by switching from using `lapply` to `mclapply`. (Note I'm using `system.time` instead of `profvis` here because I only care about running time, not profiling.)

```
library(lme4)
```

```
## Loading required package: Matrix
```

```
f <- function(i) {  
  lmer(Petal.Width ~ . - Species + (1 | Species), data = iris)  
}
```

```
system.time(save1 <- lapply(1:100, f))
```

```
##      user  system elapsed  
##    2.048    0.019    2.084
```

```
system.time(save2 <- mclapply(1:100, f))
```

```
##      user  system elapsed  
##    1.295    0.150    1.471
```

Using sockets with `parLapply`

As promised, the sockets approach to parallel processing is more complicated and a bit slower, but works on Windows systems. The general process we'll follow is

1. Start a cluster with n nodes.
2. Execute any pre-processing code necessary in each node (e.g. loading a package)
3. Use `par*apply` as a replacement for `*apply`. Note that unlike `mapply`, this is *not* a drop-in replacement.
4. Destroy the cluster (not necessary, but best practices).

Starting a cluster

The function to start a cluster is `makeCluster` which takes in as an argument the number of cores:

```
numCores <- detectCores()  
numCores
```

```
## [1] 4
```

```
cl <- makeCluster(numCores)
```

The function takes an argument `type` which can be either `PSOCK` (the socket version) or `FORK` (the fork version). Generally, `mclapply` should be used for the forking approach, so there's no need to change this.

```
### lapply
library(parallel)
f <- function(i) {
  lmer(Petal.Width ~ . - Species + (1 | Species), data = iris)
}

system.time({
  library(lme4)
  save1 <- lapply(1:100, f)
})

### mclapply
library(parallel)
f <- function(i) {
  lmer(Petal.Width ~ . - Species + (1 | Species), data = iris)
}

system.time({
  library(lme4)
  save2 <- mclapply(1:100, f)
})

### mclapply
library(parallel)
f <- function(i) {
  lmer(Petal.Width ~ . - Species + (1 | Species), data = iris)
}

system.time({
  cl <- makeCluster(detectCores())
  clusterEvalQ(cl, library(lme4))
  save3 <- parLapply(cl, 1:100, f)
  stopCluster(cl)
})
```

lapply	mclapply	parLapply
4.237	4.087	6.954


```
library(parallel)
library(MASS)

starts <- rep(100, 40)
fx <- function(nstart) kmeans(Boston, 4, nstart=nstart)
numCores <- detectCores()
numCores
```

```
## [1] 8
```

```
system.time(
  results <- lapply(starts, fx)
)
```

```
##      user  system elapsed
##  1.346    0.024    1.372
```

```
system.time(
  results <- mclapply(starts, fx, mc.cores = numCores)
)
```

```
##      user  system elapsed
##  0.801    0.178    0.367
```

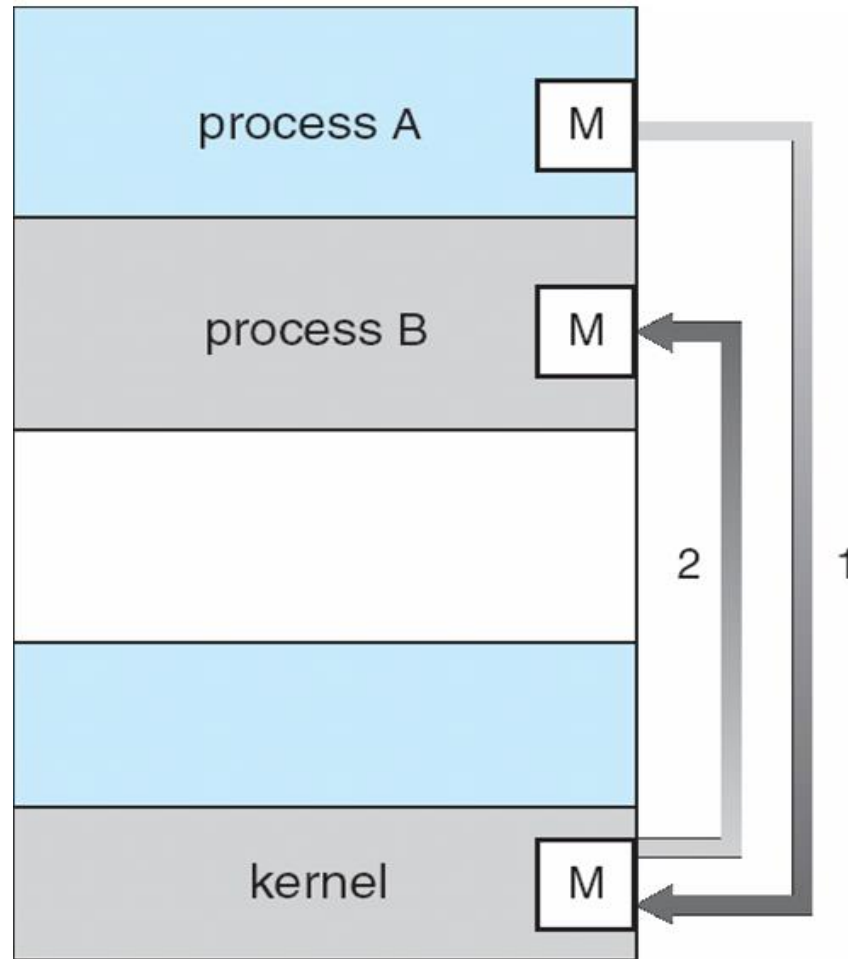
Introduction to Cooperating Process

Introduction to Cooperating Processes


- Processes within a system may be independent or cooperating.
- Independent process cannot affect or be affected by the execution of another process.
- Cooperating process can affect or be affected by other processes, including sharing data.
- Reasons for cooperating processes:
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience



Cooperation Models



(a)



(b)

Cooperation among Processes by Sharing

- Processes use and update shared data such as shared variables, memory, files, and databases.
- Writing must be mutually exclusive to prevent a race condition leading to inconsistent data views.
- Critical sections are used to provide this data integrity.
- A process requiring the critical section must not be delayed indefinitely; no deadlock or starvation.

Cooperation among Processes by Communication

- Communication by messages provides a way to synchronize, or coordinate, the various activities.
- Possible to have deadlock –
 - each process waiting for a message from the other process.
- Possible to have starvation –
 - two processes sending a message to each other while another process waits for a message.

Producer/Consumer (P/C) Problem (1)


- Paradigm for cooperating processes – *Producer* process produces information that is consumed by a *Consumer* process.
 - Example 1: a print program produces characters that are consumed by a printer.
 - Example 2: an assembler produces object modules that are consumed by a loader.
 - Example 3: a client produces a message that a server could consume it.

Producer/Consumer (P/C) Problem (2)

- We need a buffer to hold items that are produced and later consumed:
 - *unbounded-buffer* places no practical limit on the size of the buffer.
 - *bounded-buffer* assumes that there is a fixed buffer size.



Multiple Producers and Consumers




Producer/Consumer (P/C) Dynamics

- A producer process produces information that is consumed by a consumer process.
- At any time, a producer activity may create some data.
- At any time, a consumer activity may want to accept some data.
- The data should be saved in a buffer until they are needed.
- If the buffer is finite, we want a producer to block if its new data would overflow the buffer.
- We also want a consumer to block if there are no data available when it wants them.

Idea for Producer/Consumer Solution

- The bounded buffer is implemented as a circular array with 2 logical pointers: **in** and **out**.
- The variable **in** points to the next free position in the buffer.
- The variable **out** points to the first full position in the buffer.



Problems with concurrent execution

- Concurrent processes (or threads) often need to share data (maintained either in shared memory or files) and resources.
- If there is no controlled access to shared data, some processes will obtain an inconsistent view of this data.
- The action performed by concurrent processes will then depend on the order in which their execution is interleaved.

Example of inconsistent view

- 3 variables: A, B, C which are shared by thread T1 and thread T2.
- T1 computes $C = A+B$.
- T2 transfers amount X from A to B
 - T2 must do: $A = A-X$ and $B = B+X$ (so that $A+B$ is unchanged).
- But if T1 computes $A+B$ after T2 has done $A = A-X$ but before $B = B+X$.
- Then T1 will not obtain the correct result for $C = A+B$.

Data Consistency

- Concurrent access to shared data may result in data inconsistency.
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills all the buffer. We can do so by having an integer count that keeps track of the number of items in the buffer.
- Initially, the count is set to 0. It is incremented by the producer after it produces a new item and is decremented by the consumer after it consumes a item.

This is the Race Condition

- **Race condition:** The situation where several processes access and manipulate shared data concurrently. The final value of the shared data depends upon which process finishes last.
- To prevent race conditions, concurrent processes must coordinate or be **synchronized**.



Race condition updating a variable (1)

shared double balance;

Code for p1:

...

balance += amount;

...

Code for p1:

...

Load R1, balance

Load R2, amount

Add R1, R2

Store R1, balance

. . .

Code for p2:

...

balance += amount;

...

Code for p2:

...

Load R1, balance

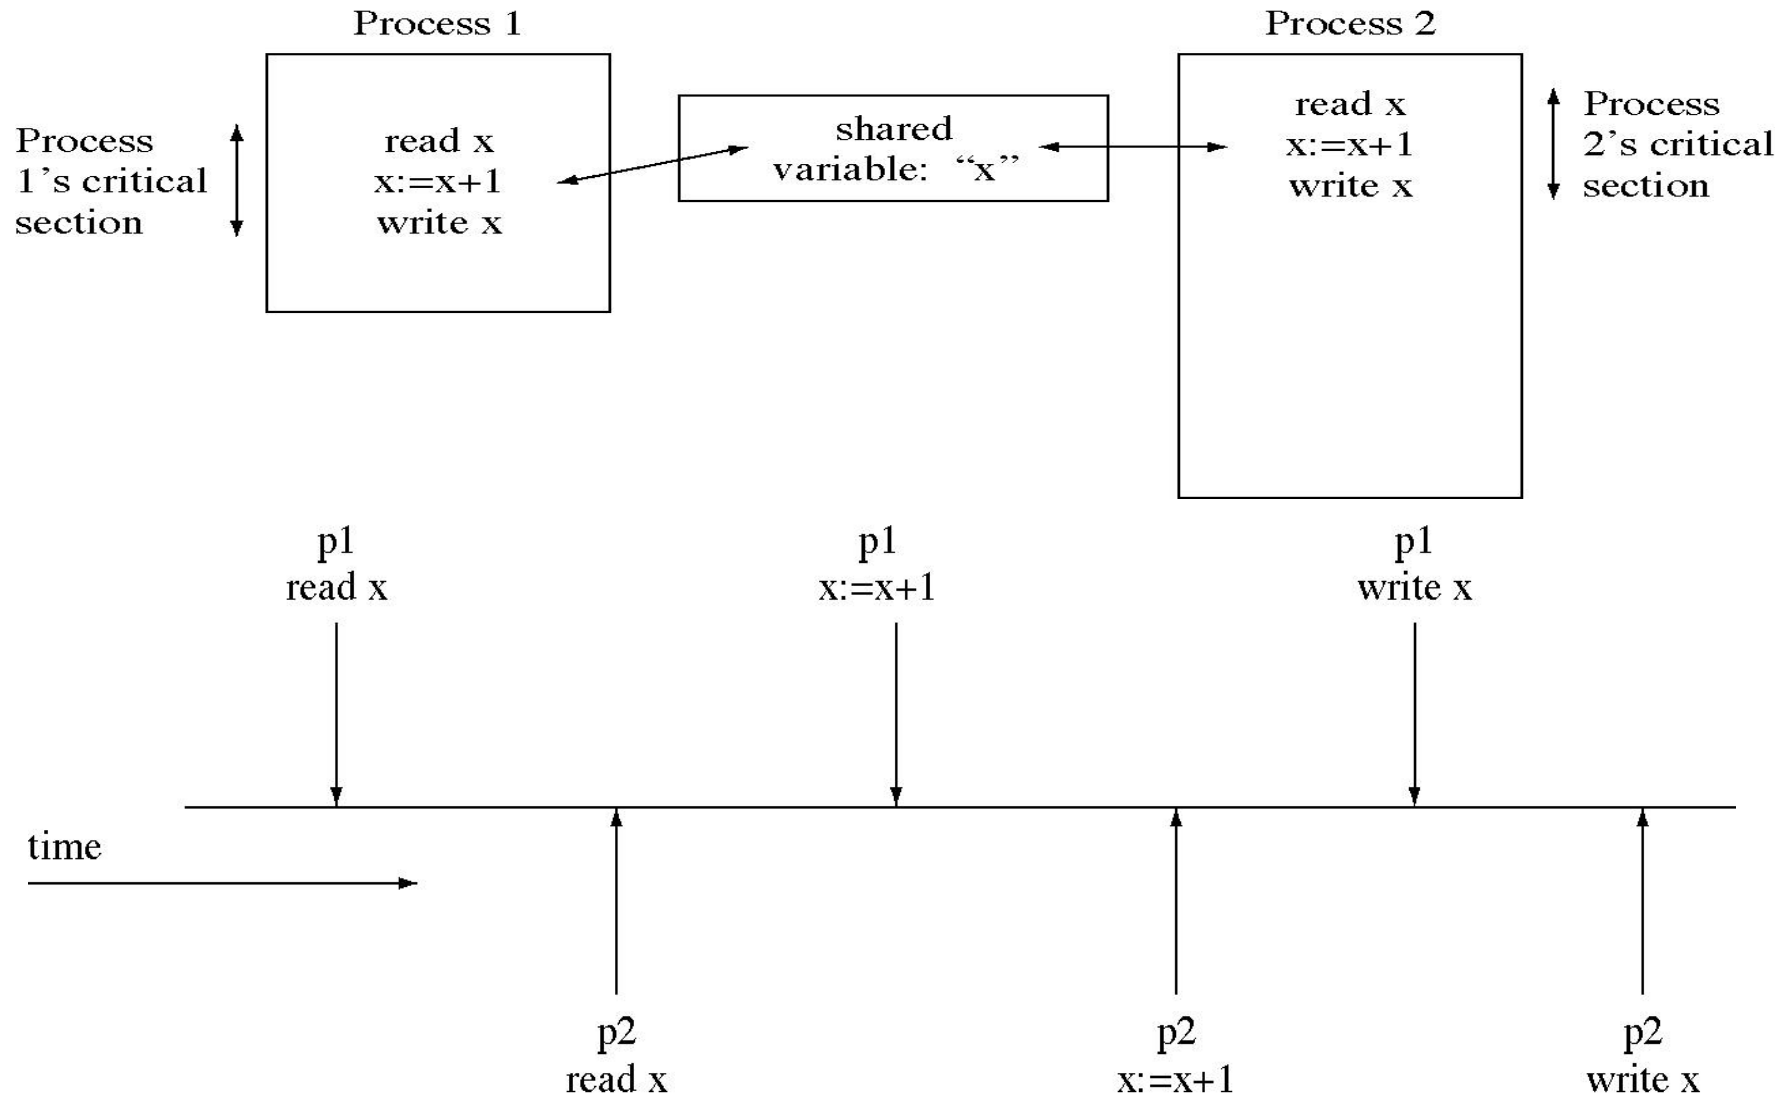
Load R2, amount

Add R1, R2

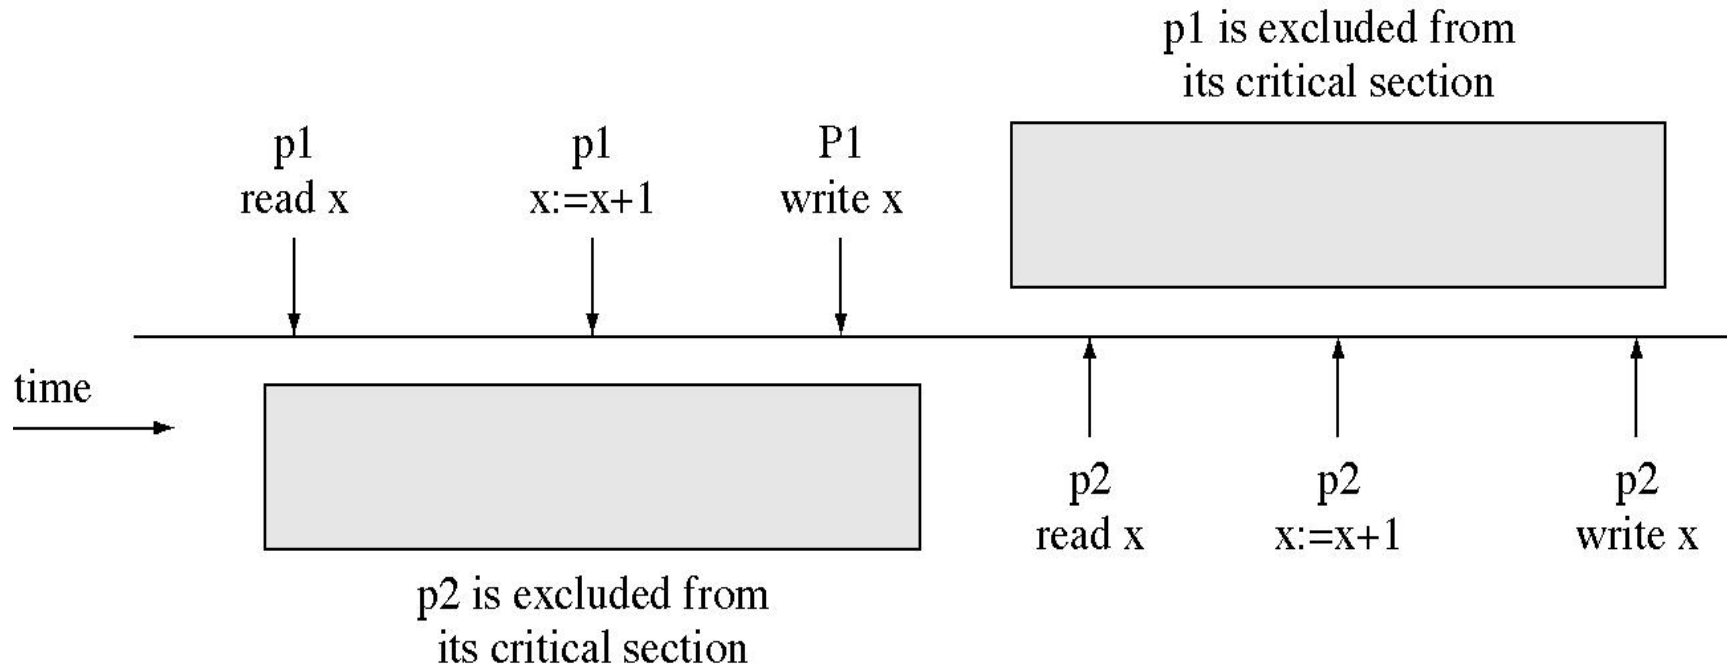
Store R1, balance

. . .

Race condition updating a variable (2)



Critical section to prevent a race condition



- Multiprogramming allows logical parallelism, uses devices efficiently but we lose correctness when there is a race condition.
- So we forbid logical parallelism inside critical section so we lose some parallelism but we regain correctness.

