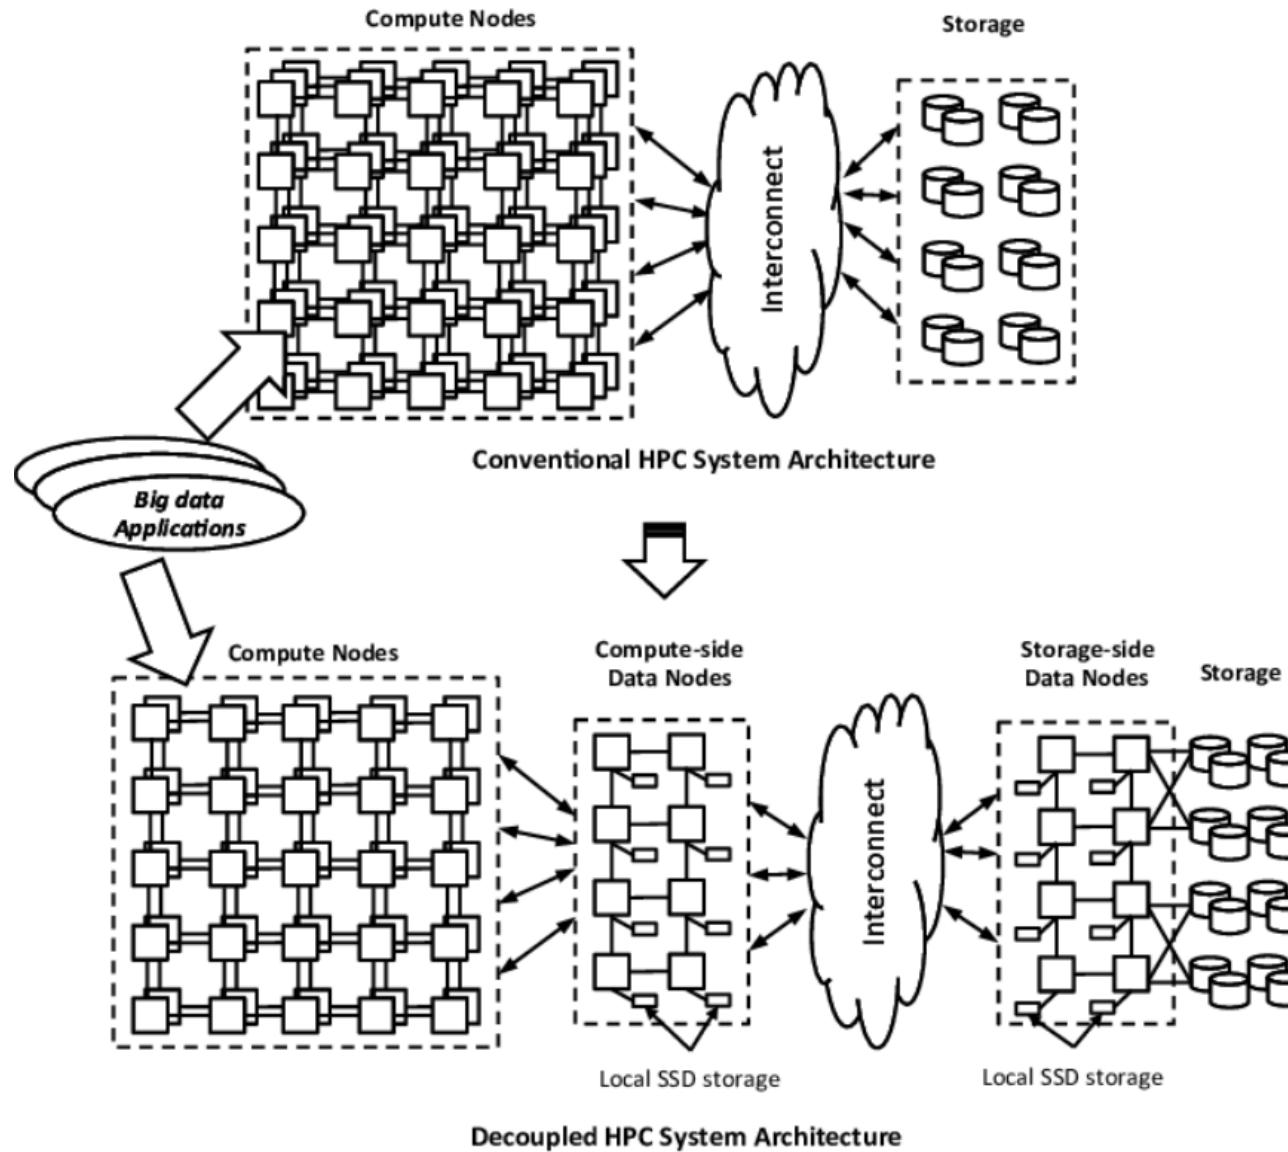


OPERATING SYSTEMS & PARALLEL COMPUTING

Big Data Platform

Decoupled High Performance Computing System Architecture



BigData Overview

Big Data

- History
- Google
 - Started with My SQL for search engine (scalability was major issue)
 - Developed solution from scratch
 - Distributed file system – GFS
 - Distributed processing – Map Reduce
 - Big Table

What is Big Data?

Extremely large data sets that may be analyzed computationally to reveal patterns, trends, and associations, especially relating to human behavior and interactions - Oxford Dictionary

When do I need Big Data?

Volume



Huge amounts of data

Velocity



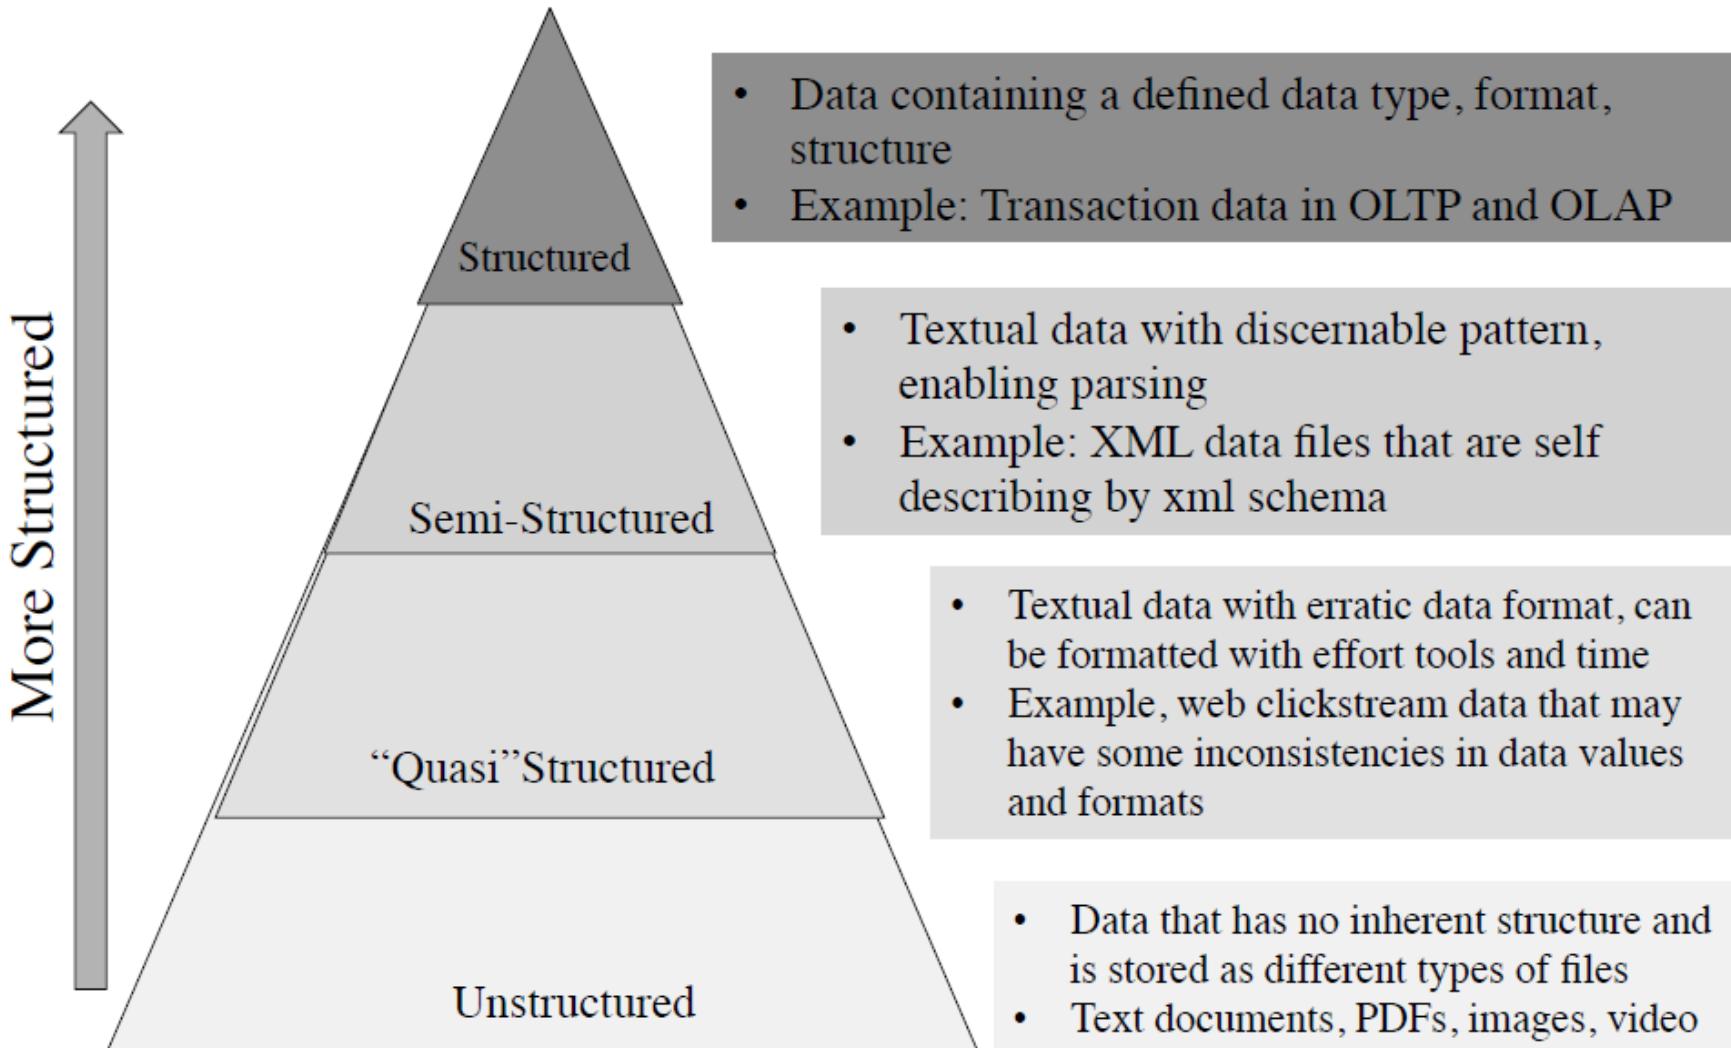
Speed at which data needs to be processed

Variety

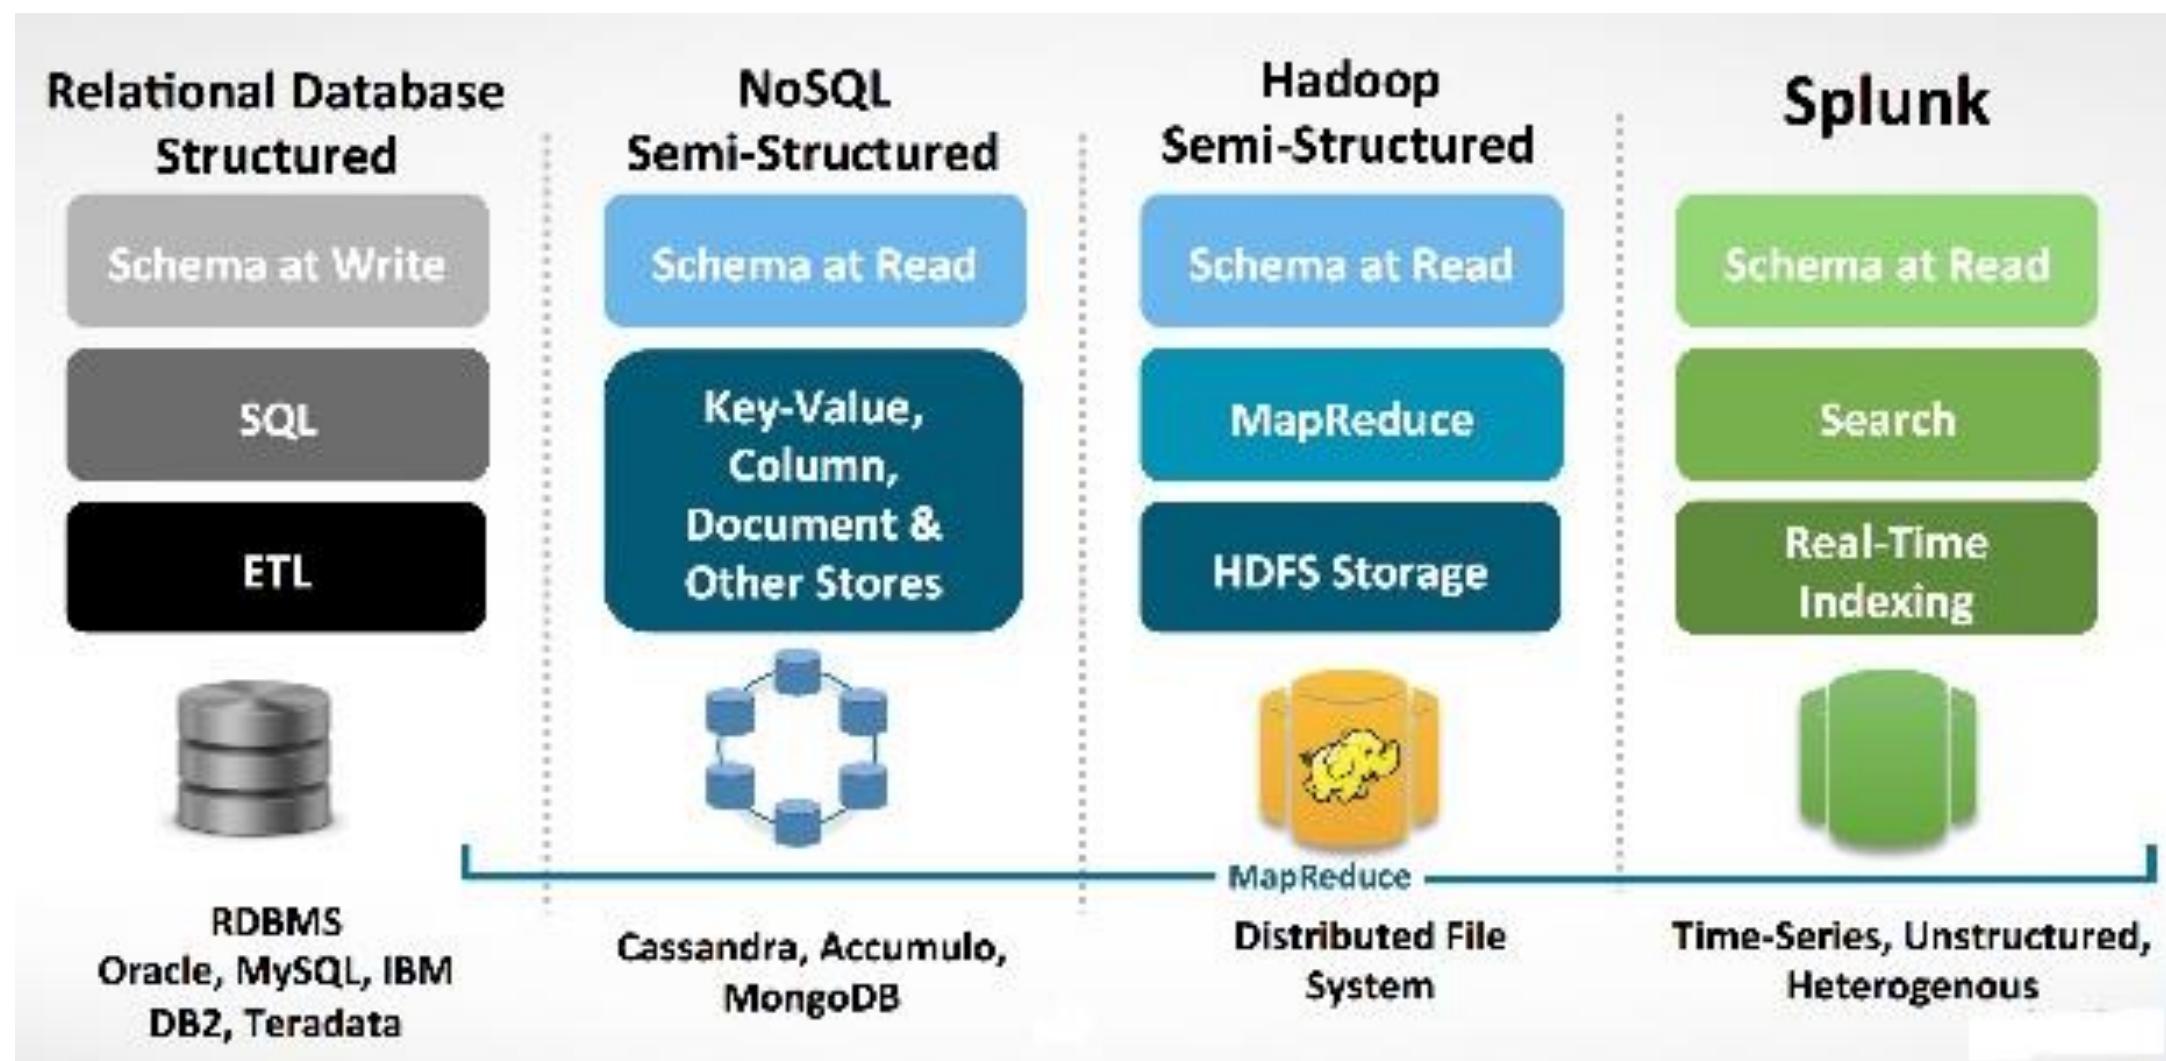


Structured and Semi Structured

Big Data Characteristics: Data Structures



Big Data Technologies



Why Hadoop is Invented?

- Storage for Large Datasets
 - The conventional RDBMS is incapable of storing huge amounts of Data. The cost of data storage in available RDBMS is very high. As it incurs the cost of hardware and software both.
- Handling data in different formats
 - The RDBMS is capable of storing and manipulating data in a structured format. But in the real world we have to deal with data in a structured, unstructured and semi-structured format.
- Data getting generated with high speed:
 - The data is oozing out in the order of tera to peta bytes daily. Hence we need a system to process data in real-time within a few seconds. The traditional RDBMS fail to provide real-time processing at great speeds.

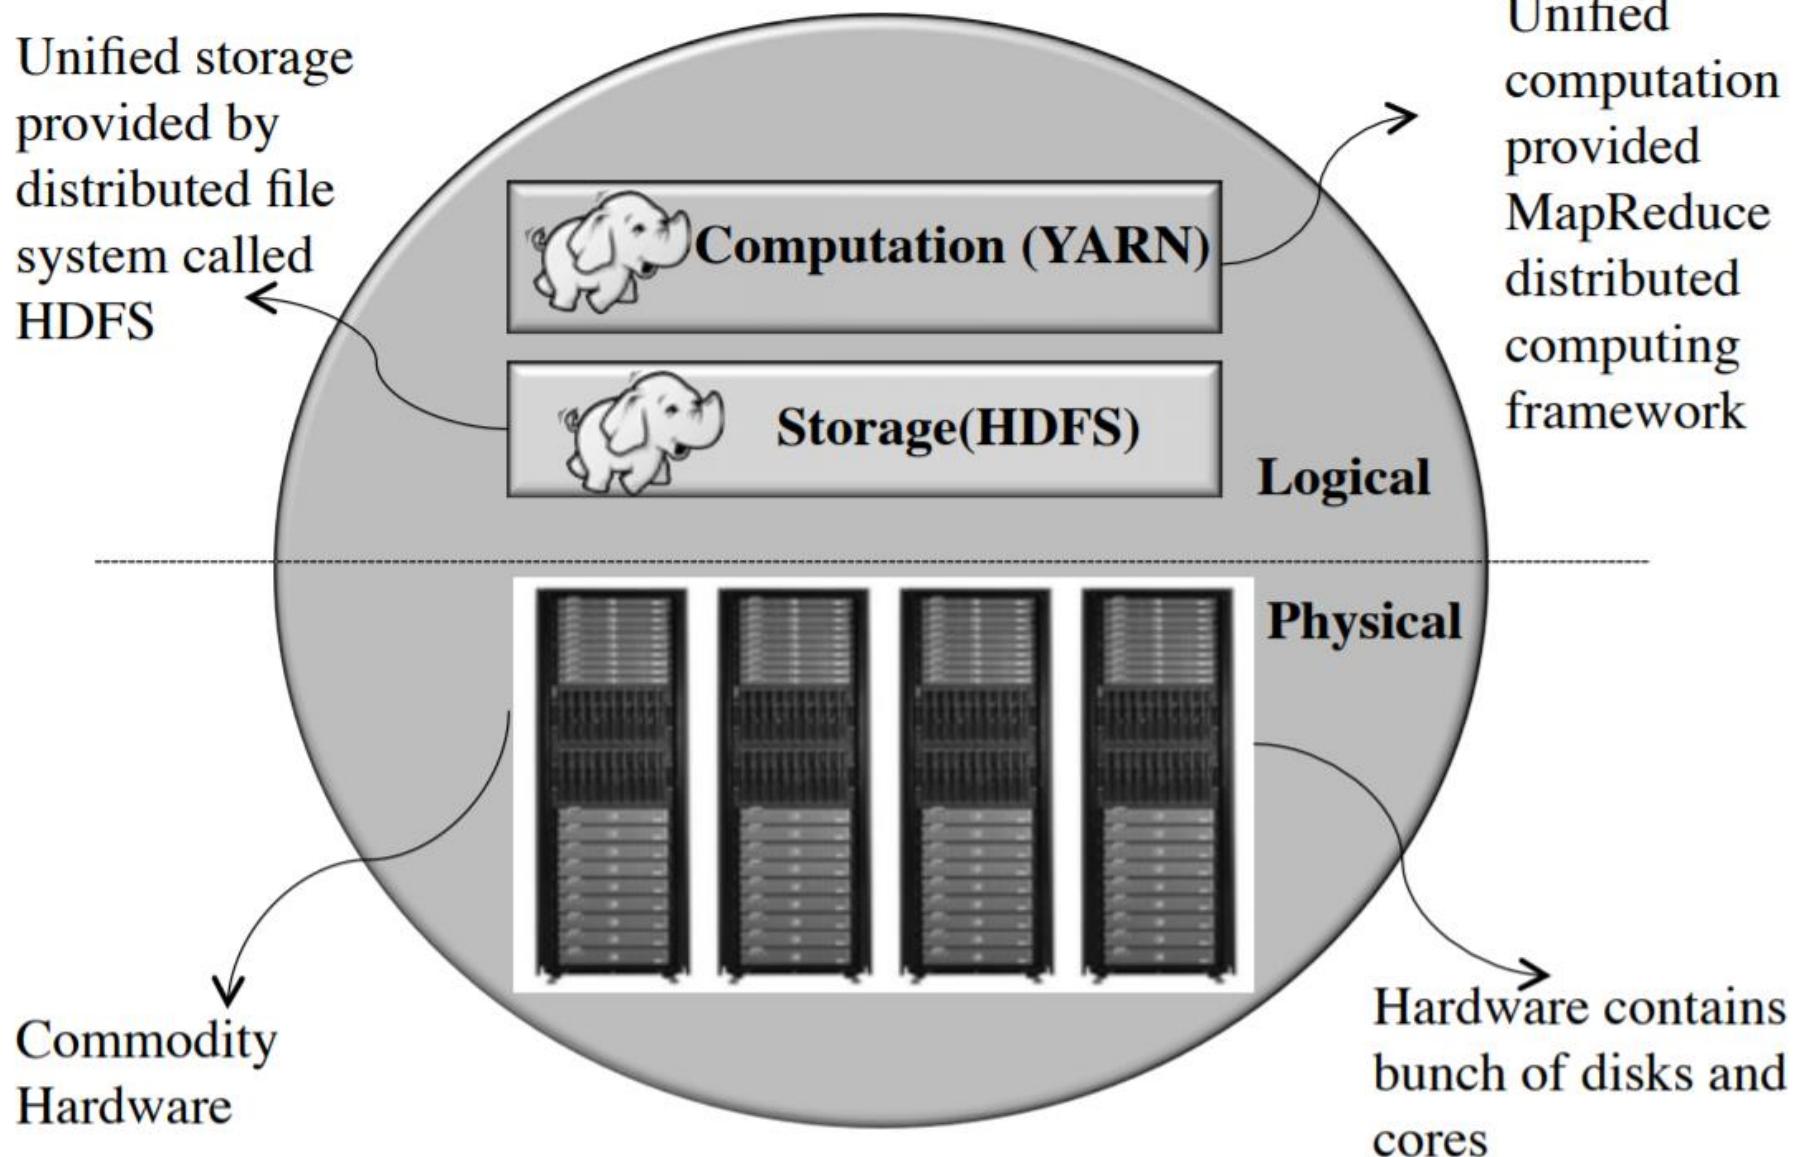
Big Data to Hadoop

- Characteristics - Volume, Variety and Velocity
- Batch
 - Hadoop with Map Reduce
 - GFS -> HDFS
 - Map Reduce -> Hadoop Map Reduce
- Operational but not Transactional
 - NoSQL
 - Google Big Table -> HBase

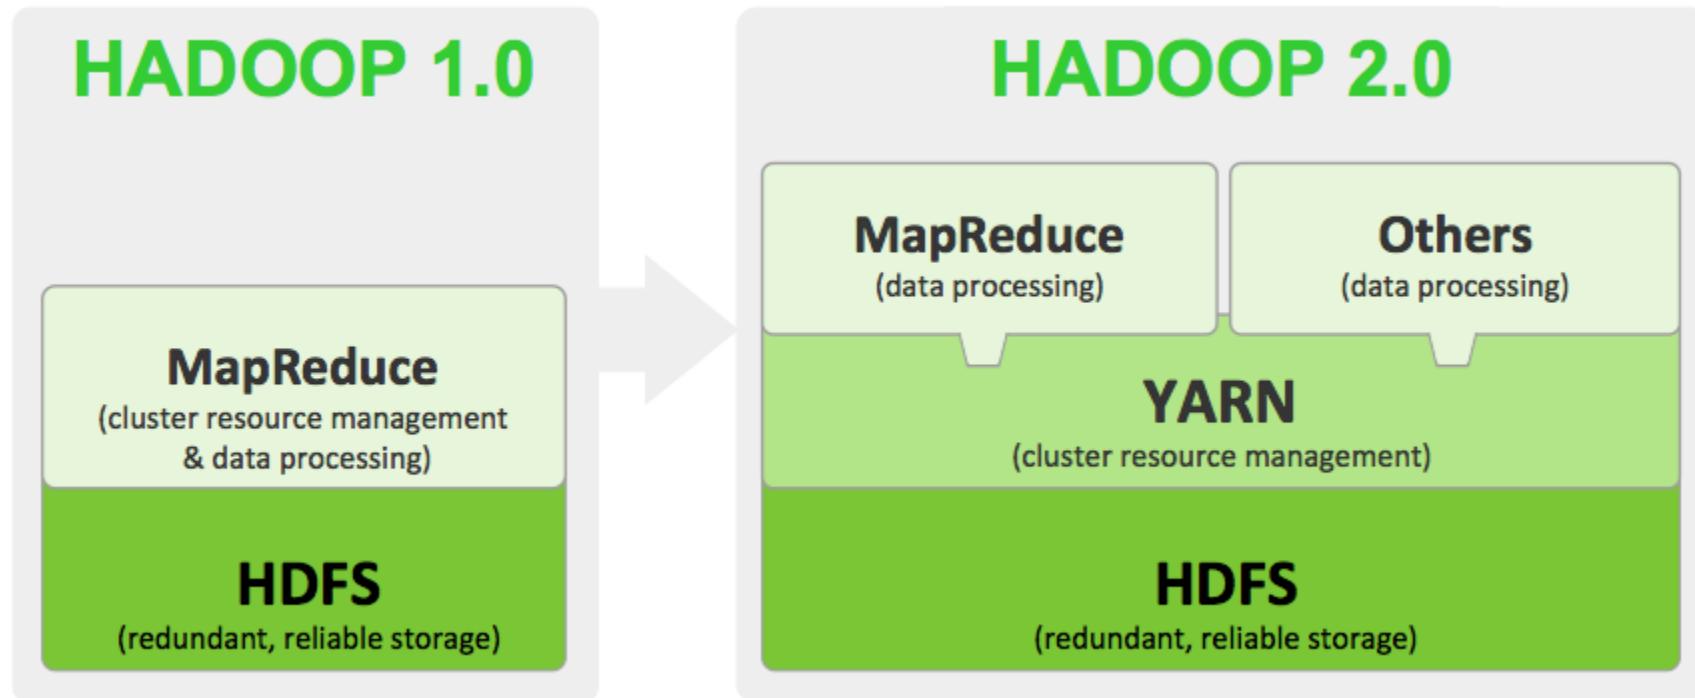
Hadoop core components

- **Hadoop Distributed File System (HDFS)** – It is the storage layer of Hadoop.
- **Map-Reduce** – It is the data processing layer of Hadoop.
- **YARN** – It is the resource management layer of Hadoop.

Hadoop – Big Picture



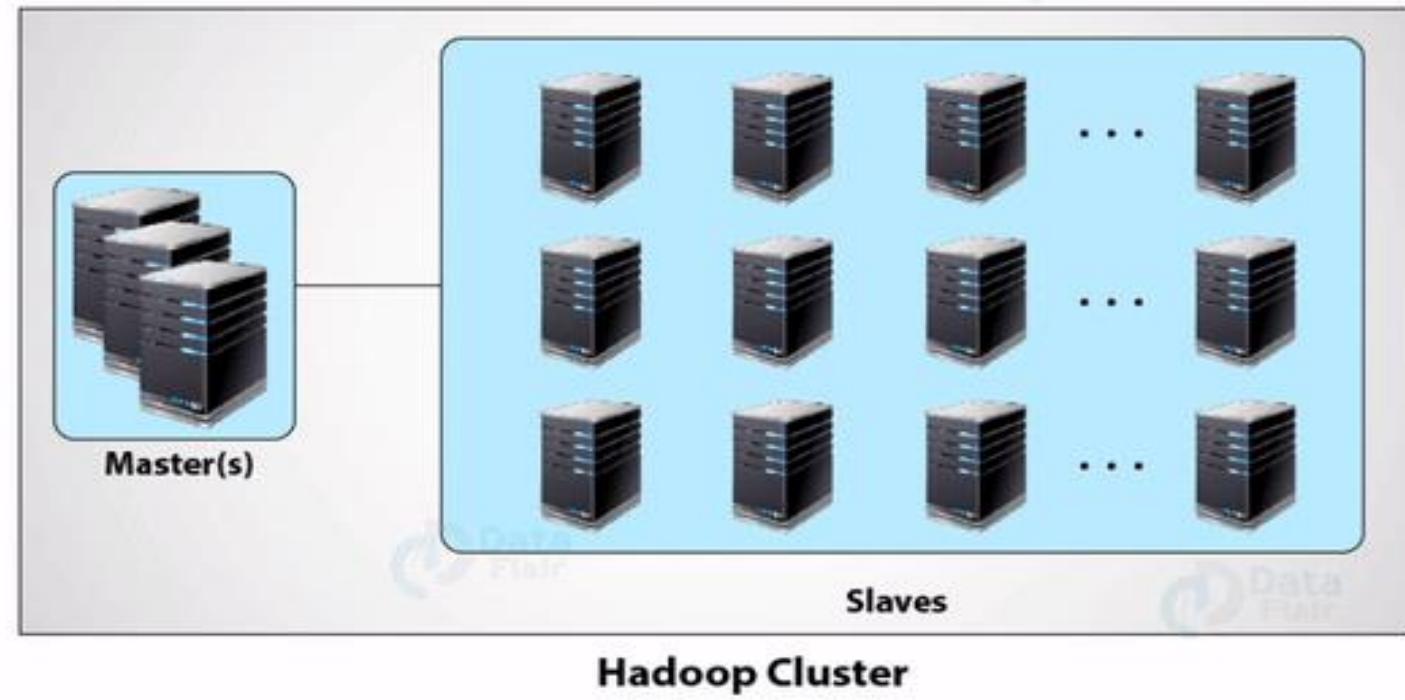
Hadoop Versions



Data Storage in HDFS

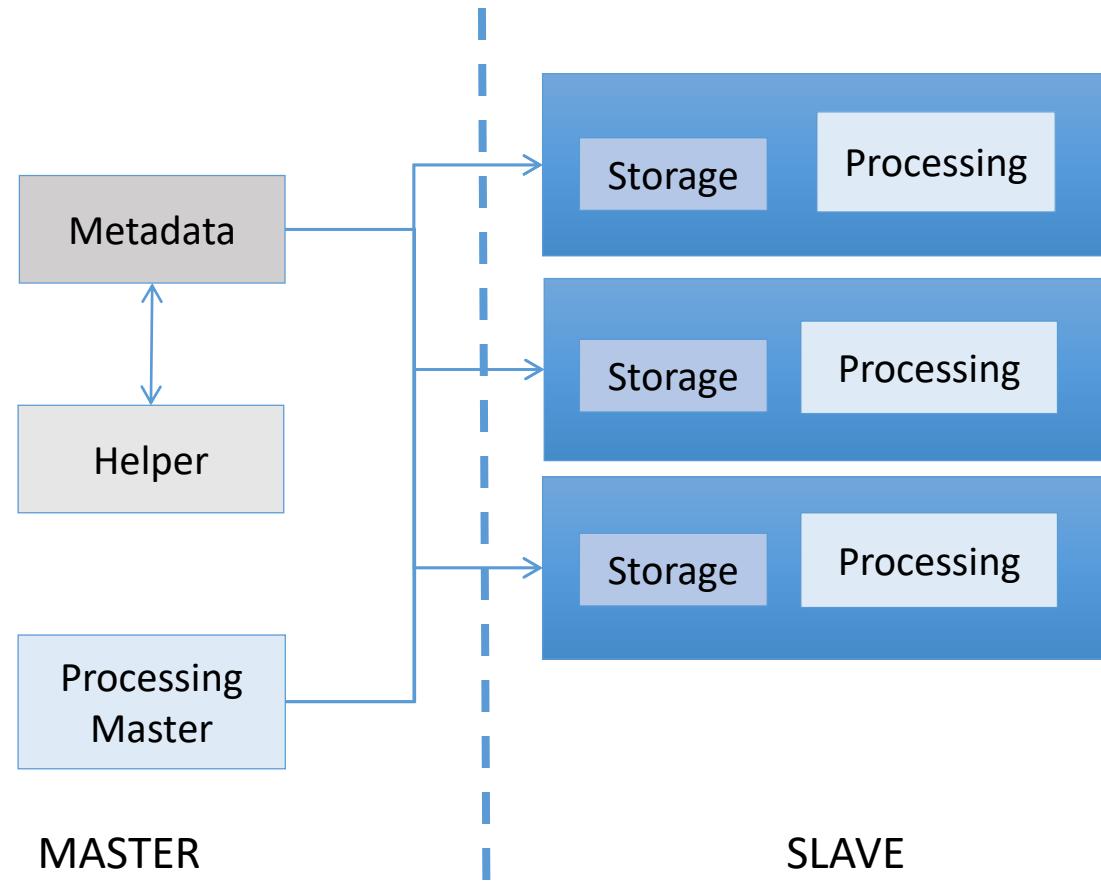


Data Storage in HDFS



Hadoop Architecture

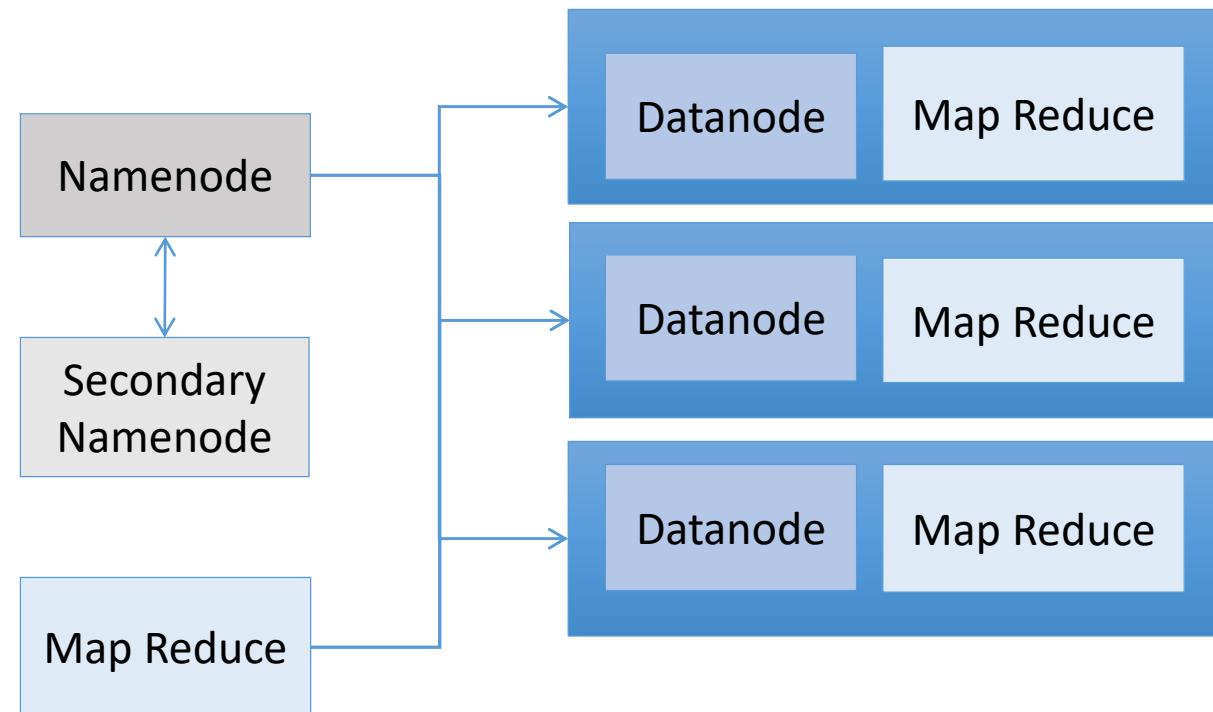
Master is a **high-end machine** where as slaves are inexpensive computers. The Big Data files get divided into the **number of blocks**. Hadoop stores these blocks in a distributed fashion on the cluster of slave nodes. On the master, we have **metadata stored**.



HDFS Architecture

- HDFS has **two daemons** running for it. They are :
- **NameNode** : NameNode performs following functions –
 - NameNode Daemon runs on the master machine.
 - It is responsible for maintaining, monitoring and managing DataNodes.
 - It records the metadata of the files like the location of blocks, file size, permission, hierarchy etc.
 - Namenode captures all the changes to the metadata like deletion, creation and renaming of the file in edit logs.
 - It regularly receives heartbeat and block reports from the DataNodes.
- **DataNode**: The various functions of DataNode are as follows –
 - DataNode runs on the slave machine.
 - It stores the actual business data.
 - It serves the read-write request from the user.
 - DataNode does the ground work of creating, replicating and deleting the blocks on the command of NameNode.
 - After every 3 seconds, by default, it sends heartbeat to NameNode reporting the health of HDFS.

HDFS Architecture



Hadoop Processing

- It is the data processing layer of Hadoop. It processes data in two phases:
- **Map Phase-** This phase applies business logic to the data. The input data gets converted into key-value pairs.
- **Reduce Phase-** The Reduce phase takes as input the output of Map Phase. It applies aggregation based on the key of the key-value pairs.
- Versions:
 - MRv1/Classic
 - MRv2/YARN

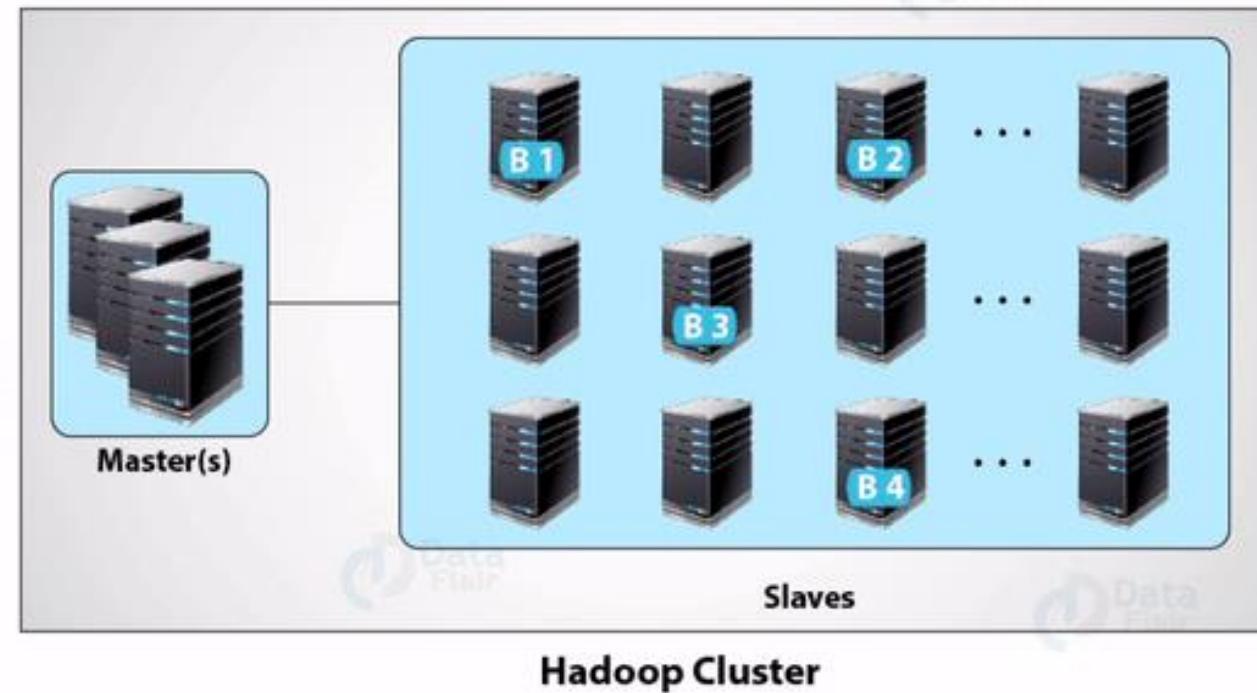
How MapReduce Works?



How MapReduce works



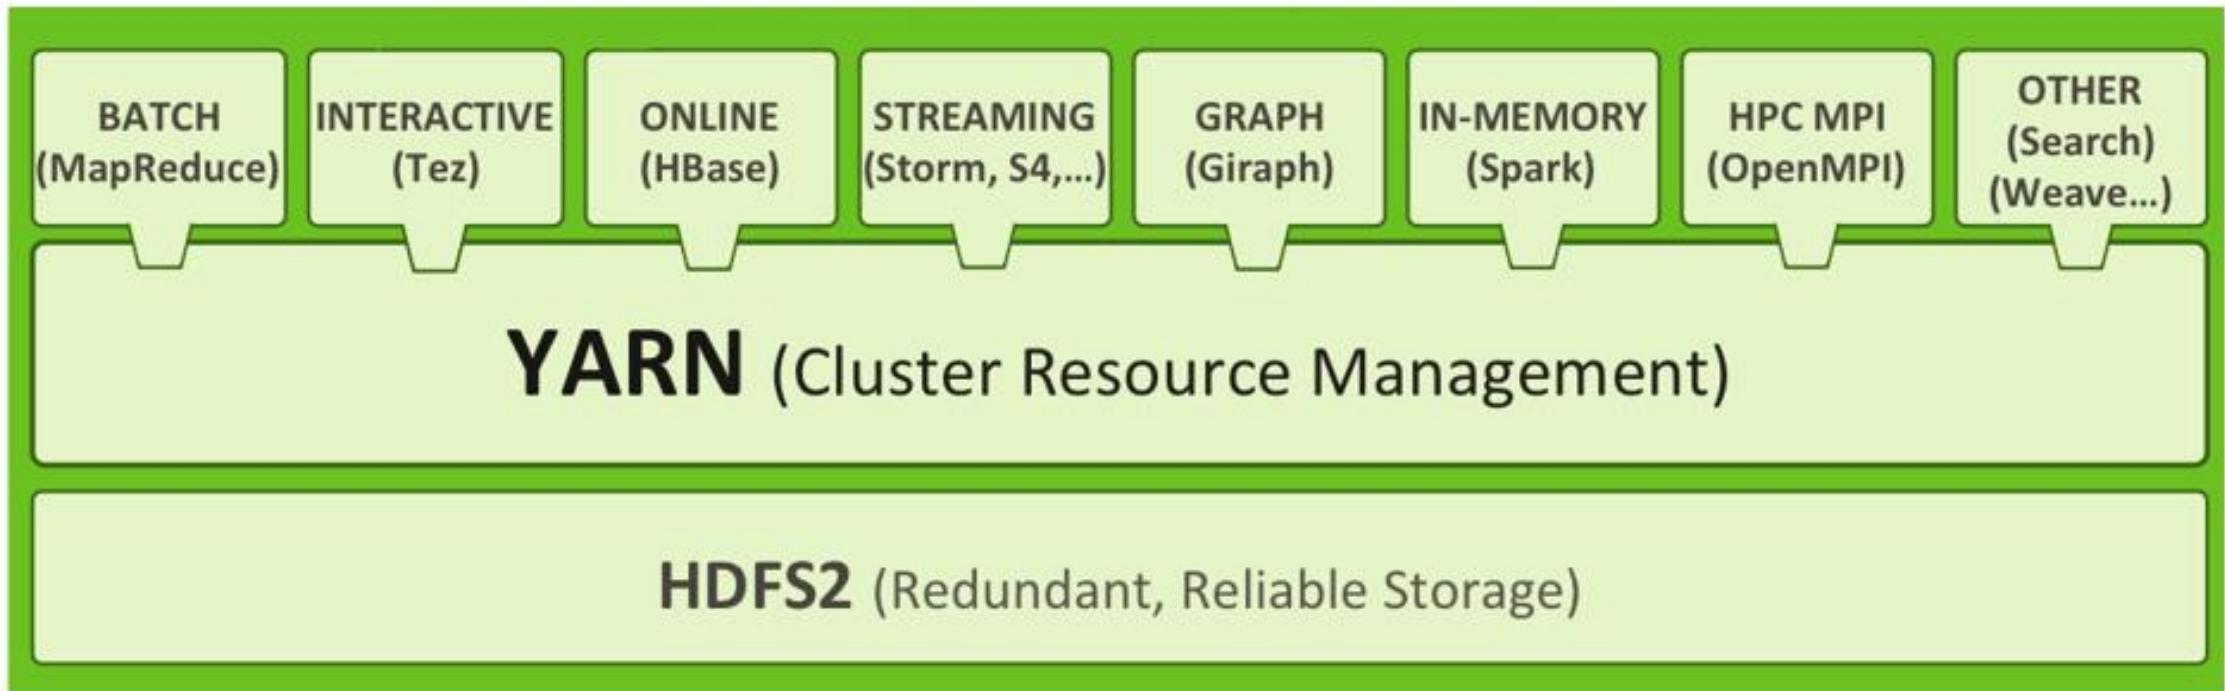
User



Map-Reduce works in the following way

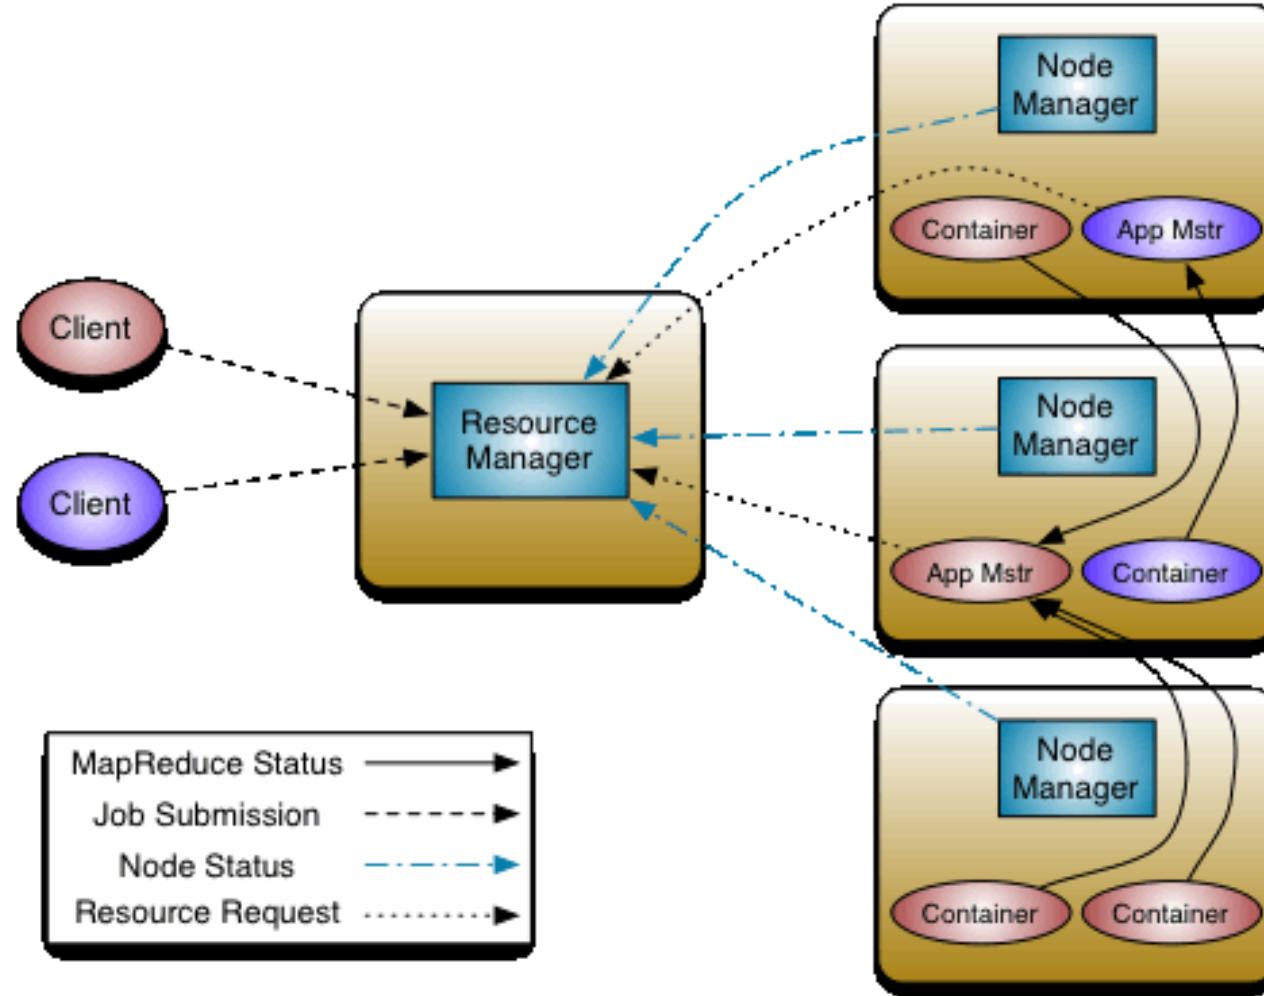
- The client specifies the file for input to the Map function. It splits it into tuples
- Map function defines key and value from the input file. The output of the map function is this key-value pair.
- MapReduce framework sorts the key-value pair from map function.
- The framework merges the tuples having the same key together.
- The reducers get these merged key-value pairs as input.
- Reducer applies aggregate functions on key-value pair.
- The output from the reducer gets written to HDFS.

YARN Architecture



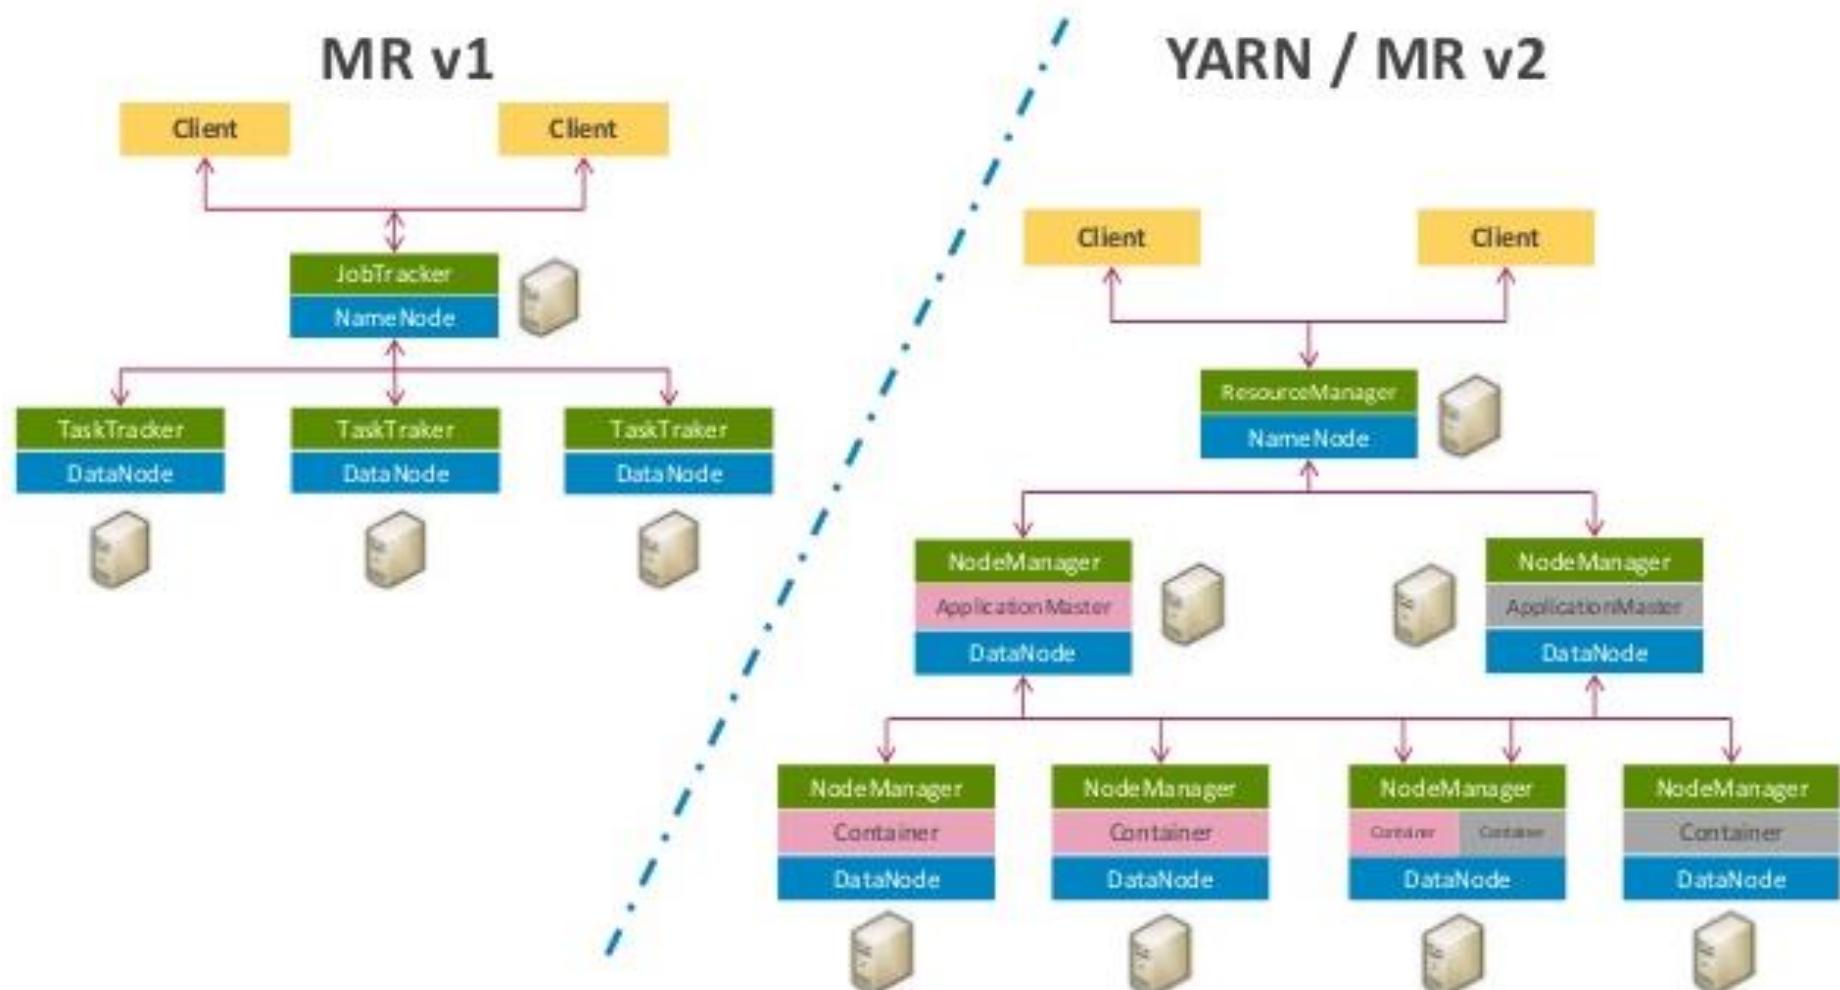
| *Apache Hadoop coupled with various frameworks to support data processing for all types of data.*

YARN Architecture MRv2

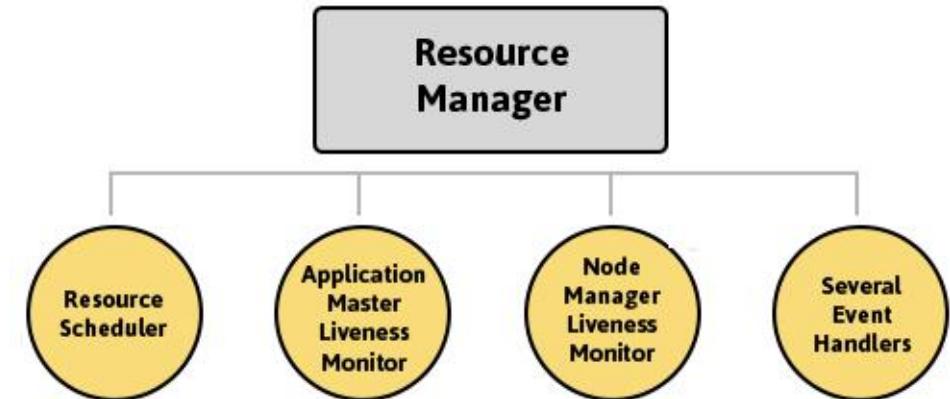


MR vs YARN Architecture

- YARN
 - Yet Another Resource Negotiator
- MR
 - MapReduce

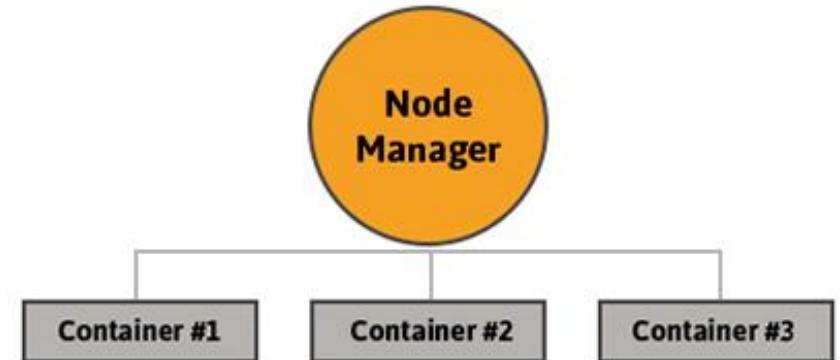


YARN Resource Manager



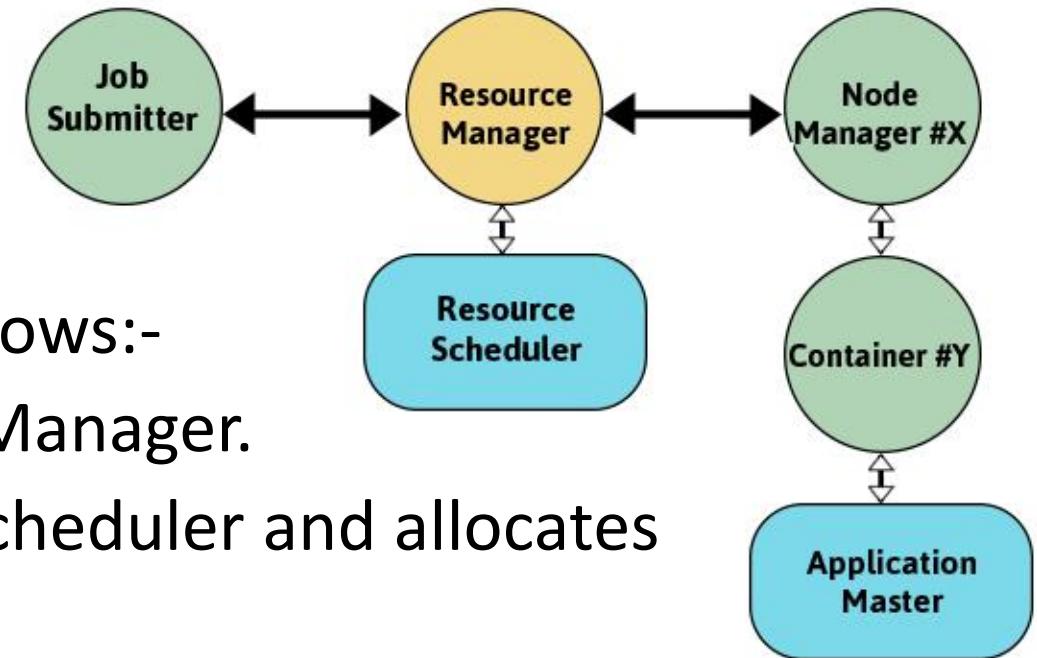
- Resource Manager runs on the master node.
- It knows where the location of slaves (Rack Awareness).
- It is aware about how much resources each slave have.
- Resource Scheduler is one of the important service run by the Resource Manager.
- Resource Scheduler decides how the resources get assigned to various tasks.
- Application Manager is one more service run by Resource Manager.
- Application Manager negotiates the first container for an application.
- Resource Manager keeps track of the heart beats from the Node Manager.

YARN Node Manager



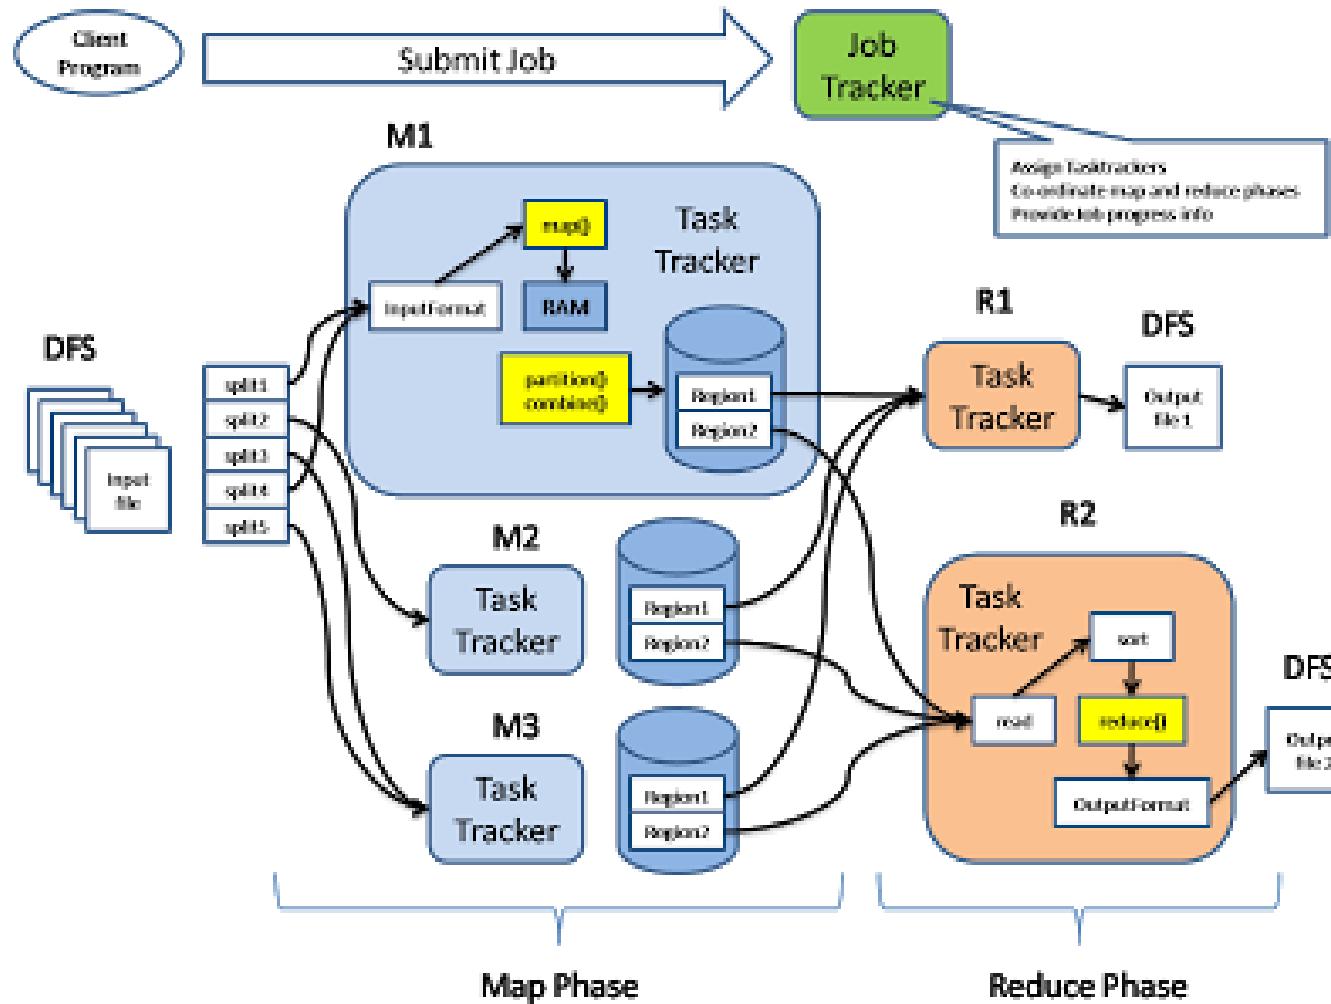
- It runs on slave machines.
- It manages containers. Containers are nothing but a fraction of Node Manager's resource capacity
- Node manager monitors resource utilization of each container.
- It sends heartbeat to Resource Manager.

YARN Job Submitter



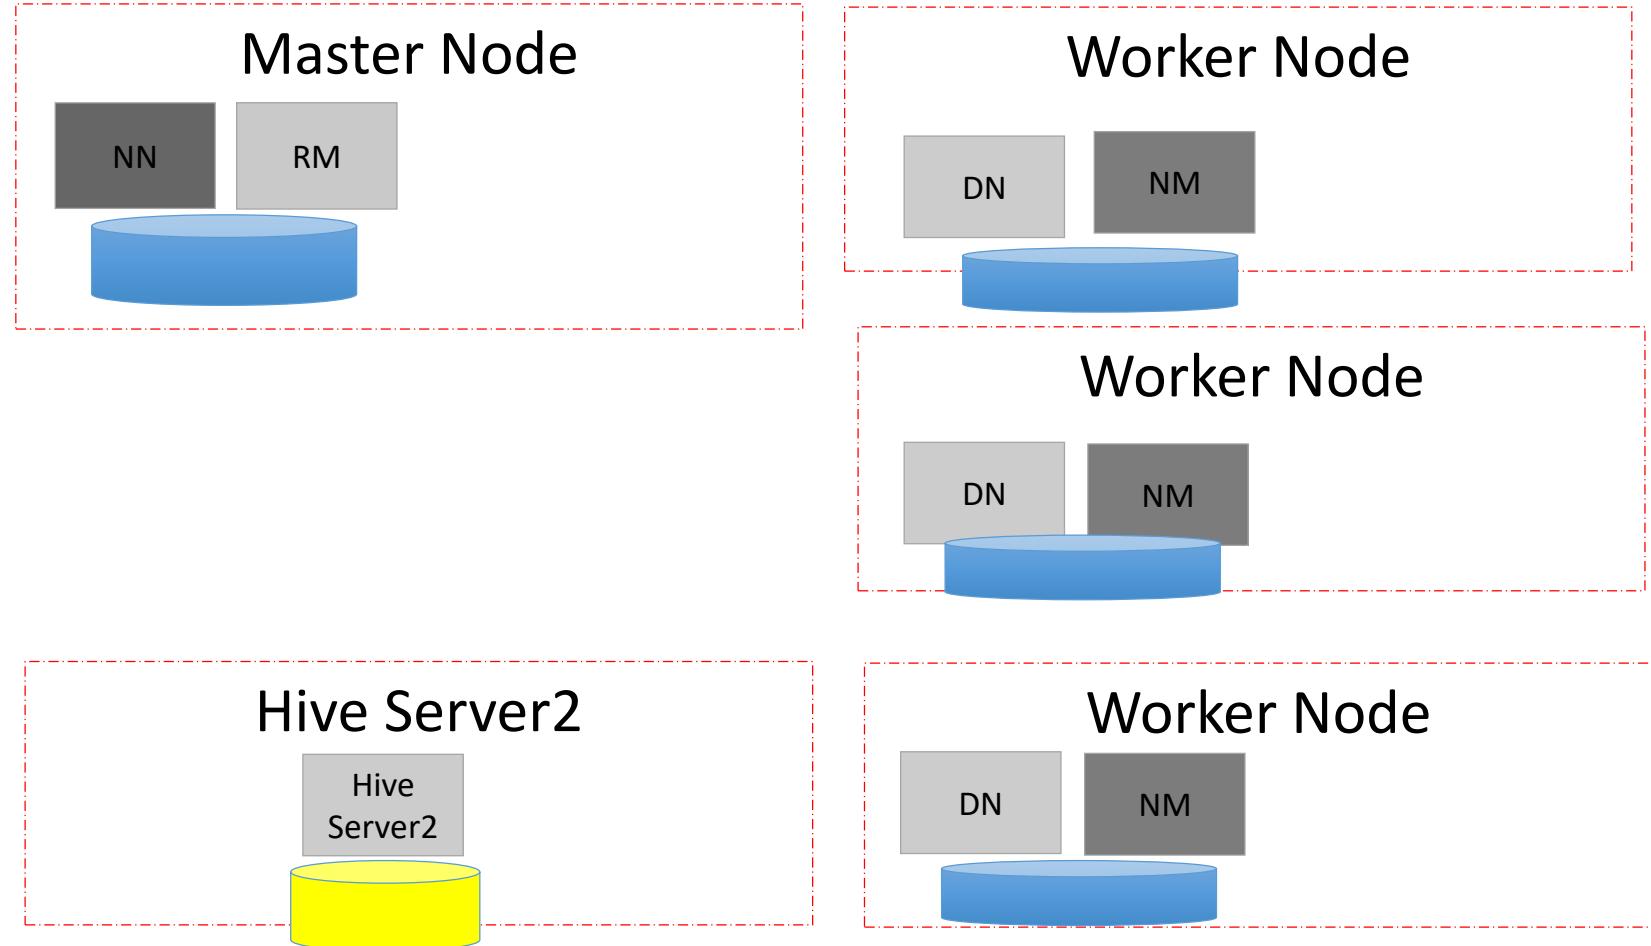
- The application startup process is as follows:-
- The client submits the job to Resource Manager.
- Resource Manager contacts Resource Scheduler and allocates container.
- Now Resource Manager contacts the relevant Node Manager to launch the container.
- Container runs Application Master.
- The basic idea of YARN was to split the task of resource management and job scheduling. It has one global Resource Manager and per-application Application Master. An application can be either one job or DAG of jobs.

YARN Job Execution

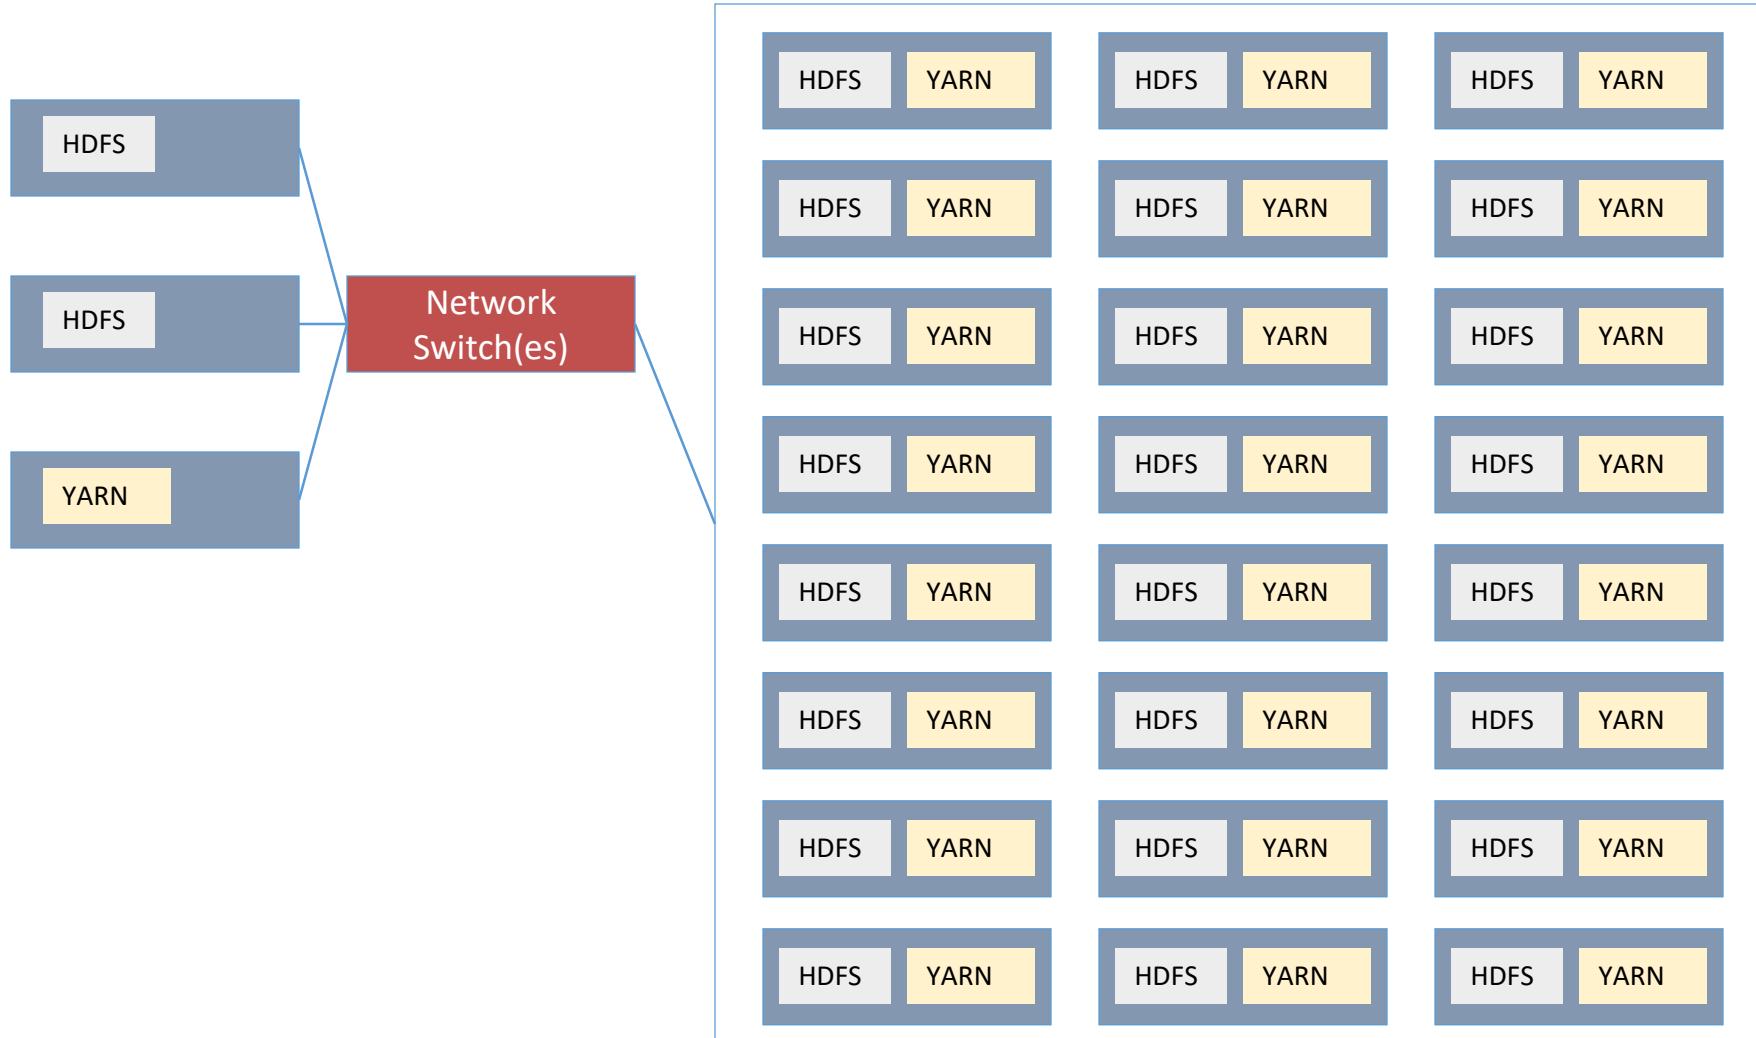


Hadoop itself is clustered....

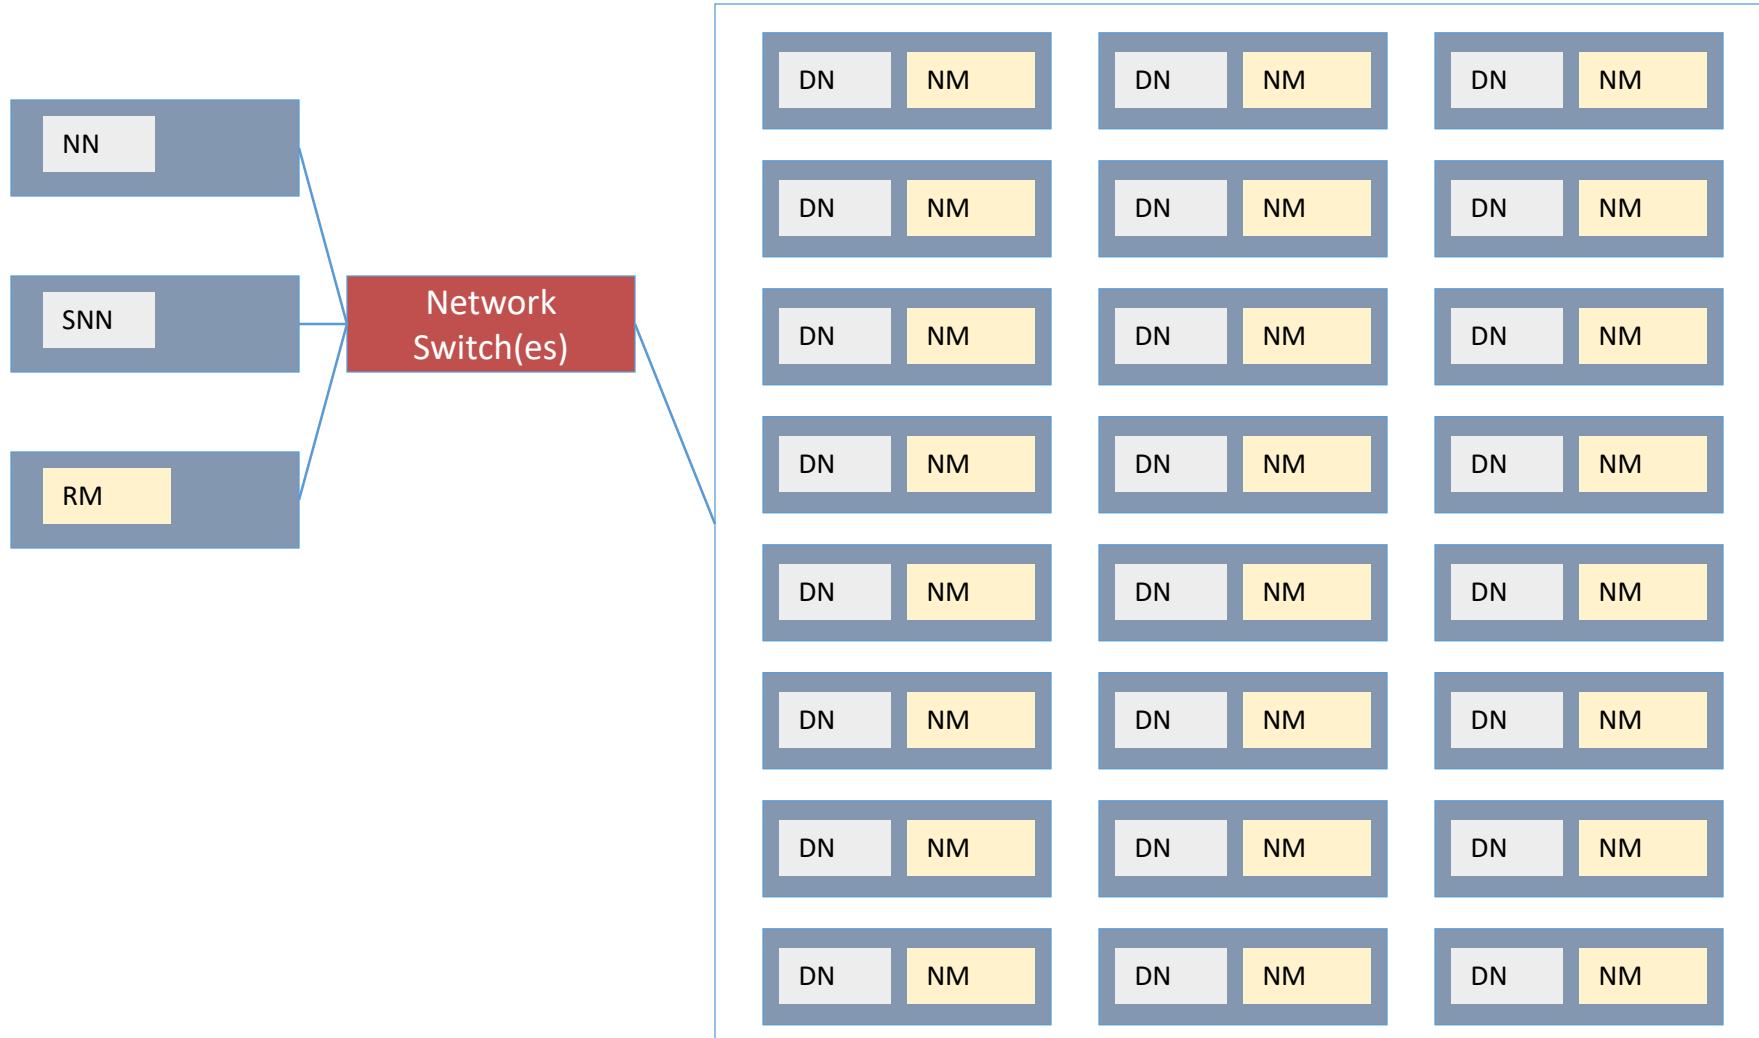
RM	YARN ResourceManager
NM	YARN NodeManager
NN	HDFS NameNode
DN	HDFS DataNode
	Data
	Metadata



Typical Hadoop Cluster



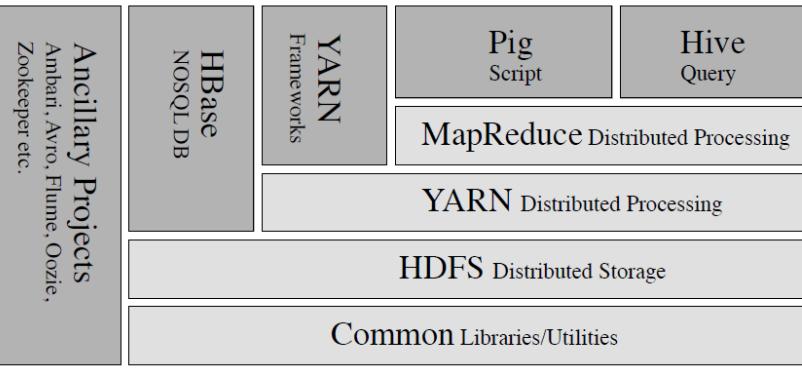
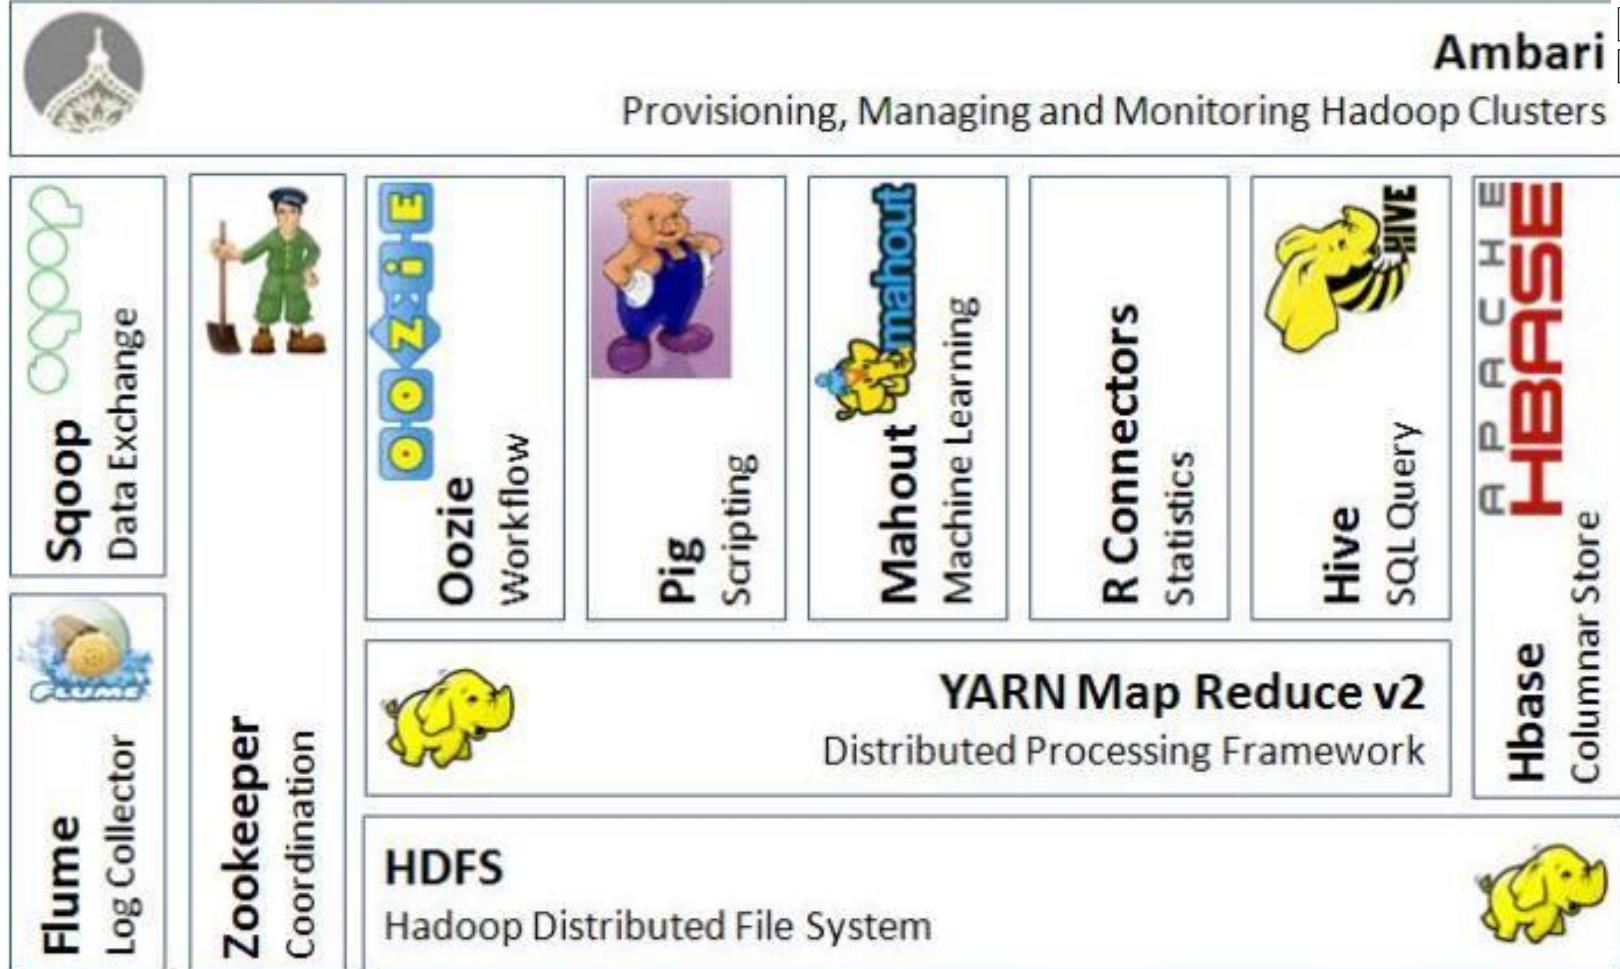
Typical Hadoop Cluster



Hadoop Technology Stack



Apache Hadoop Ecosystem



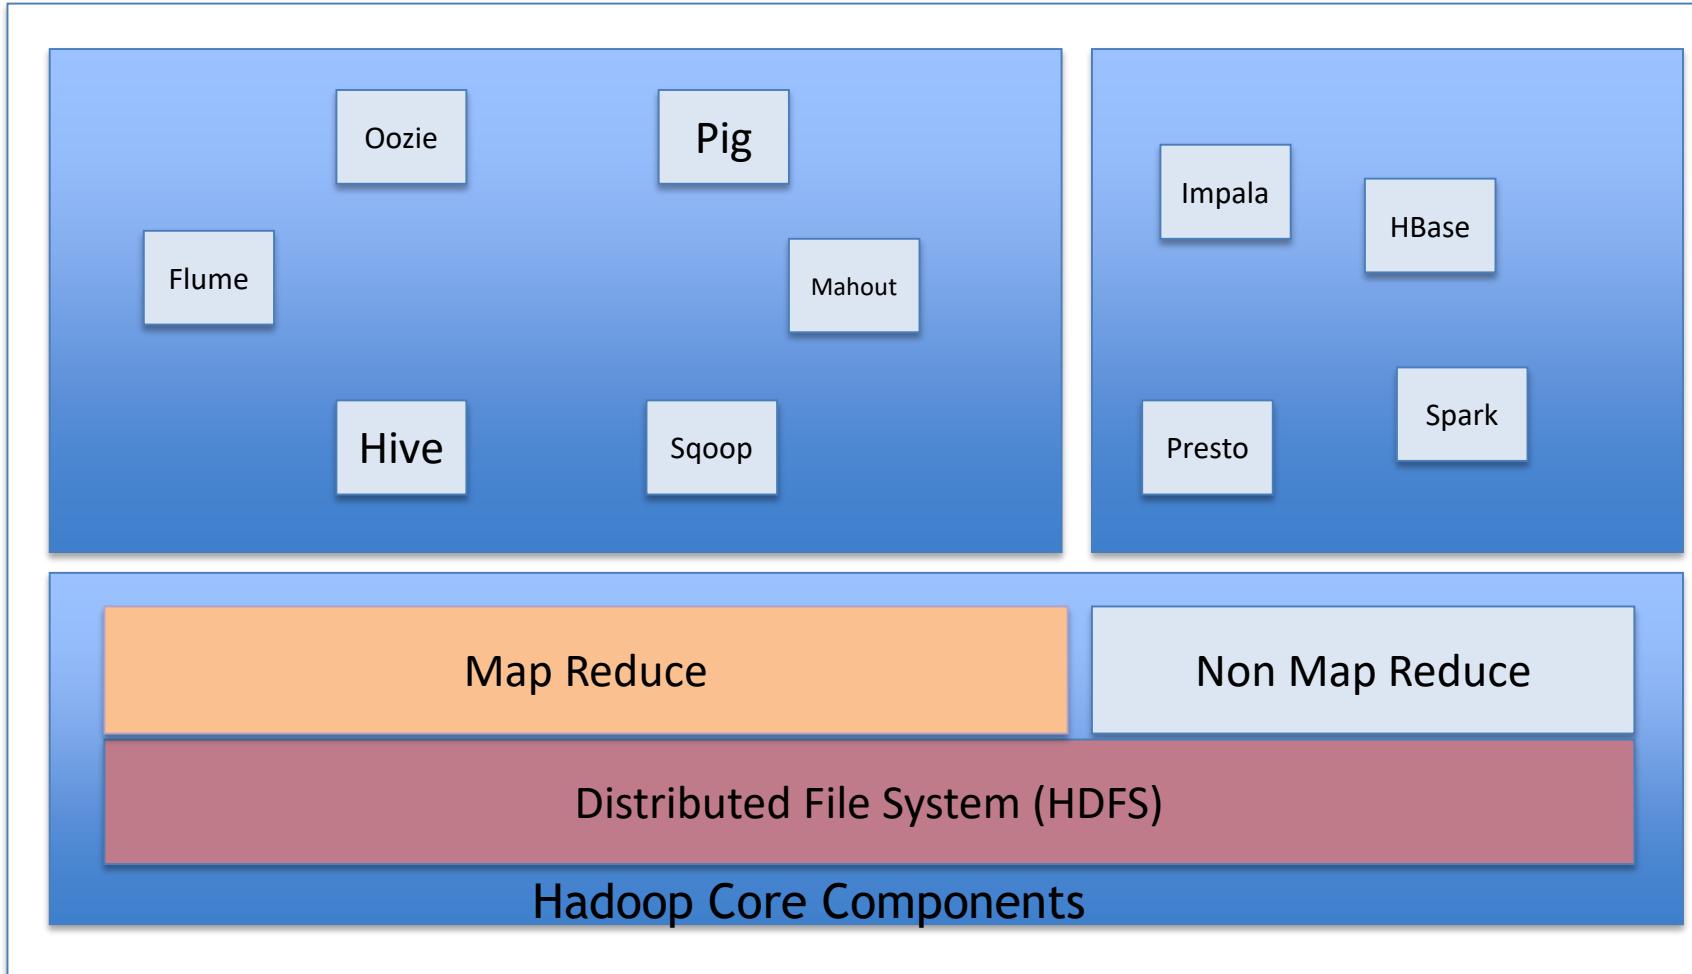
Application vs Service vs Instance



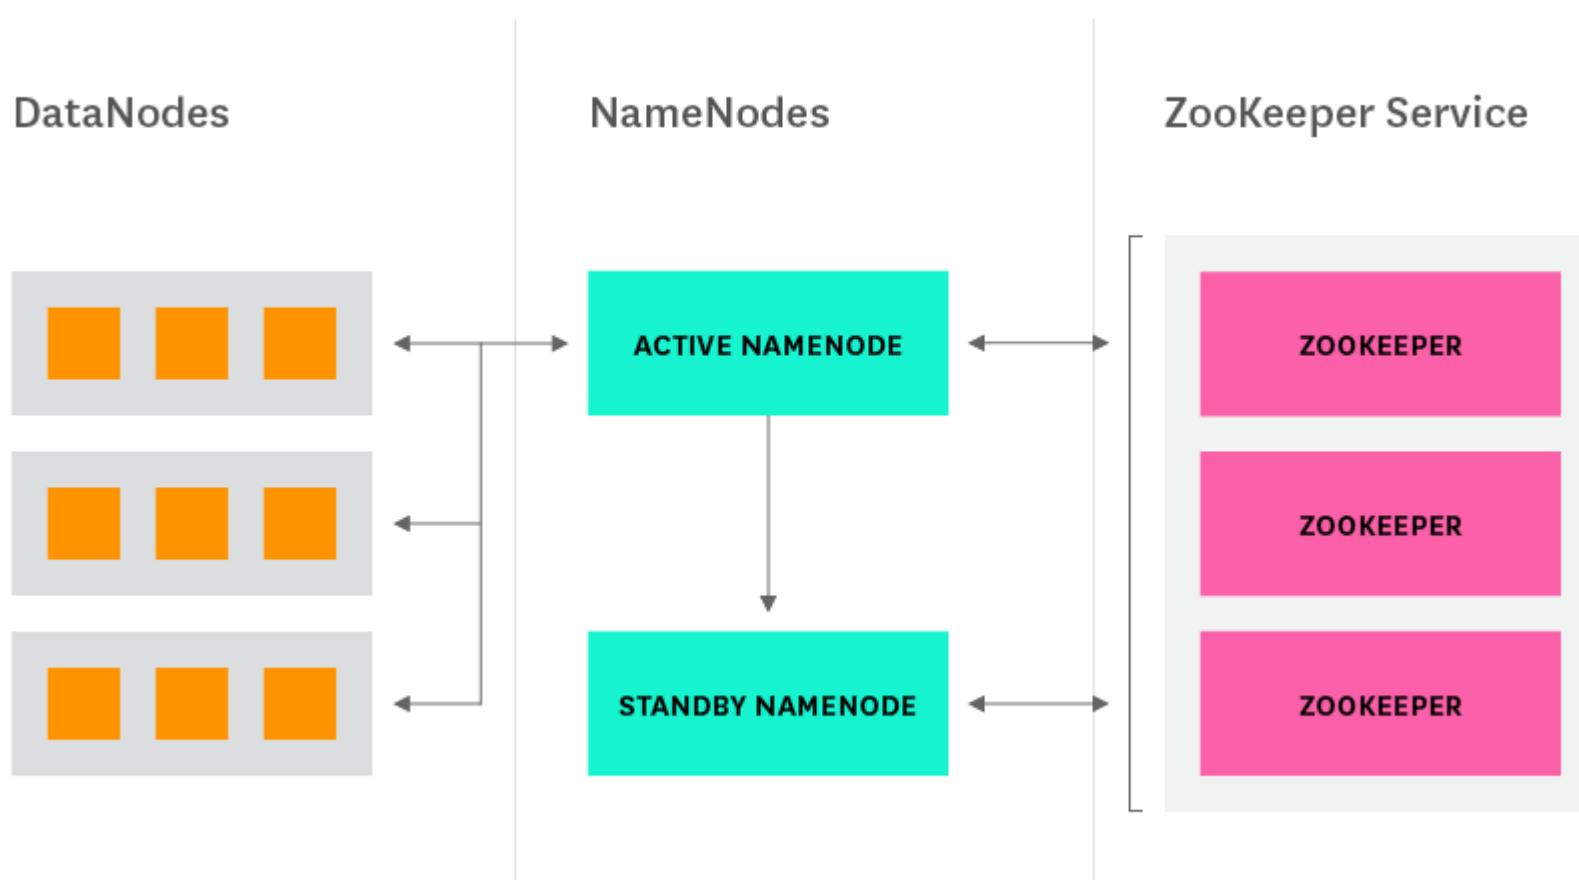
- For example, “Hadoop” is an application
- “Collection of Services” - NodeManager, DataNode, ResourceManager are application services
- The ResourceManager Service running on node host-1.example.com is an application service instance

Hadoop EcoSystem

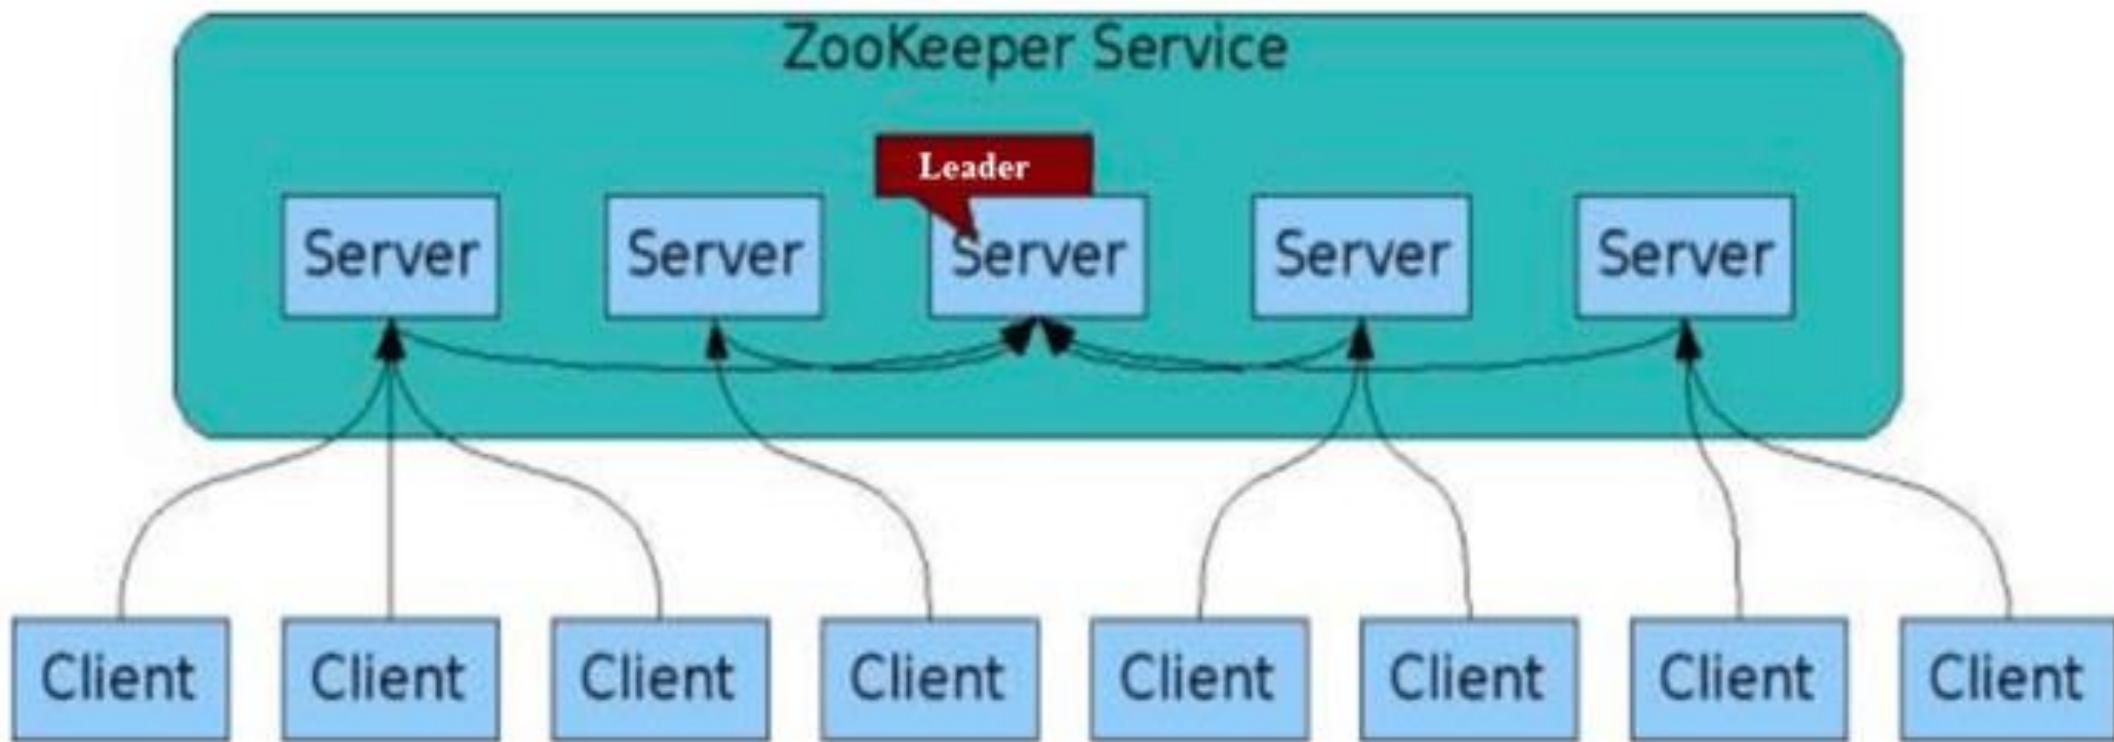
Hadoop Components



Hadoop High Availability



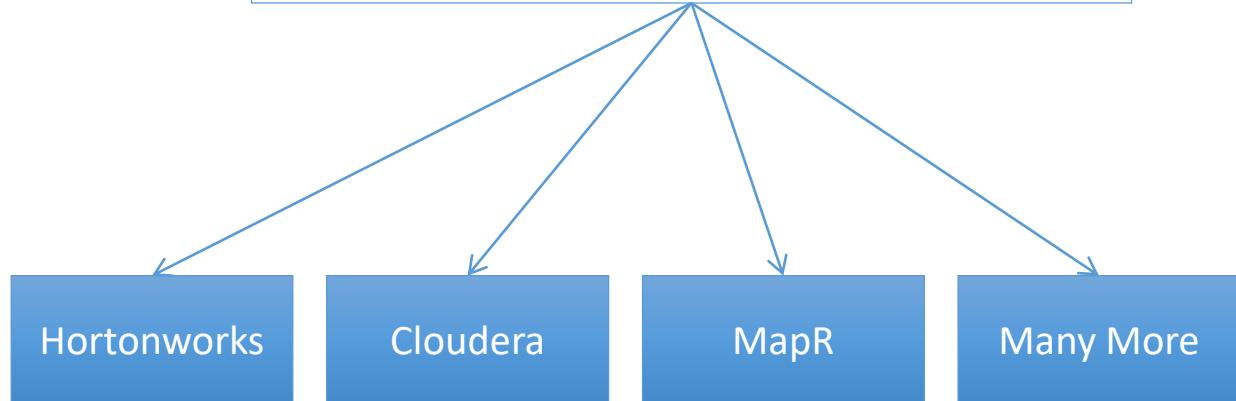
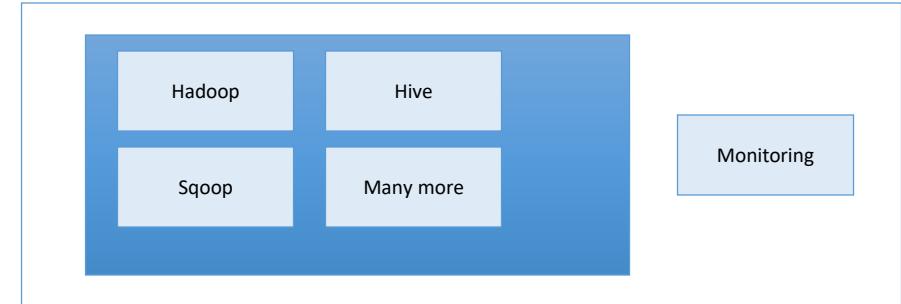
Hadoop Zookeeper Architecture



Concept of Leader / Followers

Hadoop Distributions

- **Apache** – Vanilla flavor, as the actual code is residing in Apache repositories.
- **Hortonworks** – Popular distribution in the industry.
- **Cloudera** – It is the most popular in the industry.
- **MapR** – It has rewritten HDFS and its HDFS is faster as compared to others.
- **IBM** – Proprietary distribution is known as Big Insights.



Hadoop Distributions

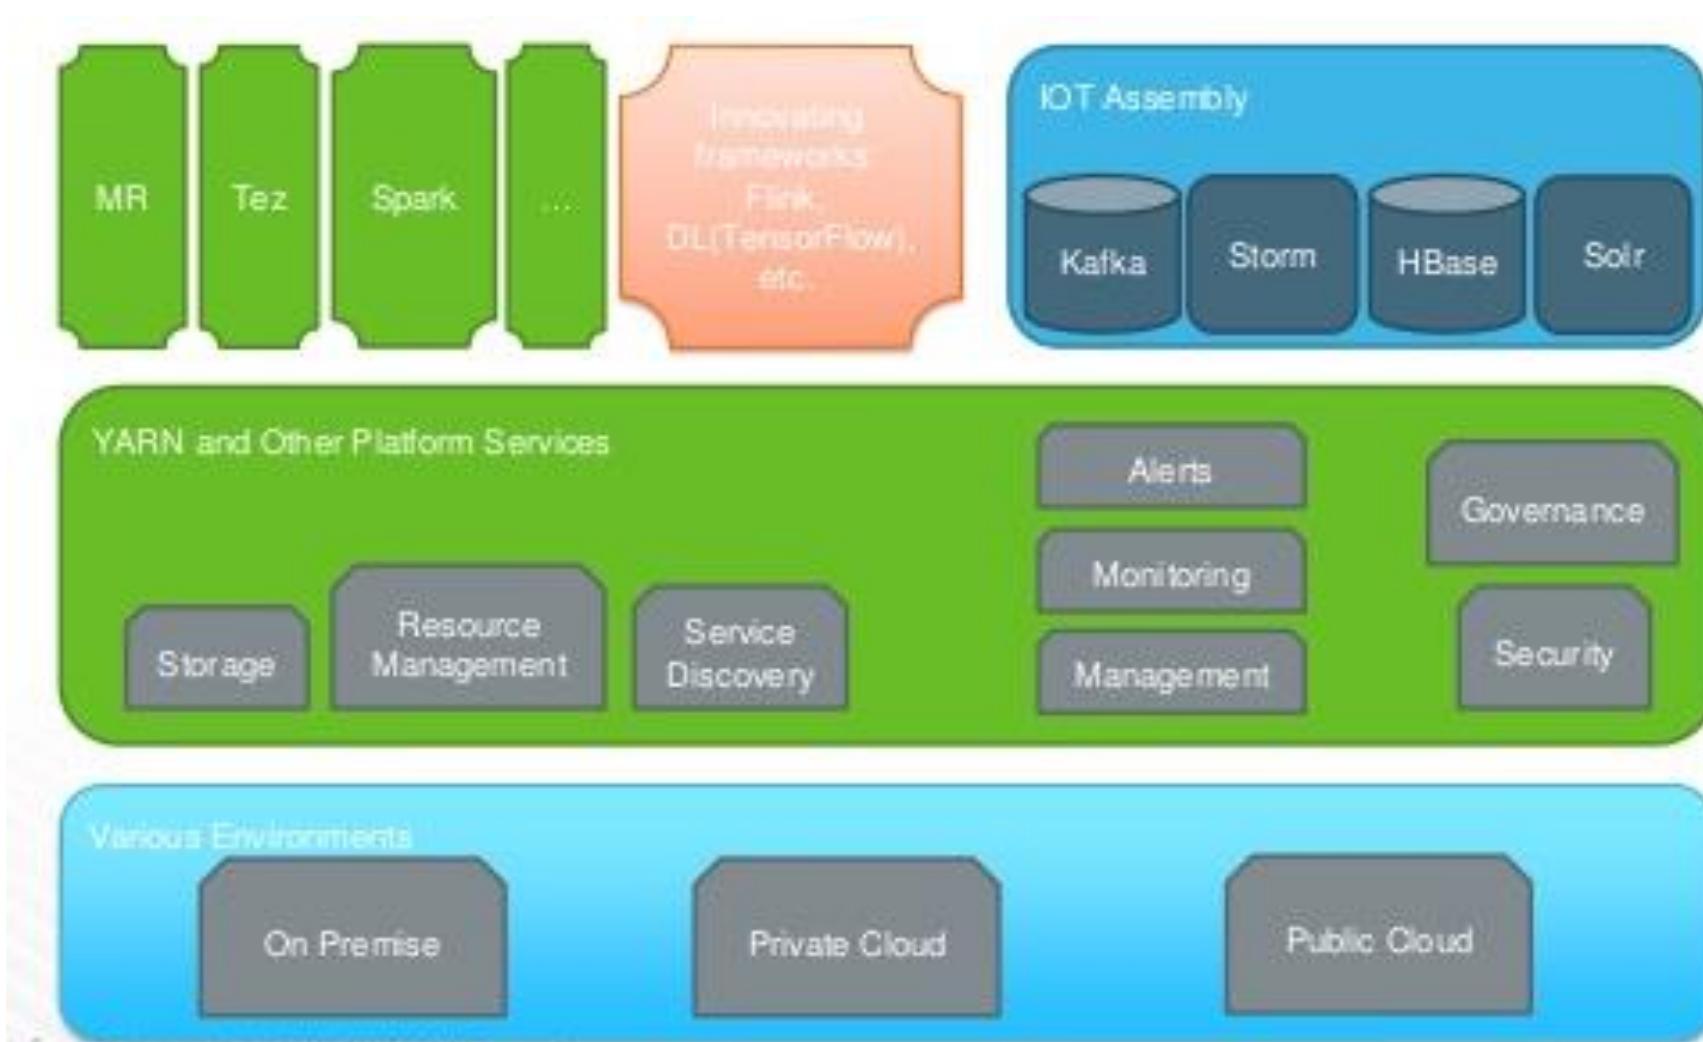
- Hadoop and eco system tools are Apache open source projects
- Cloudera, Hortonworks and other leading Hadoop based technology companies commit to these open source projects
- They provide training, support and services.
- Cloudera have proprietary monitoring tool developed for large hadoop clusters. It is free up to 50 nodes after which license fee needs to be paid.
- Hortonworks uses Ambari which is a open source monitoring tool. No license fee

Complete list of Hadoop Services?

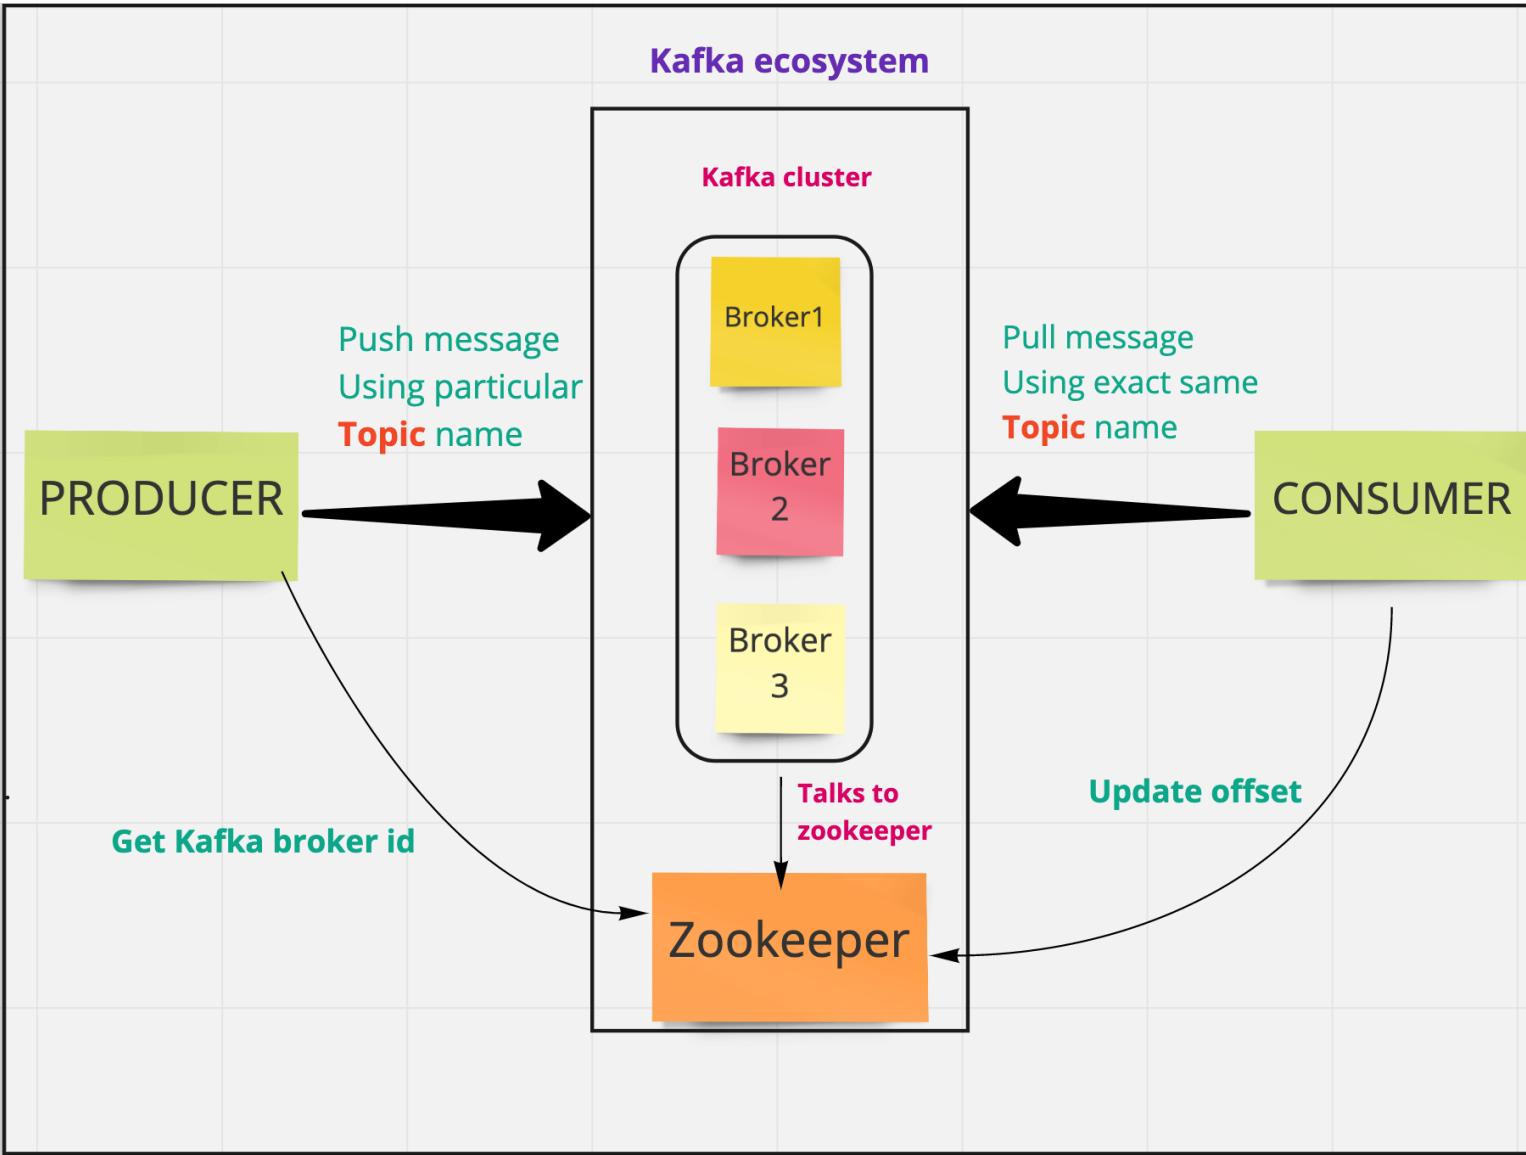
RM	YARN ResourceManager	SHS	Spark History Server	ISS	Impala State Store
NM	YARN NodeManager	Hue	Hue	ICS	Impala Catalog Server
NN	HDFS NameNode	OZ	Oozie	ID	Impala Daemon
DN	YARN DataNode	CM	Cloudera Manager	SS	Solr Server
JHS	Job History Server	DB	RDBMS	HS	Hive Server
HFS	HttpFS Service	GW	Gateway	HSS	Hive Metastore Service
JN	Journal Node	FA	Flume Agent		...
ZK	ZooKeeper				
HM	Hbase Master				
HRS	Hbase Region Server				

ACK! Seemingly no end to the Big Data services.

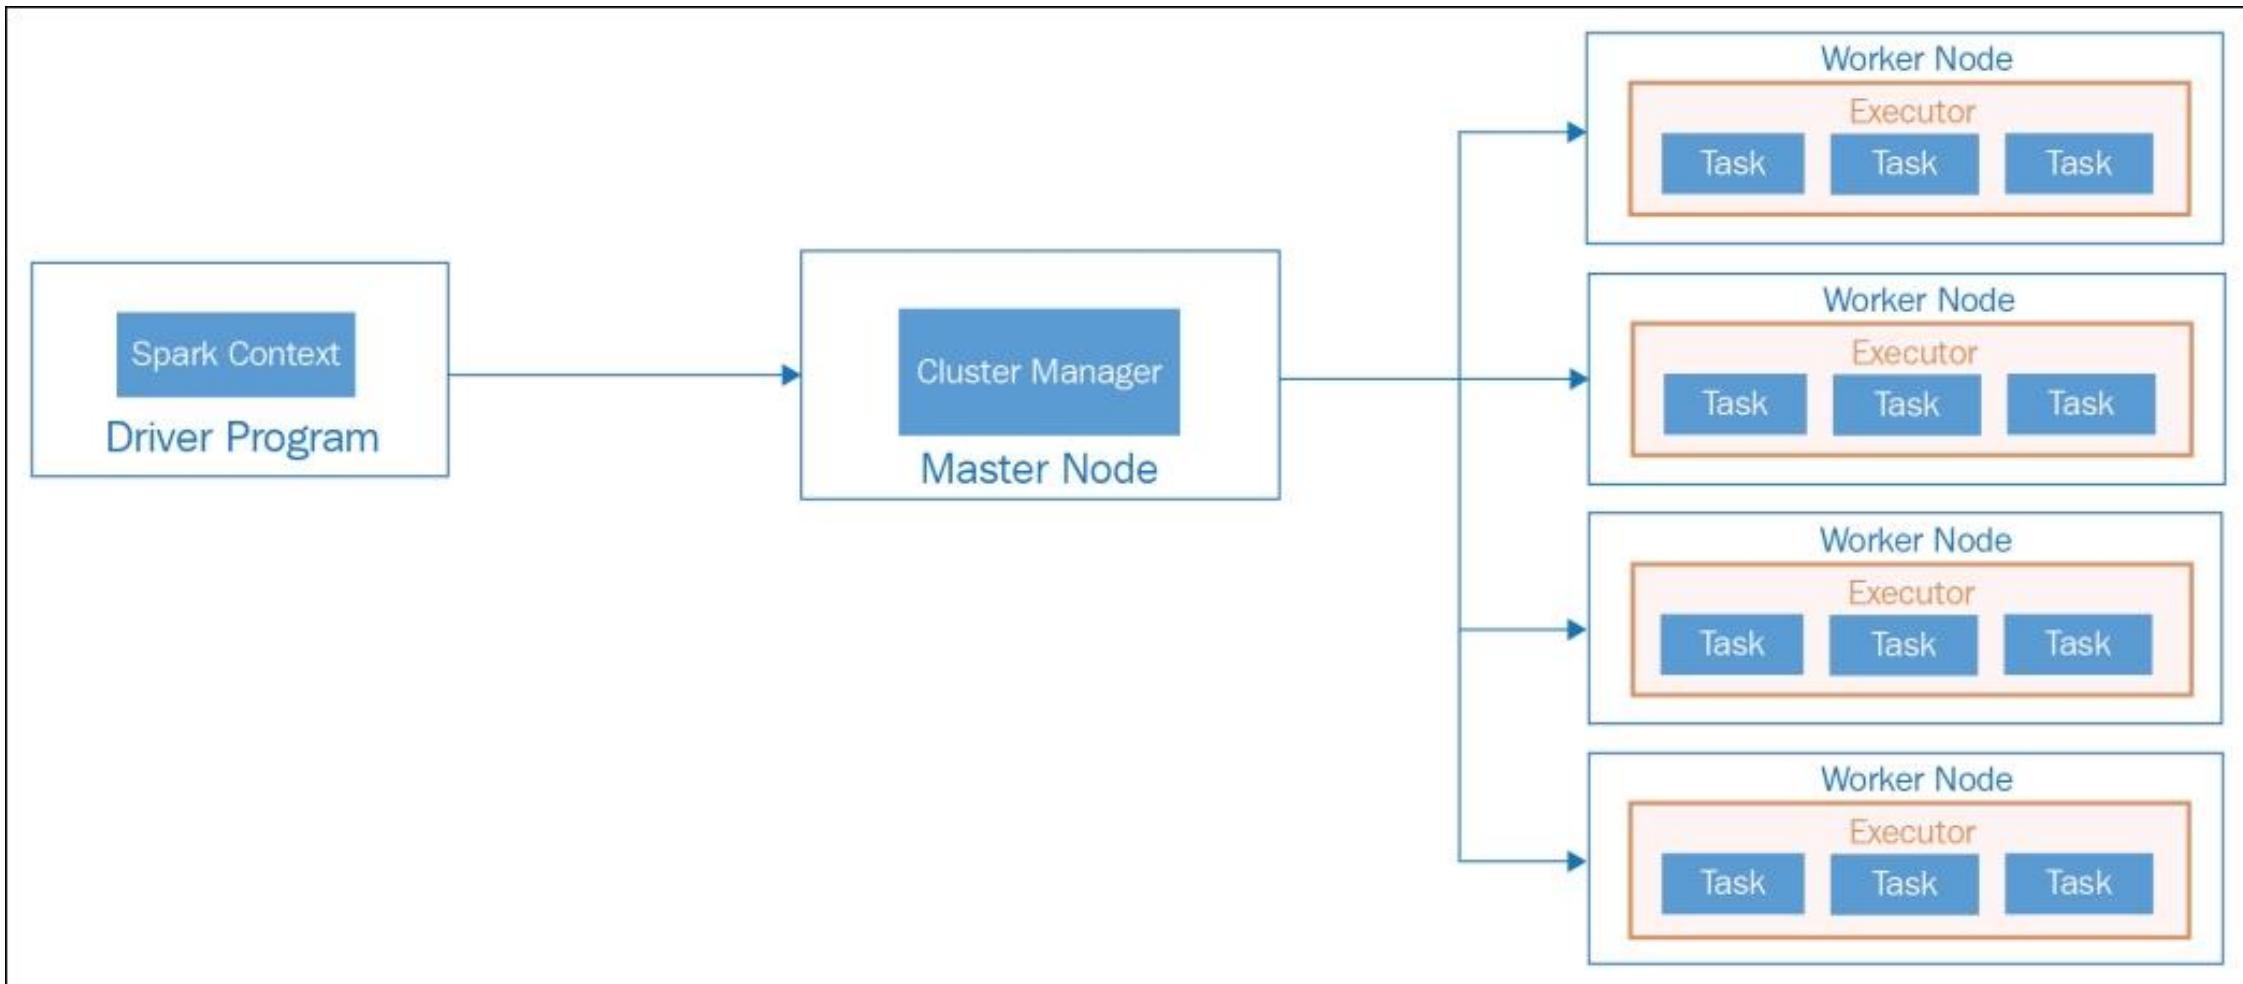
Hadoop 3.0 Architecture



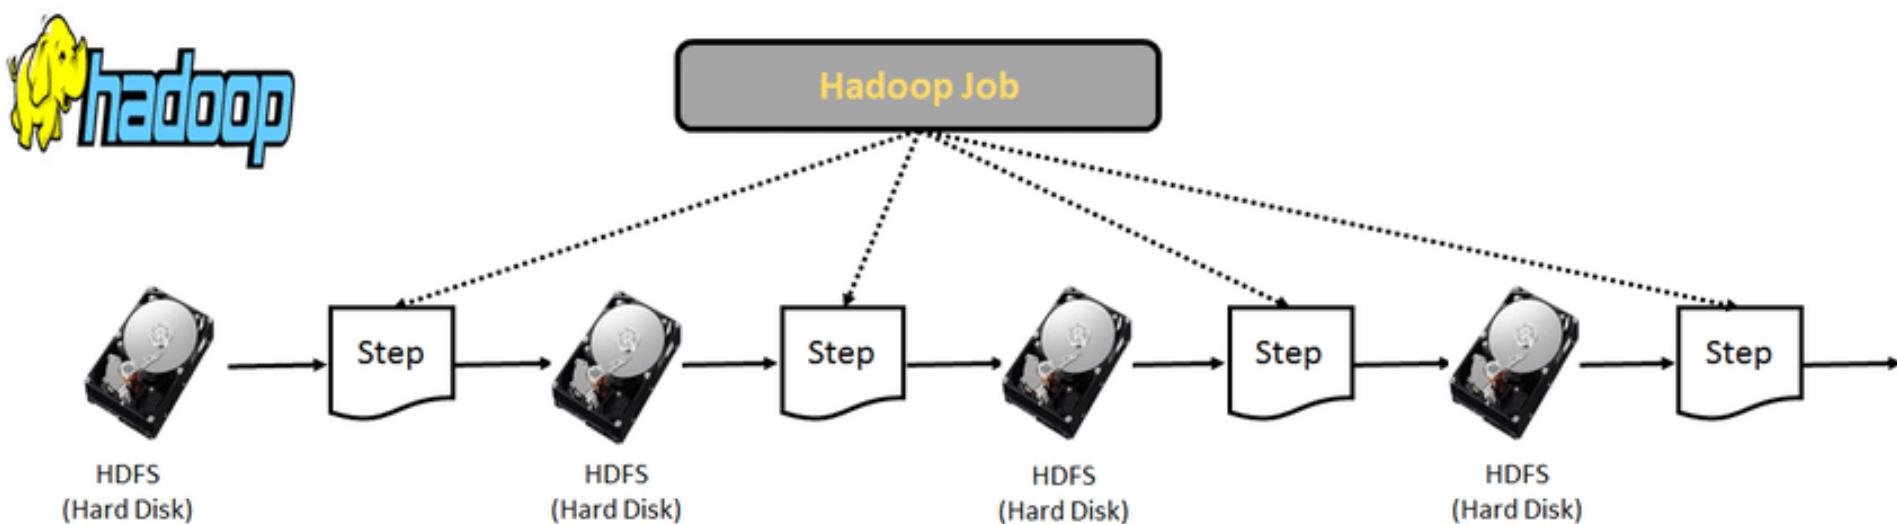
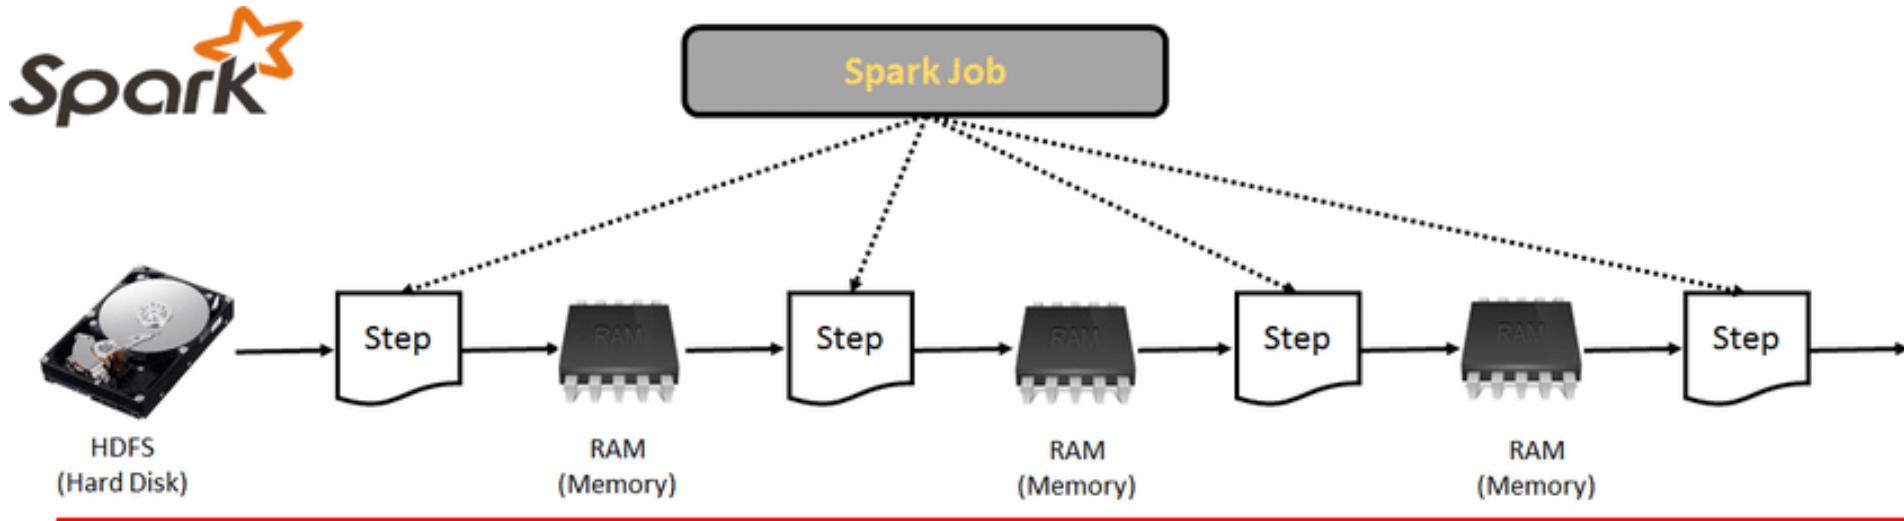
Kafka Architecture for Streaming



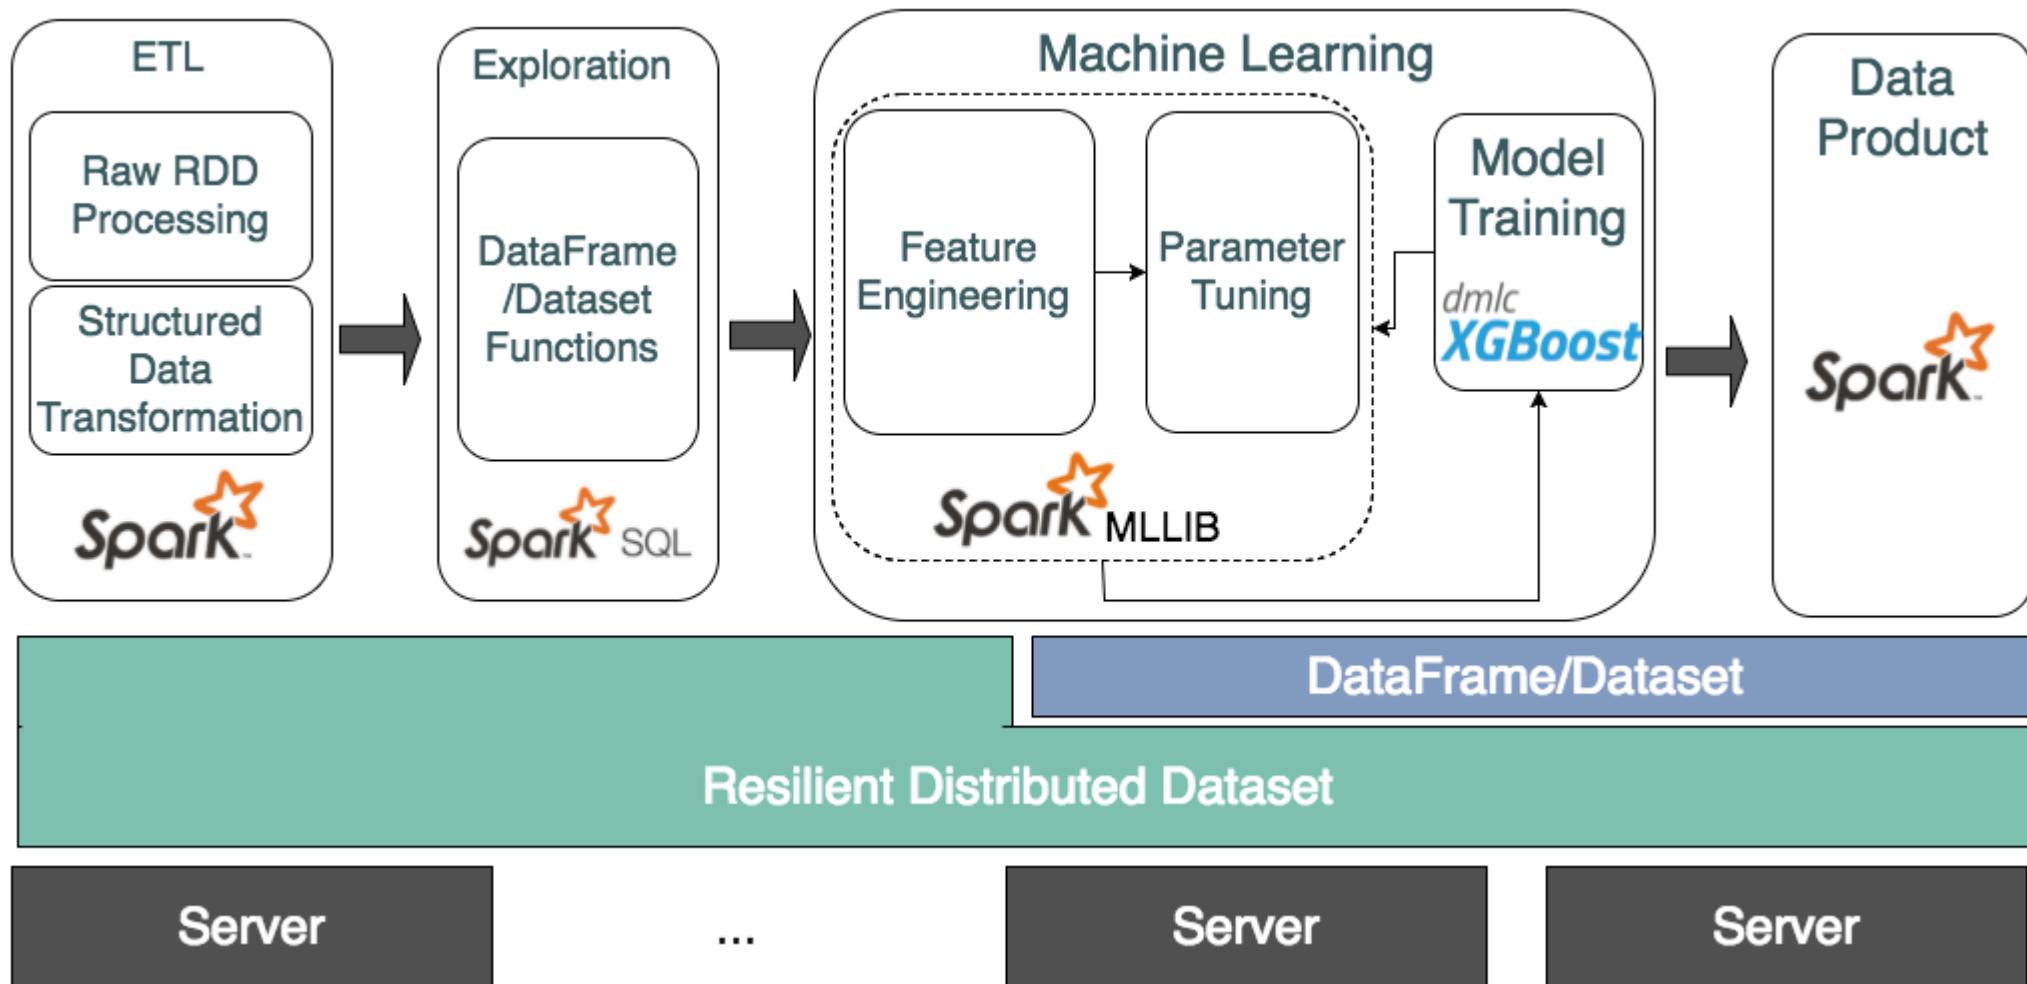
Spark Architecture



Spark vs Hadoop MapReduce



Data analytics with Spark



PySpark implementation (Keras)

```
def predict(iterator):
    model = InceptionV3(weights=None)
    model.load_weights(FLAGS.weights_file)
    return [(x[0], run_inference_on_image(model, x[1])) for x in iterator]

def main():
    sc = SparkContext(conf=conf)
    hbase_io = common.HbaseIO(FLAGS)
    out_format = common.OutputFormatter(FLAGS, MODEL_NAME)

    hbase_images = hbase_io.load_from_hbase(sc)
    classified_images = hbase_images.mapPartitions(predict) \
        .map(out_format.imagenet_format)
    classified_images.foreachPartition(hbase_io.put_to_hbase)
```

PySpark implementation (Keras)

- The Python environment with tensorflow is distributed to the executors at runtime, it is not preinstalled on the nodes.
- The individual models only need to implement the following functions:
 - prepare
 - predict
 - output_format
- Conceptually very close to scikit-learn or Spark ML Pipelines approach
- Deep Learning Pipelines can be a way to streamline the implementation

Spark implementation (dl4j / Scala)

```
def predict(pairs: Iterator[(String, (INDArray, Int, Int))]) = {
    val model = ModelSerializer.restoreComputationGraph(modelLoc)
    pairs.map{ case (name, image) =>
        (name, run_inference_on_image(model, image))
    }
}

def main(args: Array[String]) = {
    val sc = SparkContext(conf=conf)
    val hbase_io = common.HbaseIO(args)
    val out_format = common.OutputFormatter(args)

    val hbase_images = hbase_io.load_from_hbase(sc)
    val classified_images = hbase_images.mapPartitions(predict) \
        .map(out_format.imagenet_format)
    val classified_images.foreachPartition(hbase_io.put_to_hbase)
}
```

Image Clustering using KMeans

- Extraction of image features
- Conversion of specific data formats into:
`org.apache.spark.mllib.linalg.{Vector, Vectors}`
- Features are persisted in this reusable format as a Parquet file
- From here we go ... apply SparkML code to a part of the compound dataset.
- Finally we feed the new labels back into the compound dataset
→ e.g., for comparison of different clustering models with known clusters

./runSpark2.sh

```
import org.apache.spark.rdd.RDD
import org.apache.spark.mllib.linalg.{Vector, Vectors}
import org.apache.spark.mllib.regression.LabeledPoint
import org.apache.spark.mllib.clustering.{KMeans, KMeansModel}
import lire.Base64ImageConverter
import lire.LireToolWrapperGF
```

```
val sqlContext = new org.apache.spark.sql.SQLContext(sc)
import sqlContext.implicits._
```

// FEATURE EXTRACTION USING OPEN SOURCE LIBRARIES ...

// THIS IS ONE SLICE OF THE “COMPOUND DATASET”: Images can be vectorized in multiple ways, e.g., using color histograms

```
val dataPath = "/GITHUB/finance/data/out/image-analysis/A_2018-09-10-21-11-49/image_MD_ALL.parquet"
```

```
val df = sqlContext.read.parquet(dataPath)
val feature = df.select( "FCTH" )
val input = feature.rdd.map( x => x.getAs[org.apache.spark.mllib.linalg.Vector]("FCTH") );
```

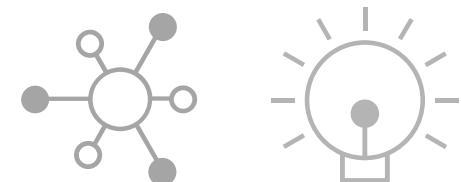
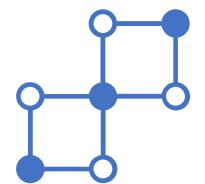
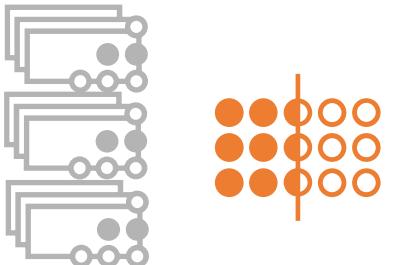
```
val kmeans = new KMeans().setK(8).setSeed(1L)
val model = kmeans.run( input )
```

```
val WSSSE = model.computeCost(vectors)
println(s"Within Set Sum of Squared Errors = $WSSSE")
```

// Shows the result: TODO: WRAP THIS MODEL INTO A “TENZING LABELER WHICH ALSO PERSISTED THE RESULTS IN MD-Collection”

```
println("Cluster Centers: ")
model.clusterCenters.foreach(println)
```

```
val dataToCluster = ....
val labeledDdata = model.predict( dataToCluster )
```



OUTLOOK: SPARK IMPLEMENTATION OF KMEANS CLUSTERING (SCALA)

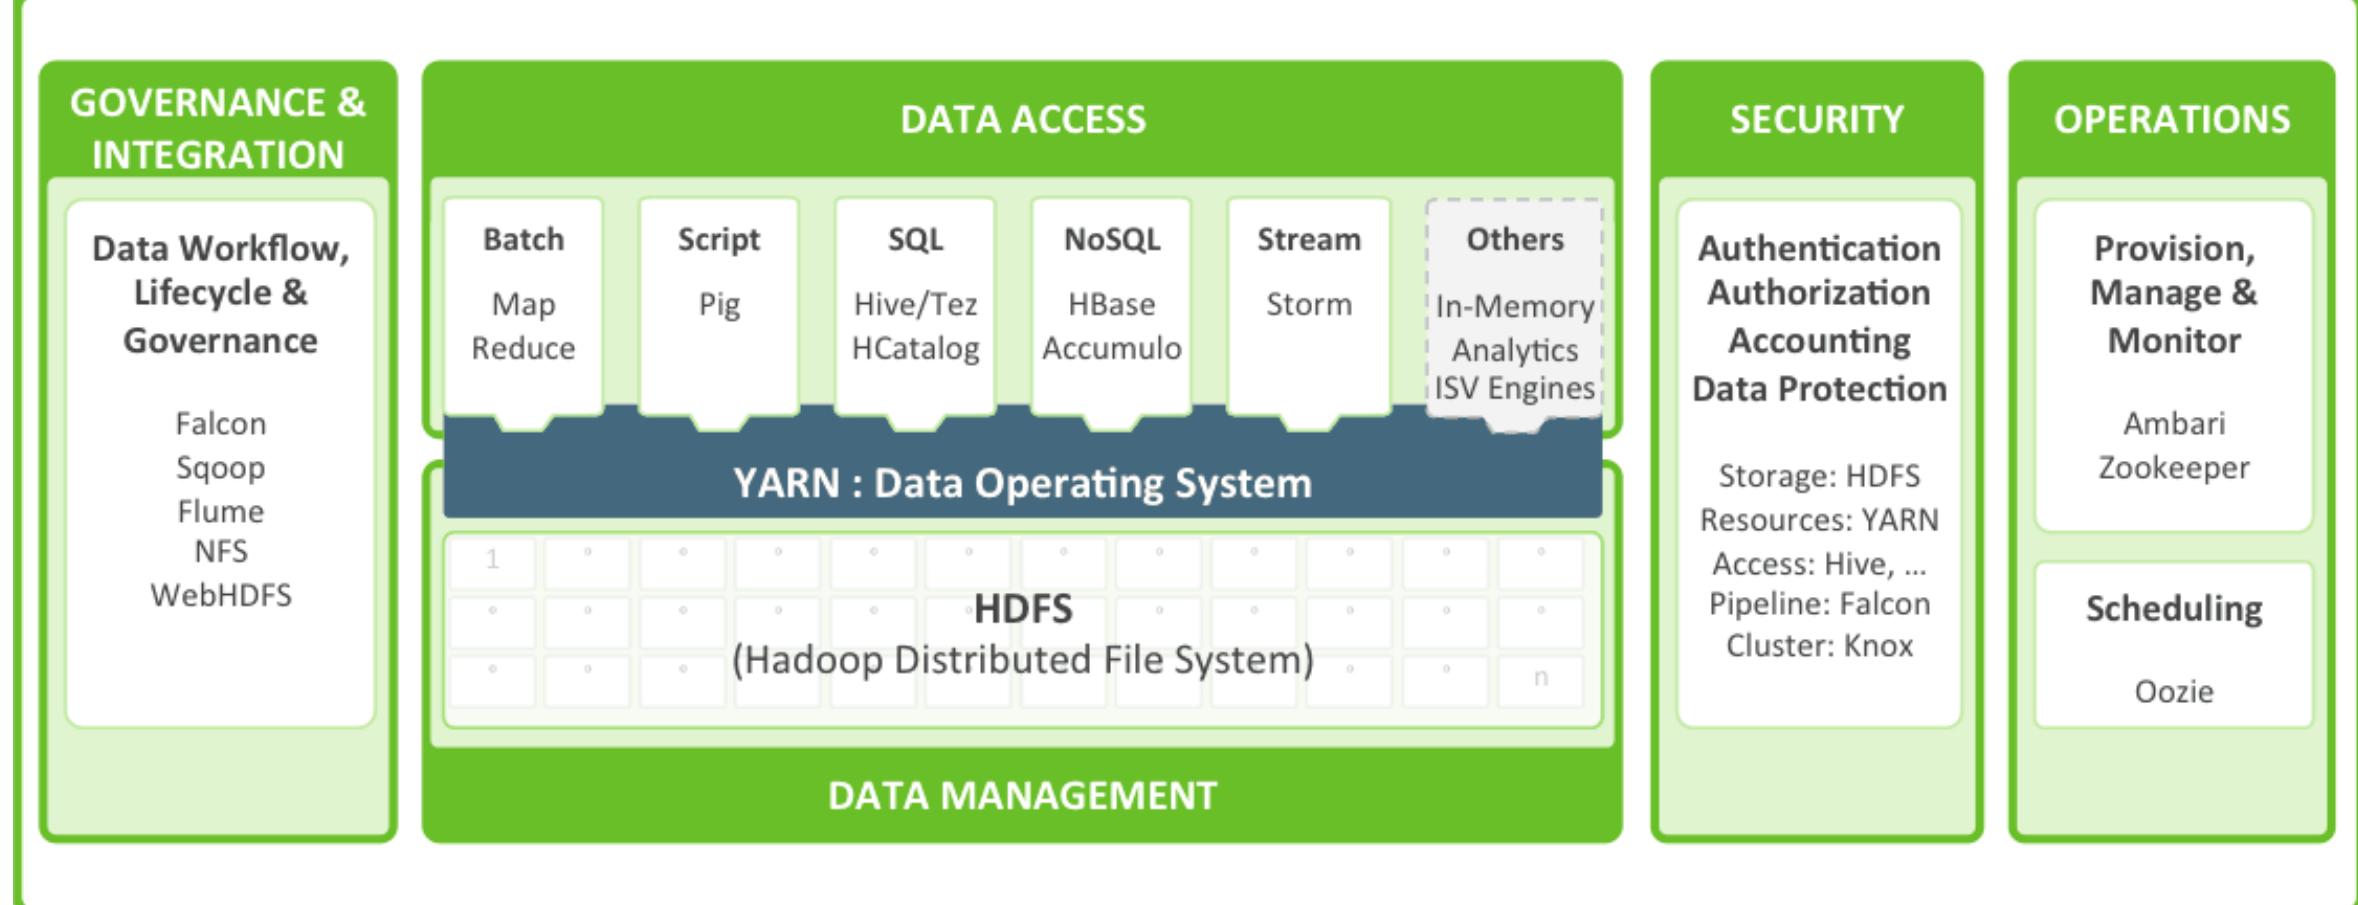
```
def predict(pairs: Iterator[(String, (INDArray, Int, Int))]) = {
    val model = ClusteringModelSerializer.restoreKMeansModel(modelLoc)
    pairs.map{ case (name, imageFeatures) =>
        (name, run_kmeans_clustering_on_image_features(model, imageFeatures))
    }
}

def main(args: Array[String]) = {
    val sc = SparkContext(conf=conf)
    val hbase_io = common.HbaseIO(args)
    val out_format = common.KMeansOutputFormatter(args)

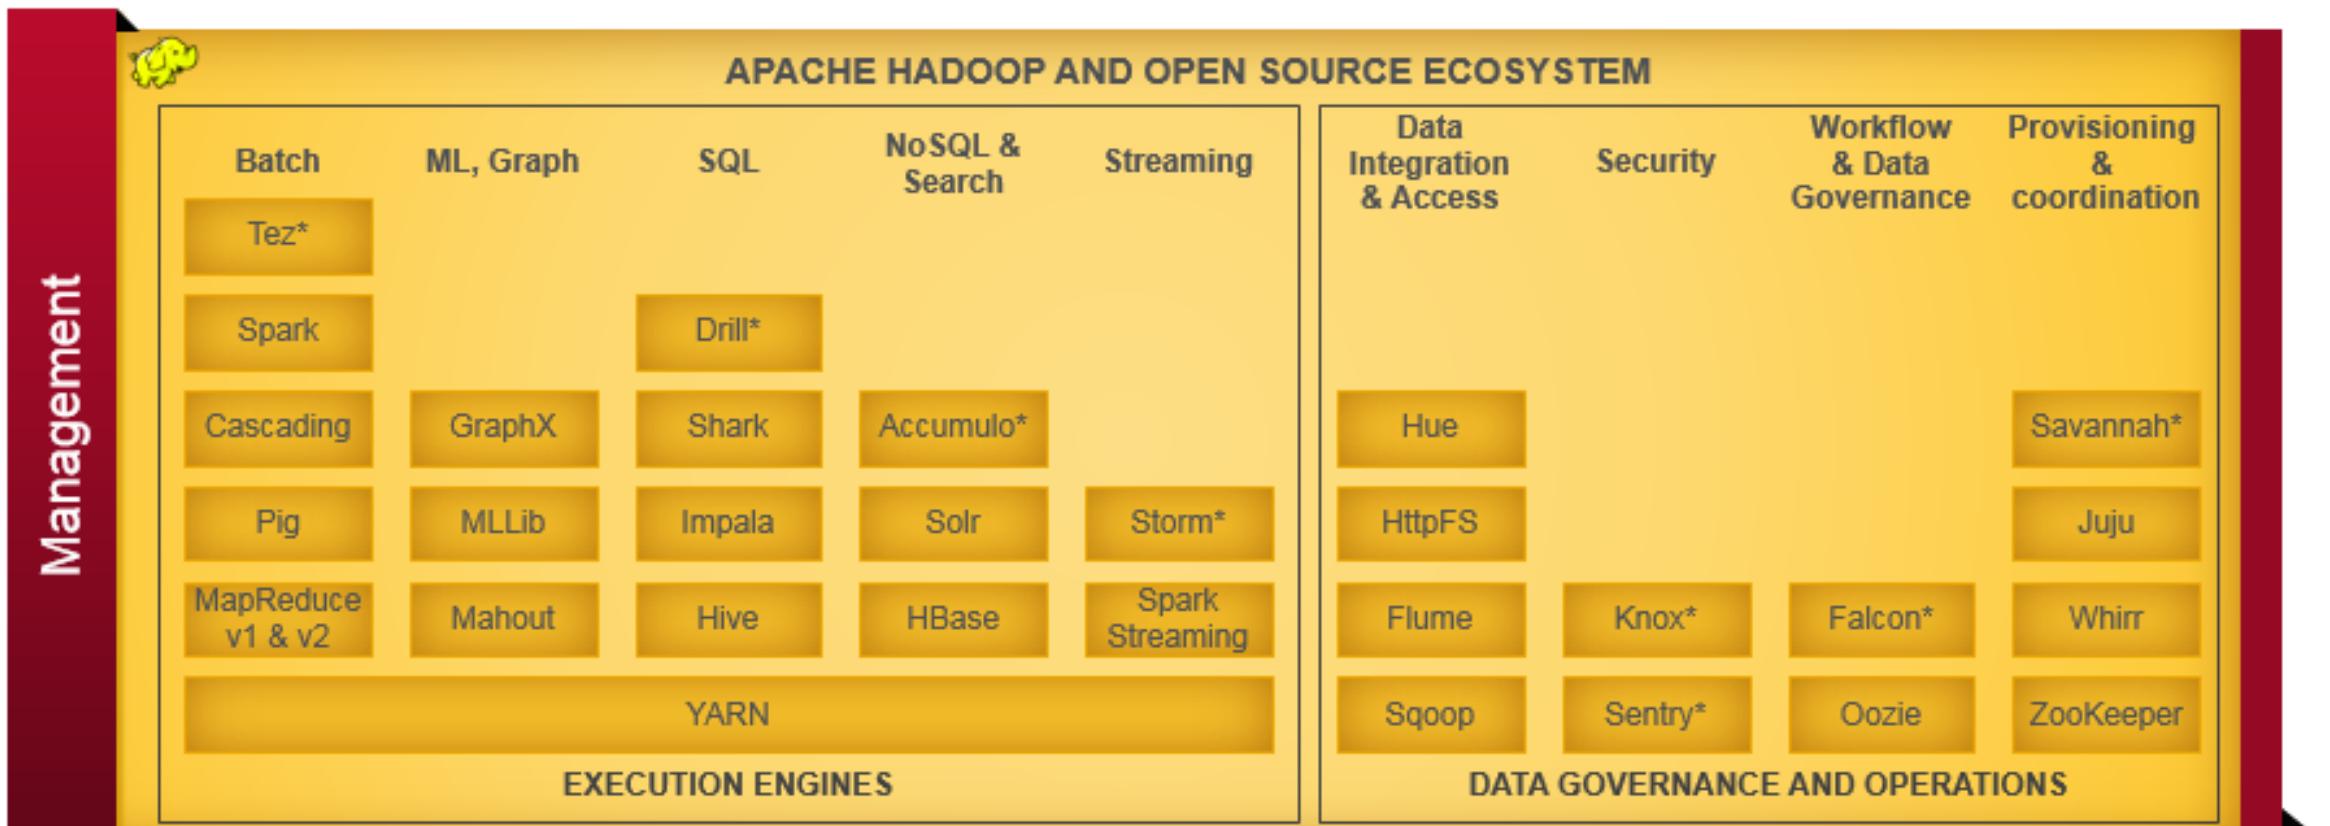
    val hbase_image_features = hbase_io.load_features_from_hbase(sc)
    val clustered_images = hbase_images.mapPartitions(predict) \
        .map(out_format.imagenet_format)
    val clustered_images.foreachPartition(hbase_io.put_to_hbase)
}
```

HortonWorks Architecture

Hortonworks Data Platform

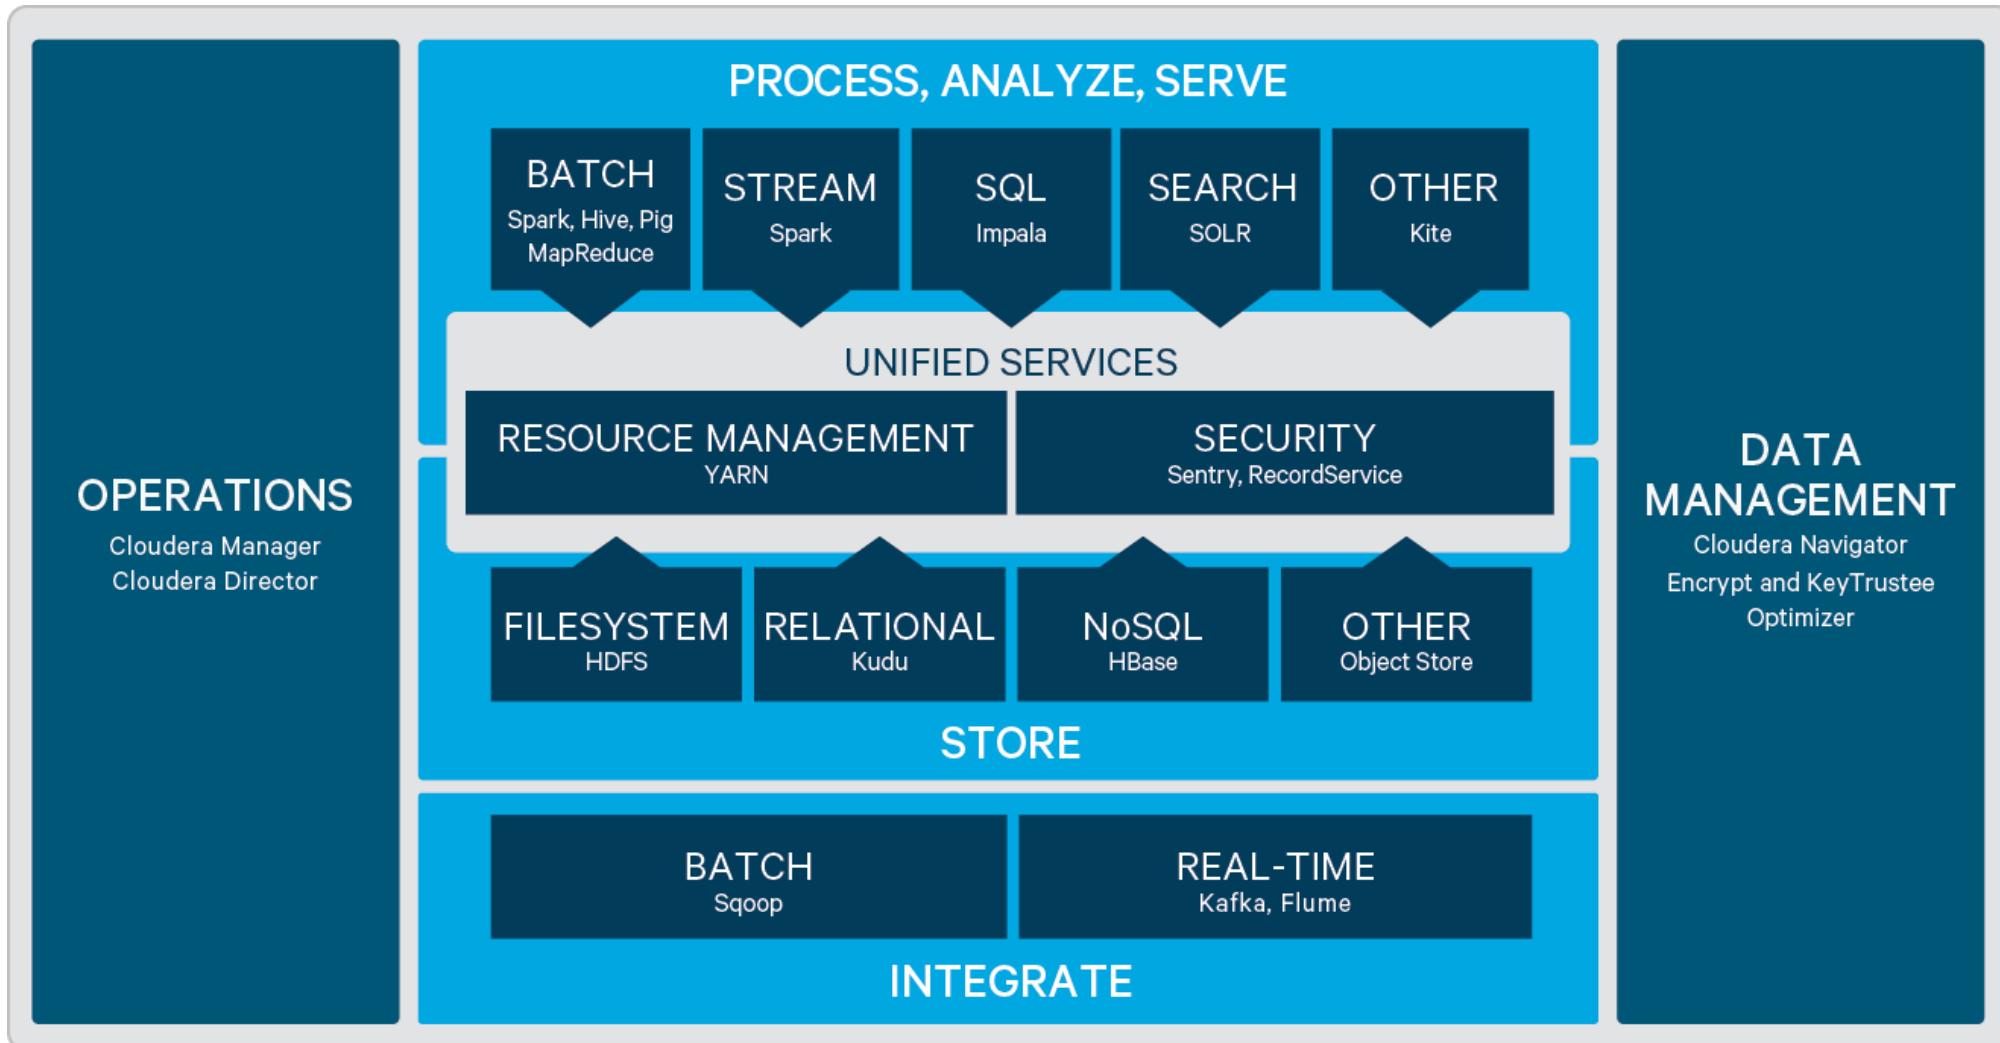


MapR Architecture

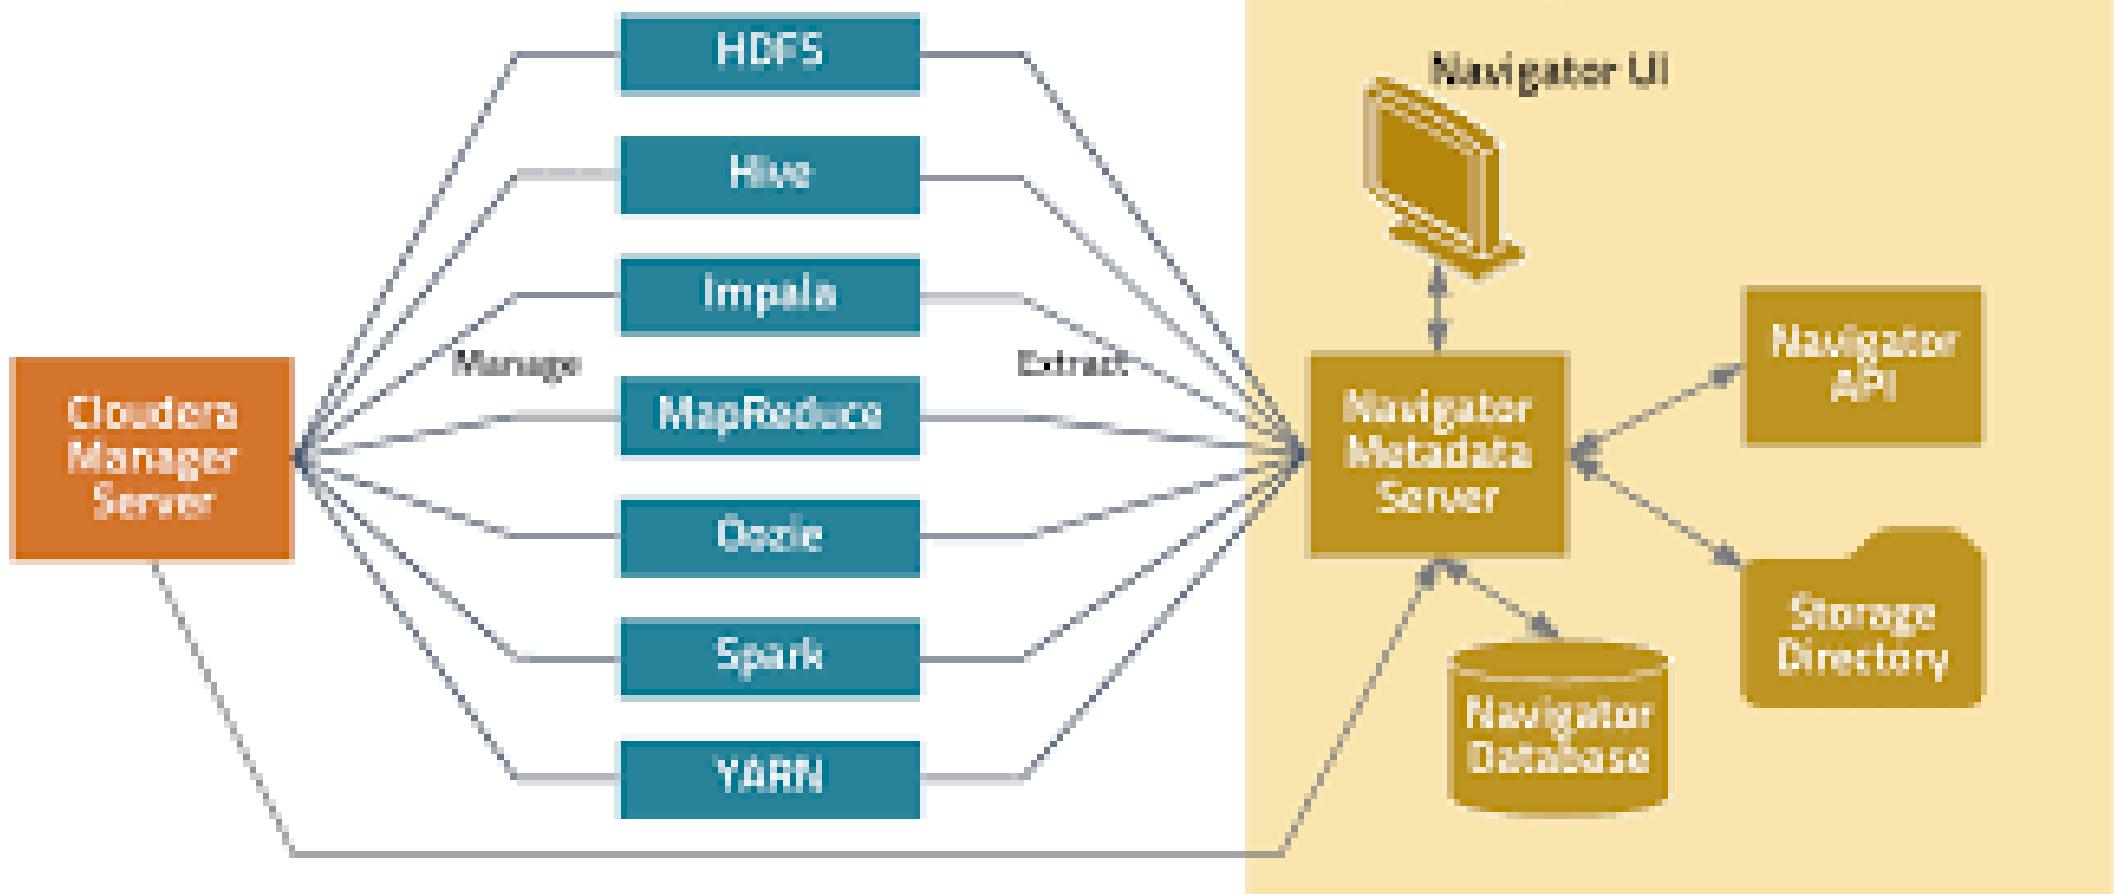


MapR Data Platform

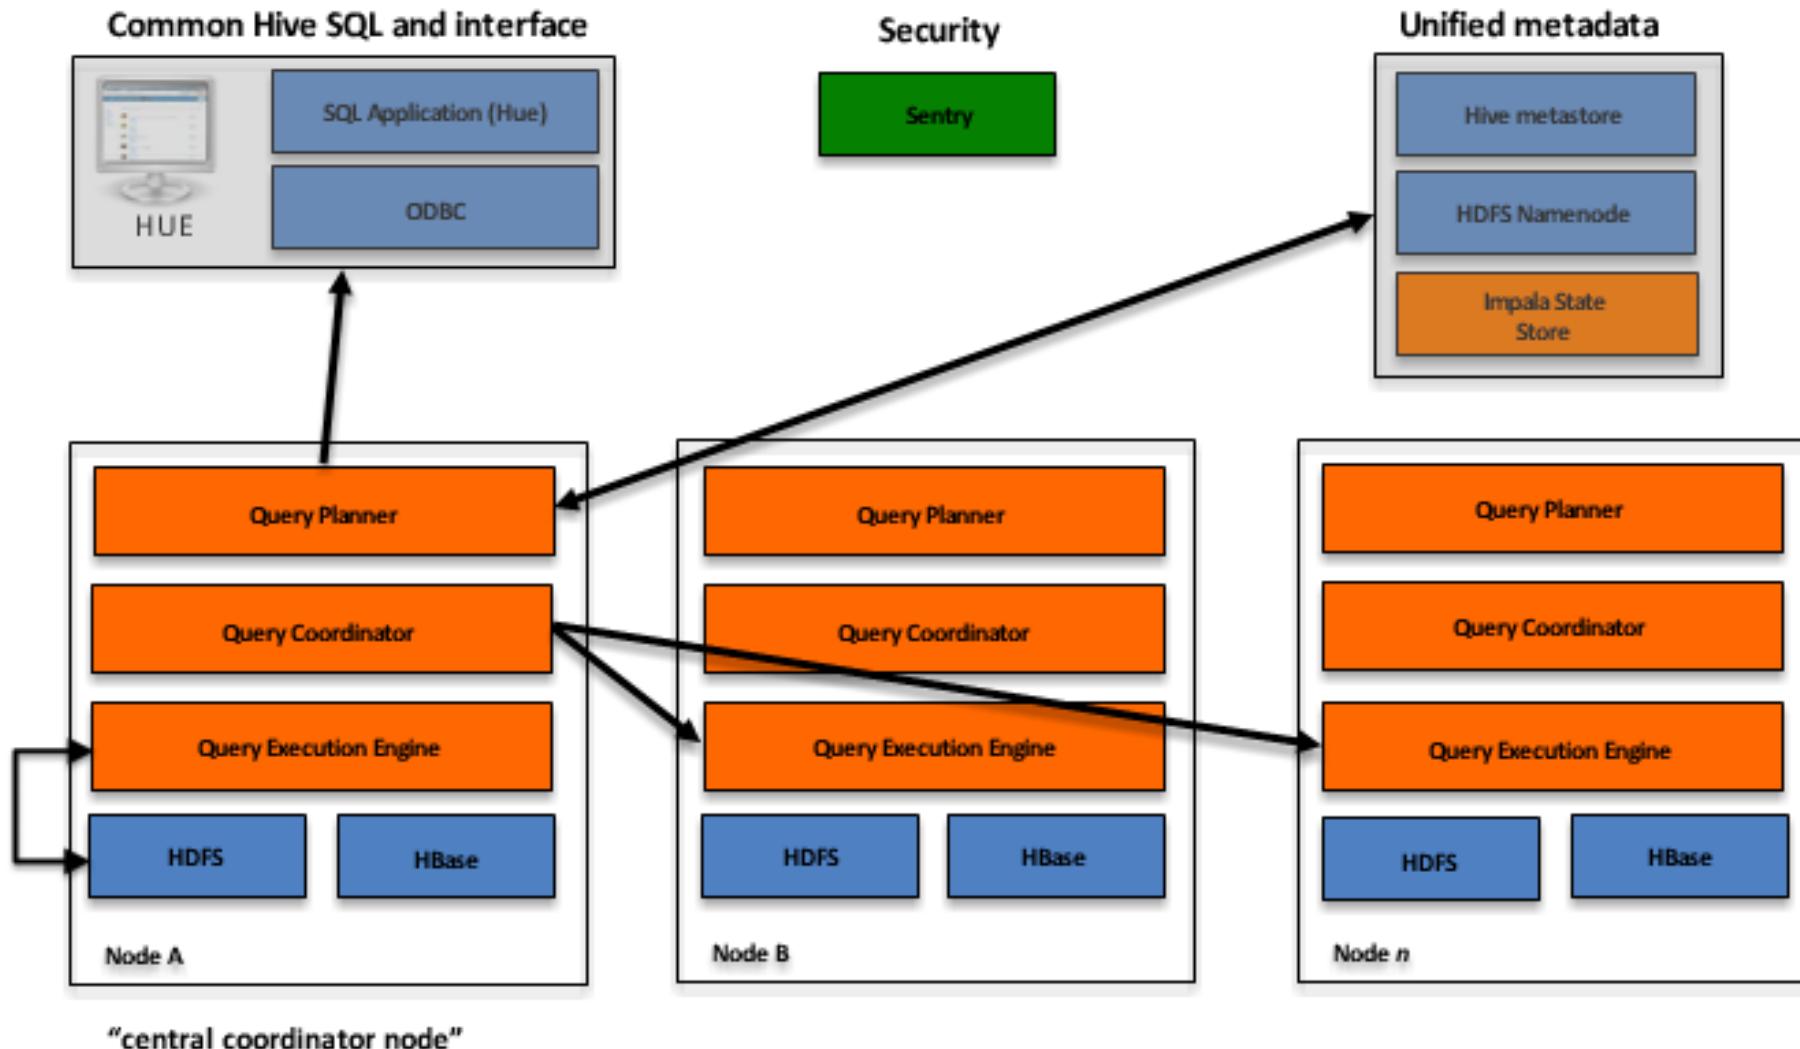
Cloudera CDH



Cloudera Architecture



Impala Architecture



BigData on AWS

<https://aws.amazon.com/es/big-data/getting-started/>



Images

Big Data Challenges:



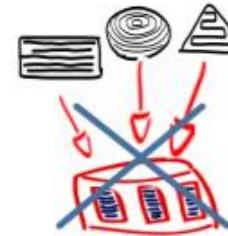
High Cost &
Commitment



Capacity
Planning &
scalability



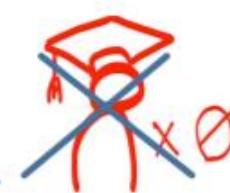
IT
Complexity



Data
Variety...



Volume,
velocity

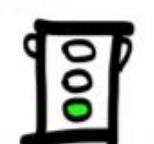


Advanced
profiles



Old Answers
& Questions

Big
Value



amazon

BIG DATA Platform:



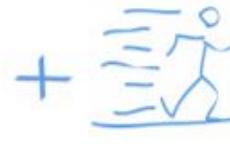
Lower Cost,
OpEx



Experiment &
learn more

Virtually
unlimited &
Elastic
Resources

No heavy lifting &
Reduced Time to
Market, parallel
processing on
demand



Managed
Services

Fully managed,
secured &
automated
services that
brings agility &
focus

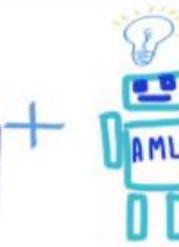


NOSQL



S3,
EMR,
Kinesis,
Redshift,
DynamoDB:

Collect all data, do
Complex
computations and
processing it, both
in Real-Time &
Batch



Machine
Learning

Easy deployment
of ML powerful
models without
the need of ML
Experts ready to
be used



New Answers/
questions &
Business Ideas

Extract the
meaning from all
your data & focus
on new business
Ideas, Models, etc..

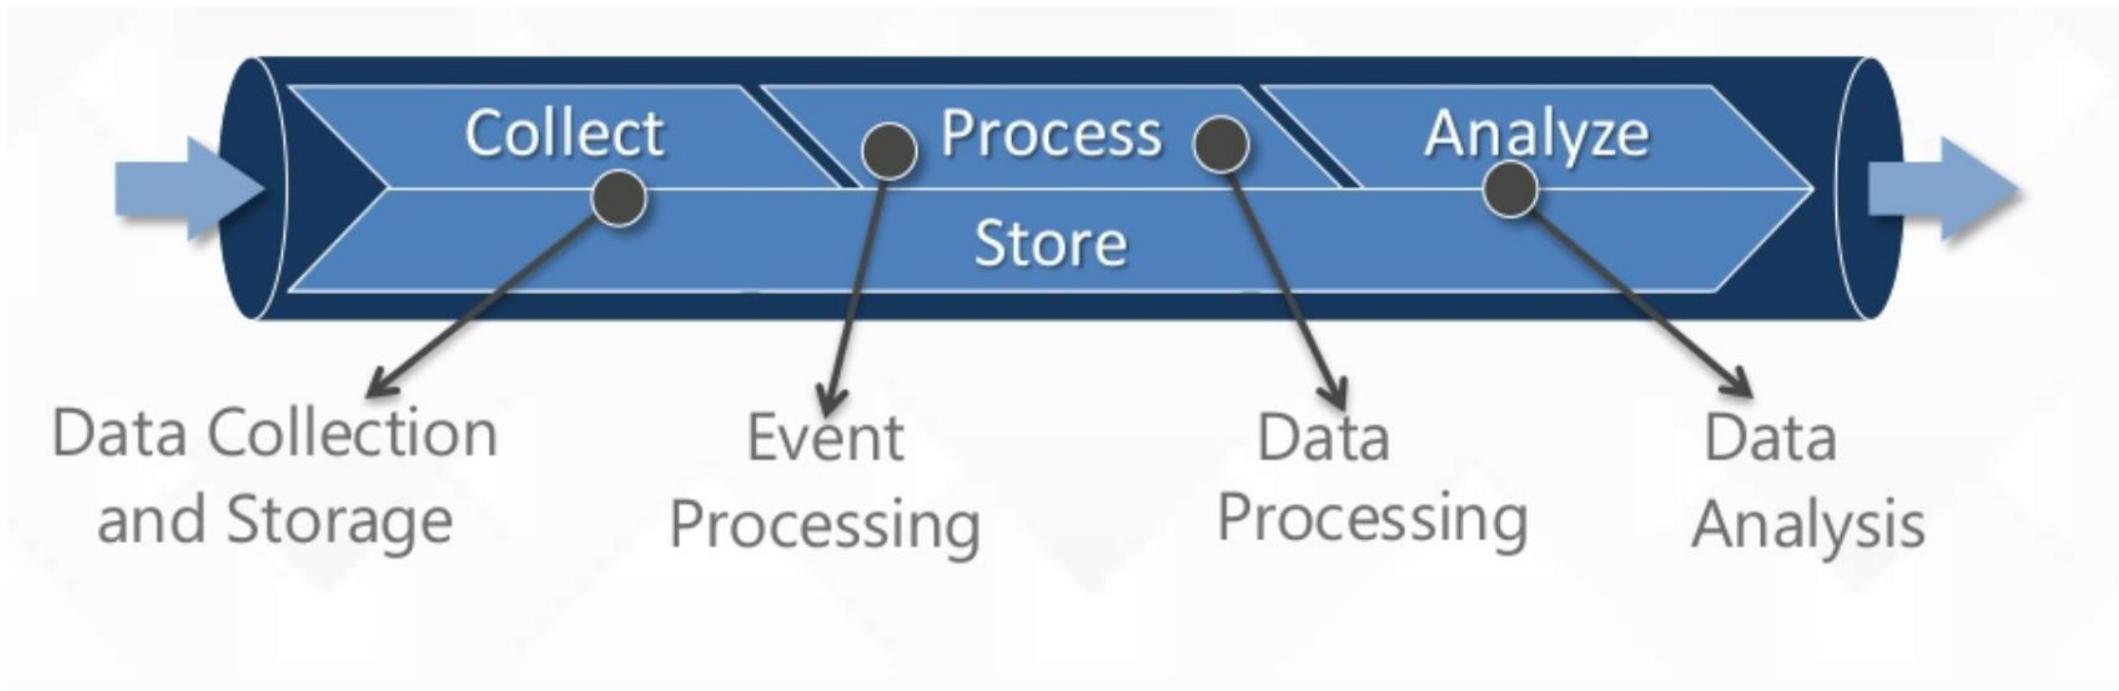
BigData Portfolio on AWS



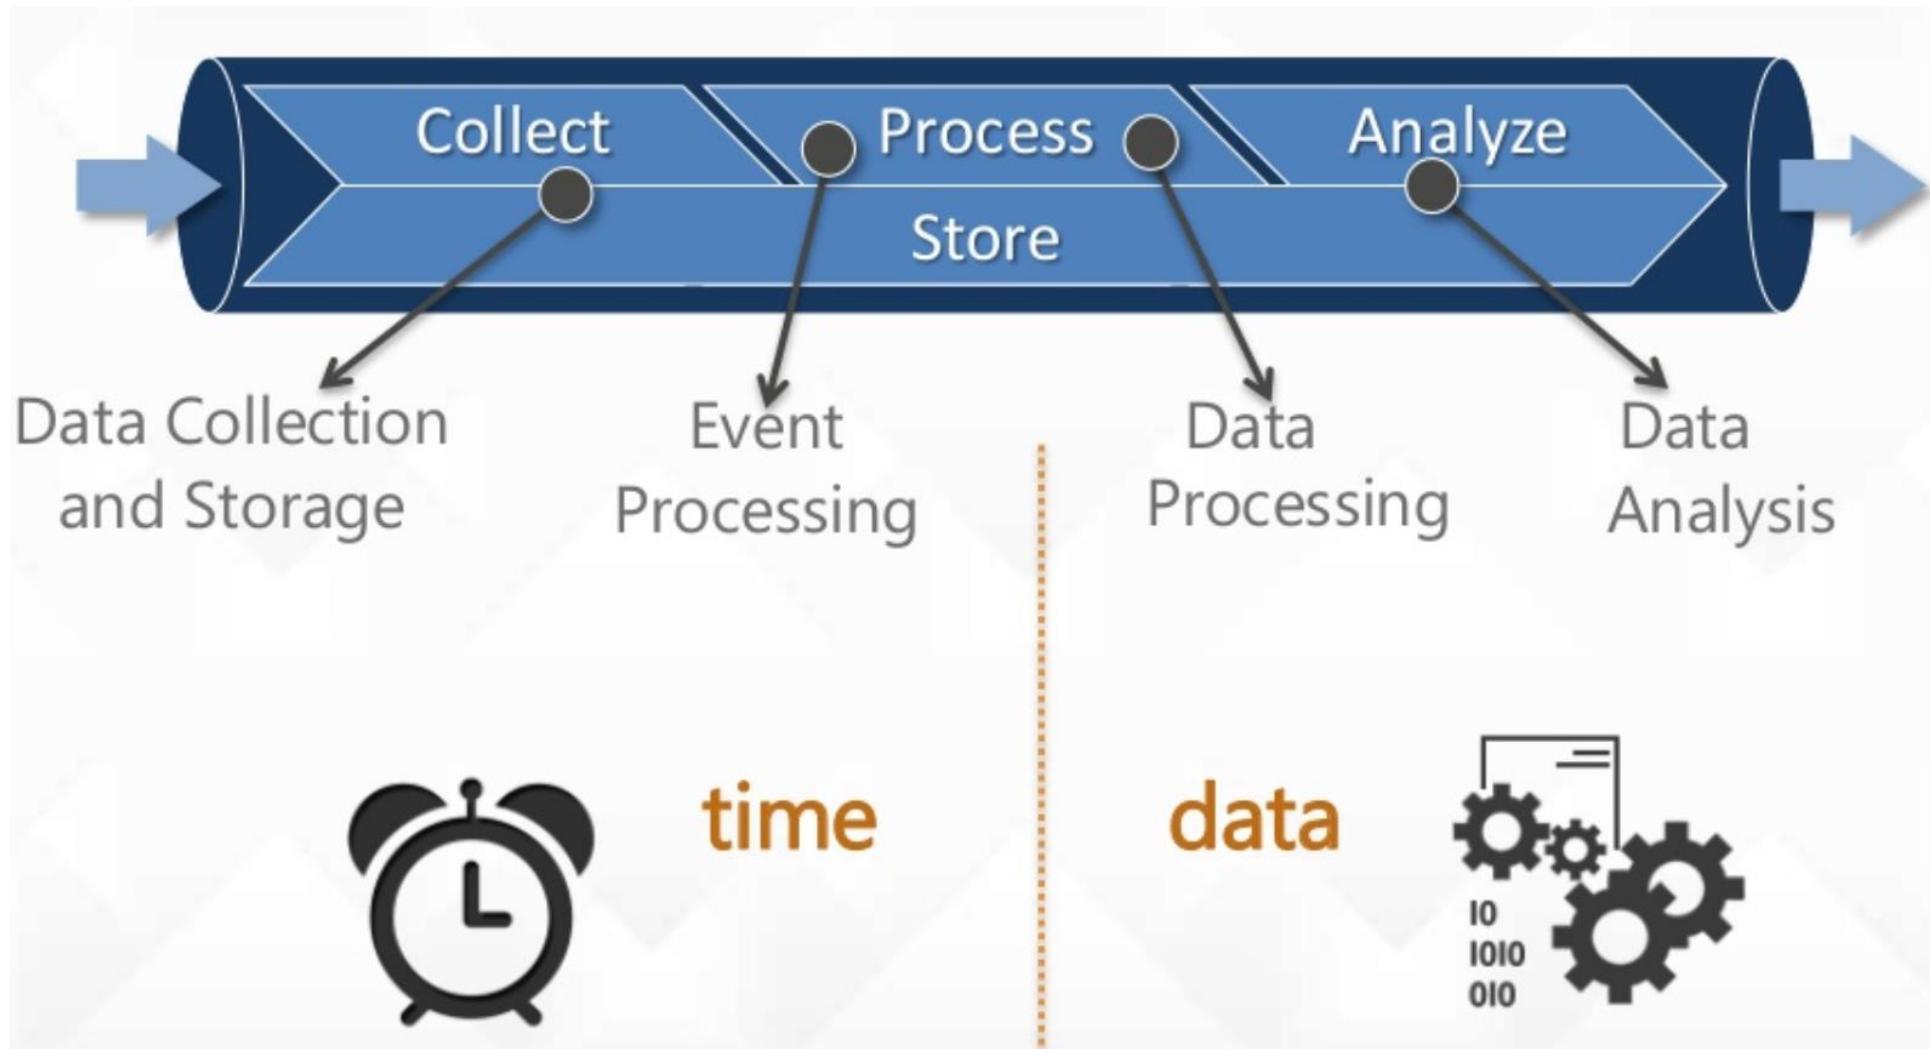
Big Data Pipeline



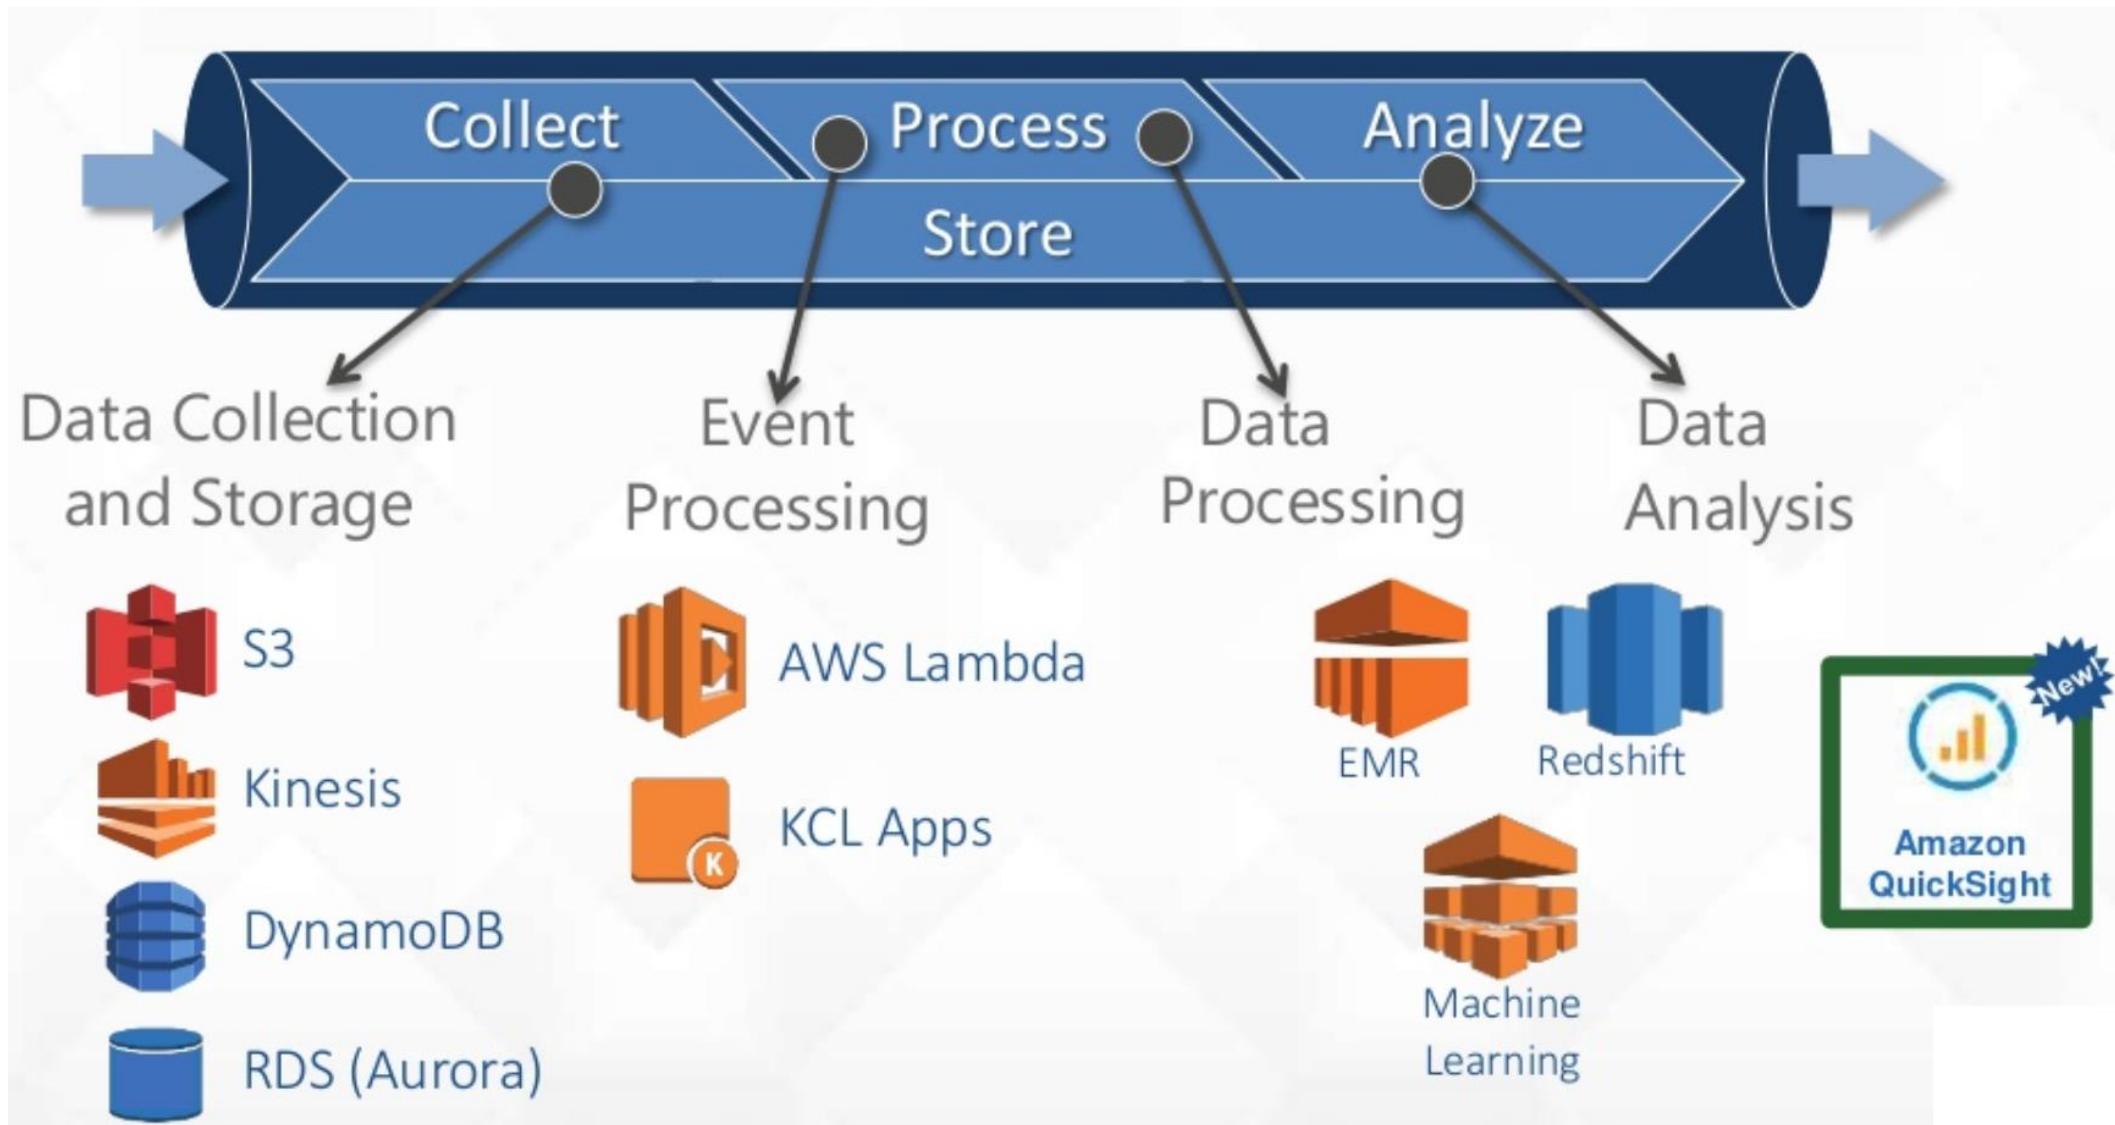
Primitive patterns



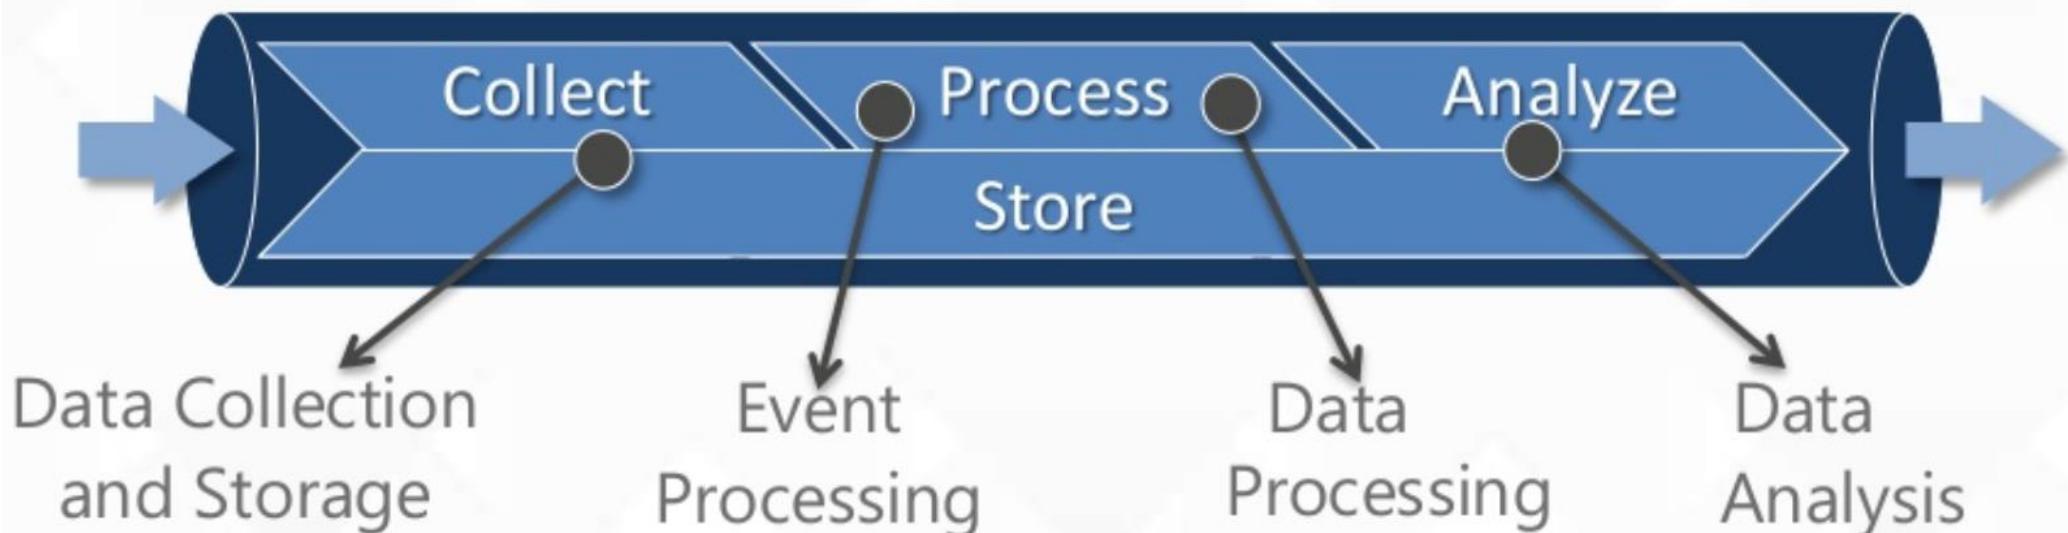
Primitive patterns



Primitive patterns



Primitive patterns



S3



Kinesis



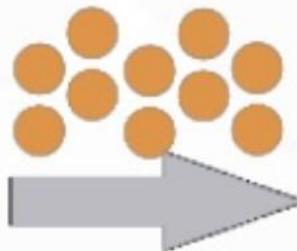
DynamoDB



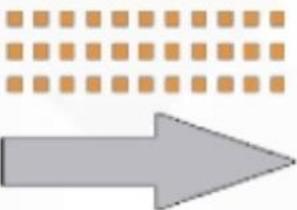
RDS (Aurora)

Data Collection and Storage

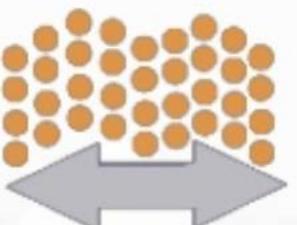
Apps Devices Logging Frameworks



File: media, log files (sets of records)



Stream: records (eg: device stats)

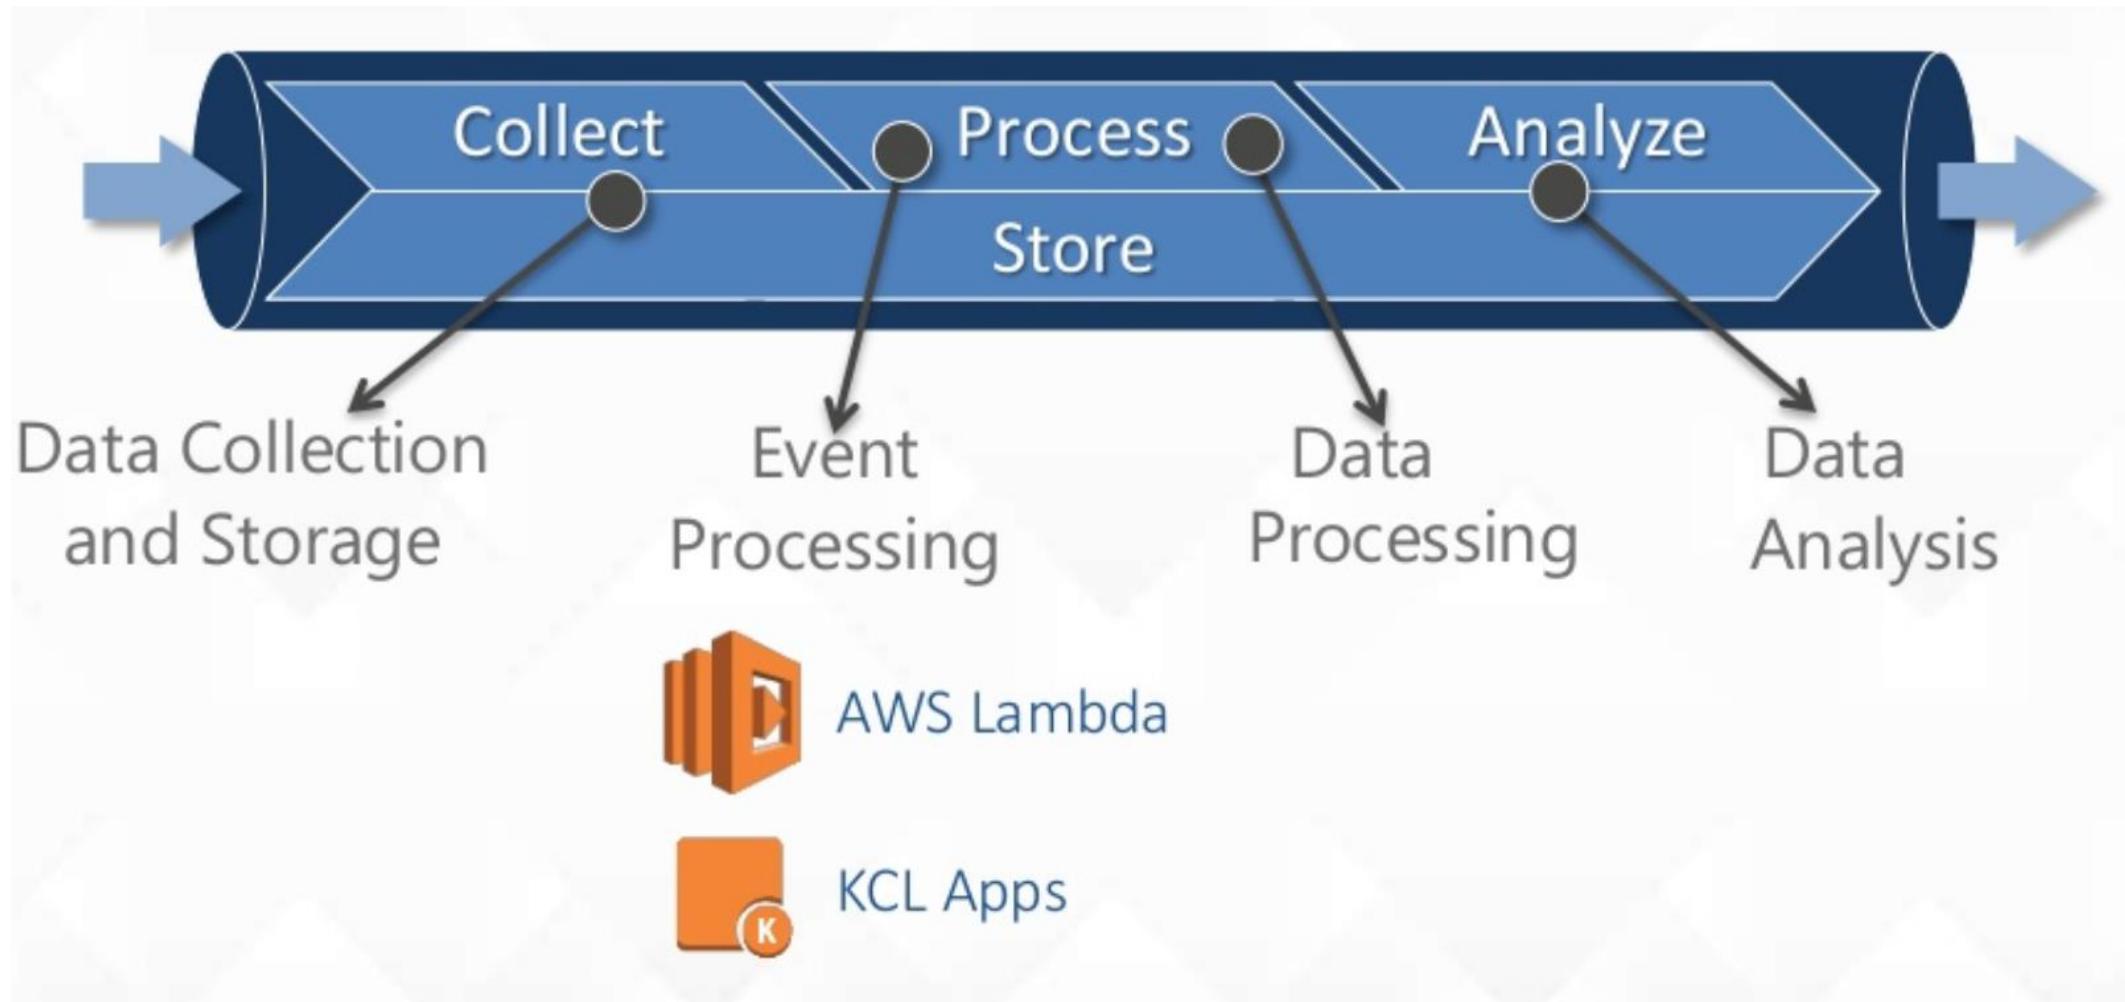


Transactional: database reads/writes

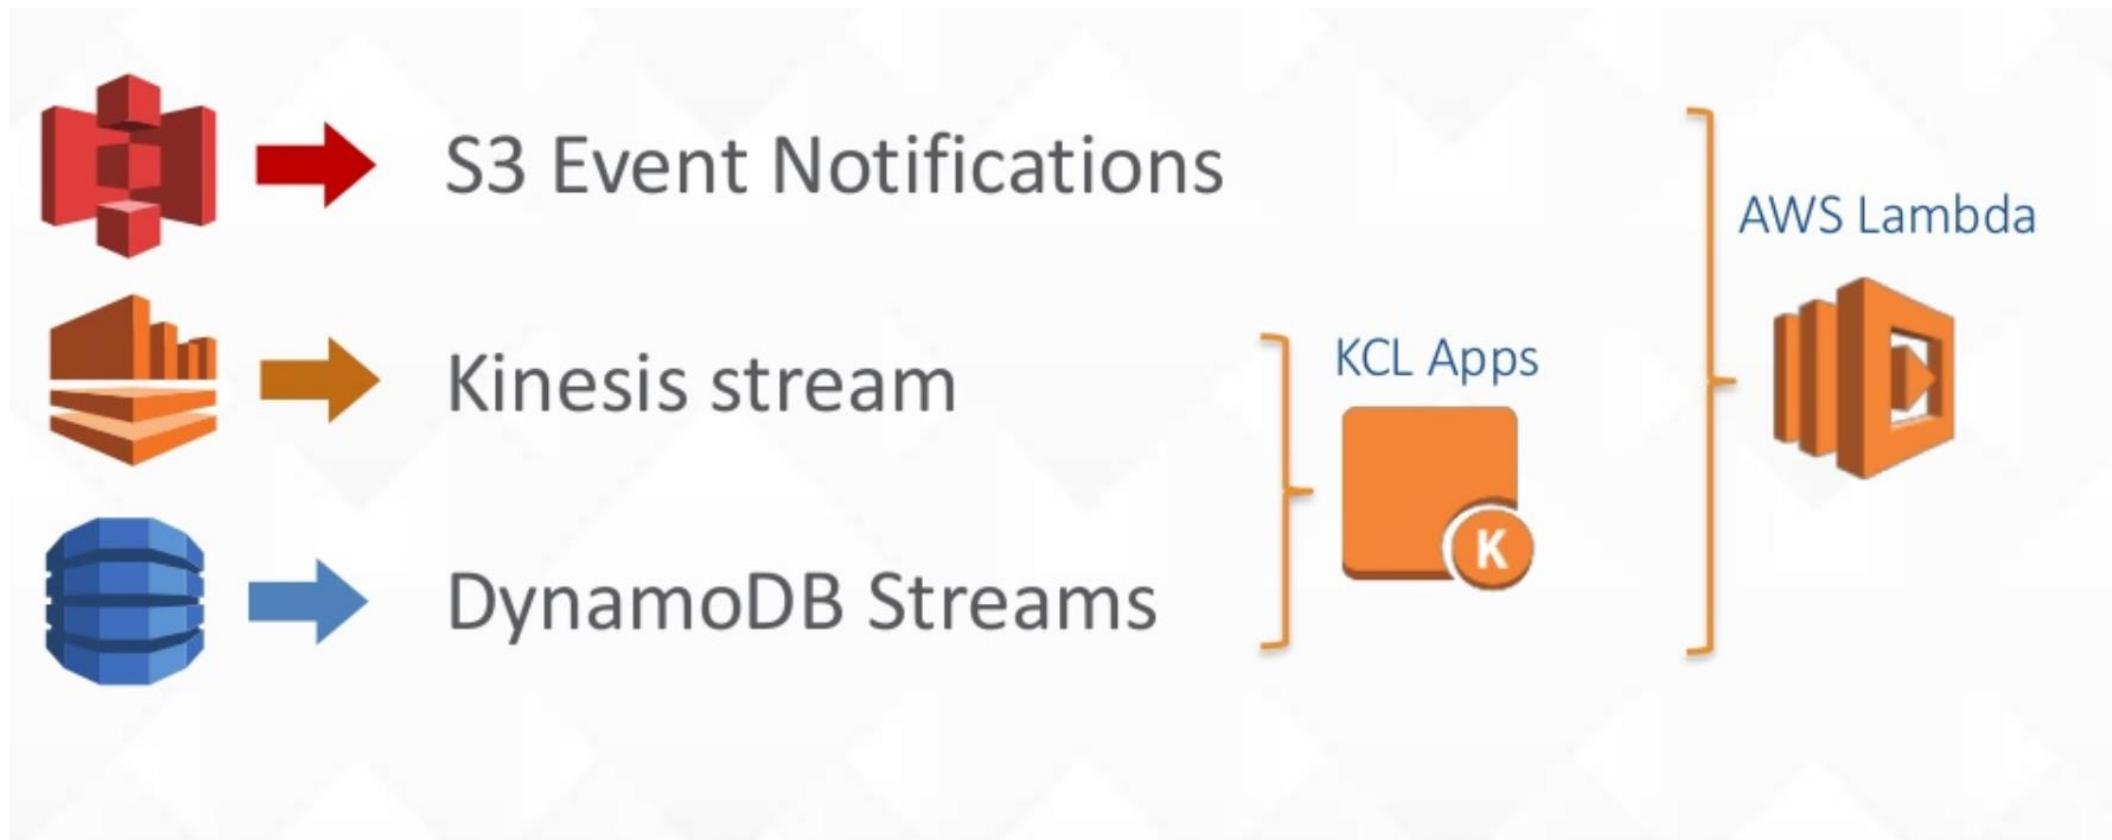
AWS Service: Data Collection and Storage



Primitive patterns



Event processing: enabling capabilities

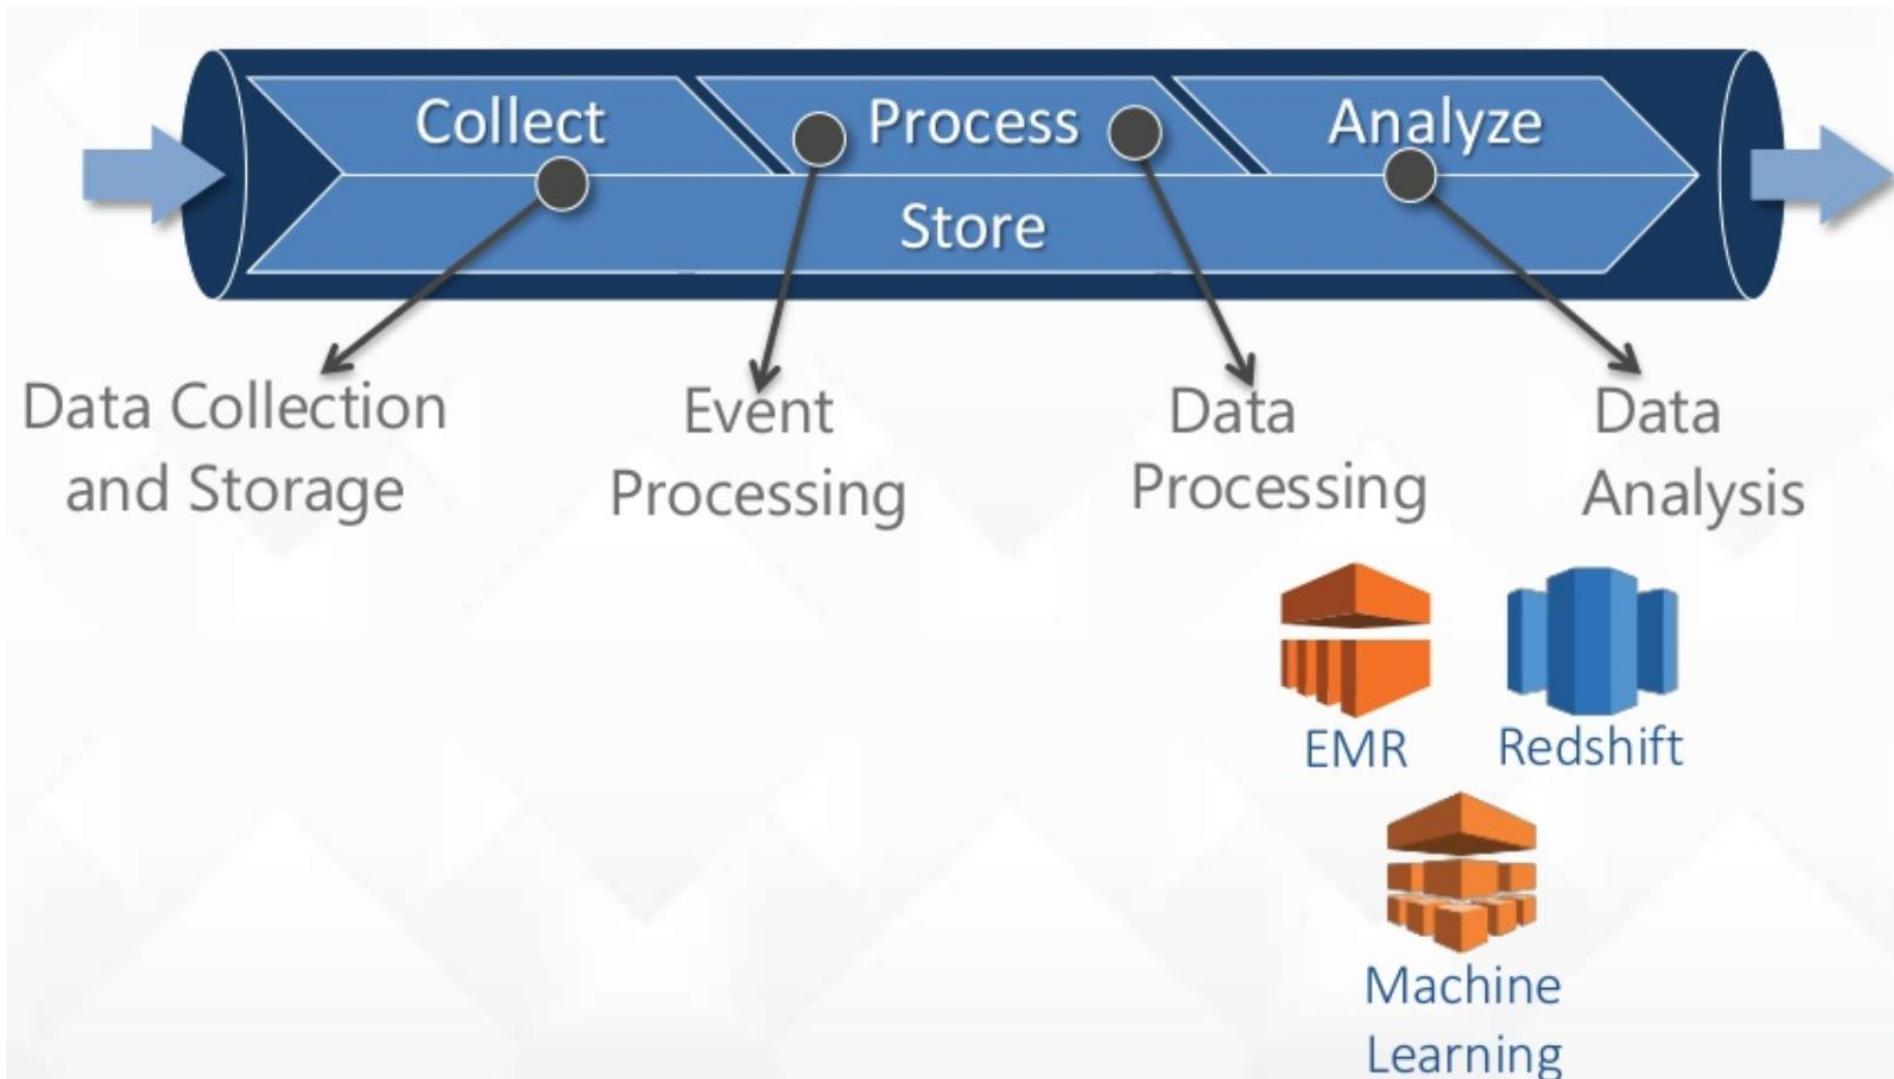


Benefits of Event Processing

- Build / add real-time events
 - Take action between data collection and analytics
 - Alerts and notifications, performance and security
 - Automated data enrichment (eg: aggregations)
- De-couple application modules
 - Streamline development and maintenance
 - Increase agility
 - MVP + iterate on discrete components



Primitive patterns



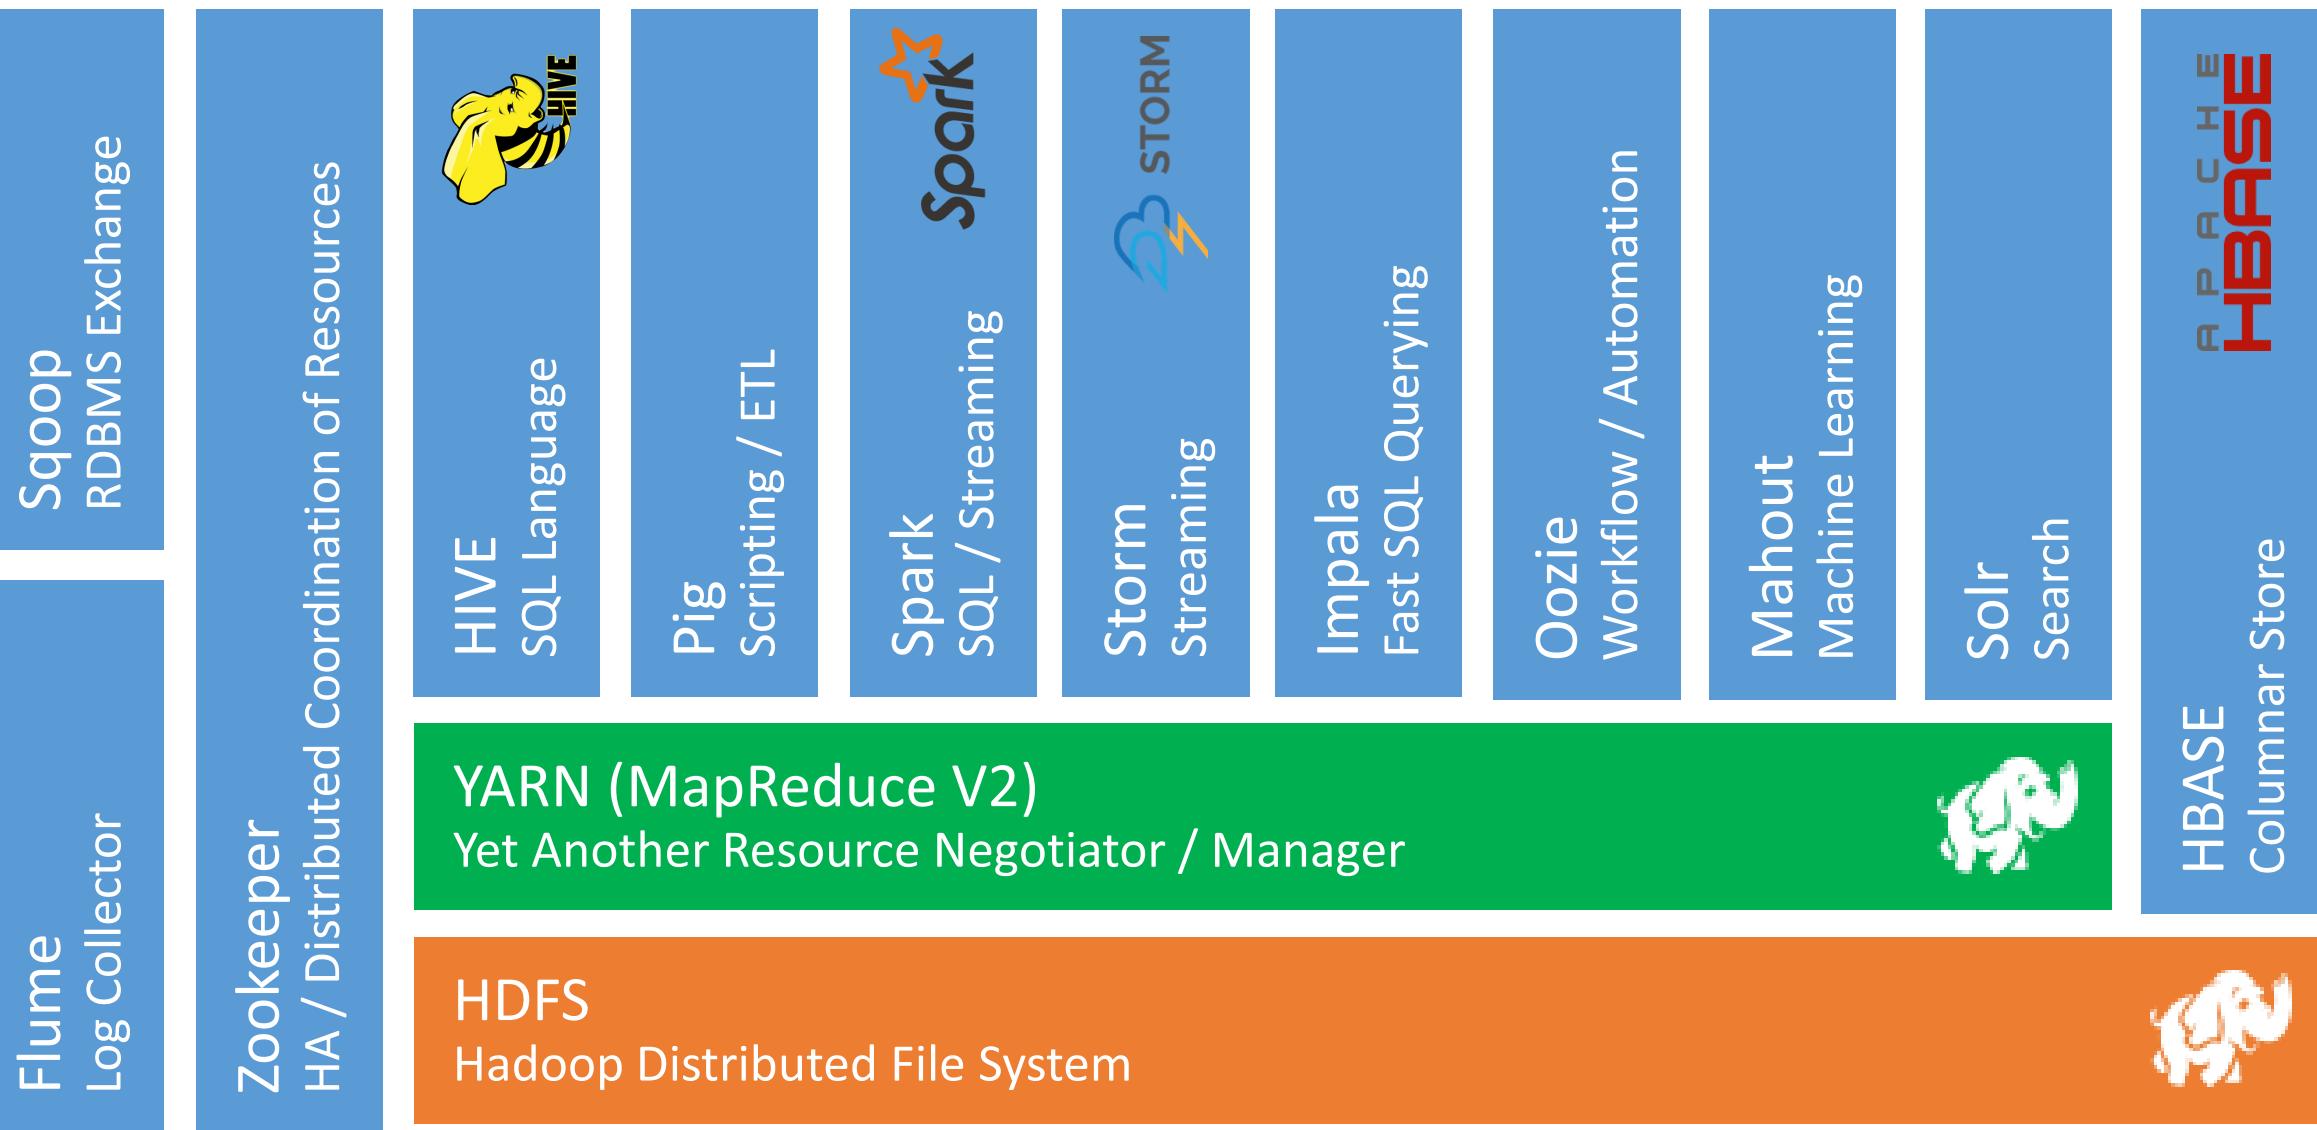
BigData on Azure

<https://azure.microsoft.com/en-us/solutions/big-data/>

Ambari / Cloudera Manager / Cloudera Director

Cluster Management and Monitoring UI

Hadoop Ecosystem



Hadoop On Your Terms

Microsoft Azure
HDInsight 

[100% Apache Hadoop-based Service in
the Cloud](#)



[Hortonworks Data Platform is now
Microsoft Azure Certified](#)

cloudera

[Cloudera Selects Microsoft Azure as a
Preferred Cloud Platform](#)

Allied Market Research estimates the global Hadoop market to grow from \$2 billion in 2013 to \$50.2 billion by the year 2020.

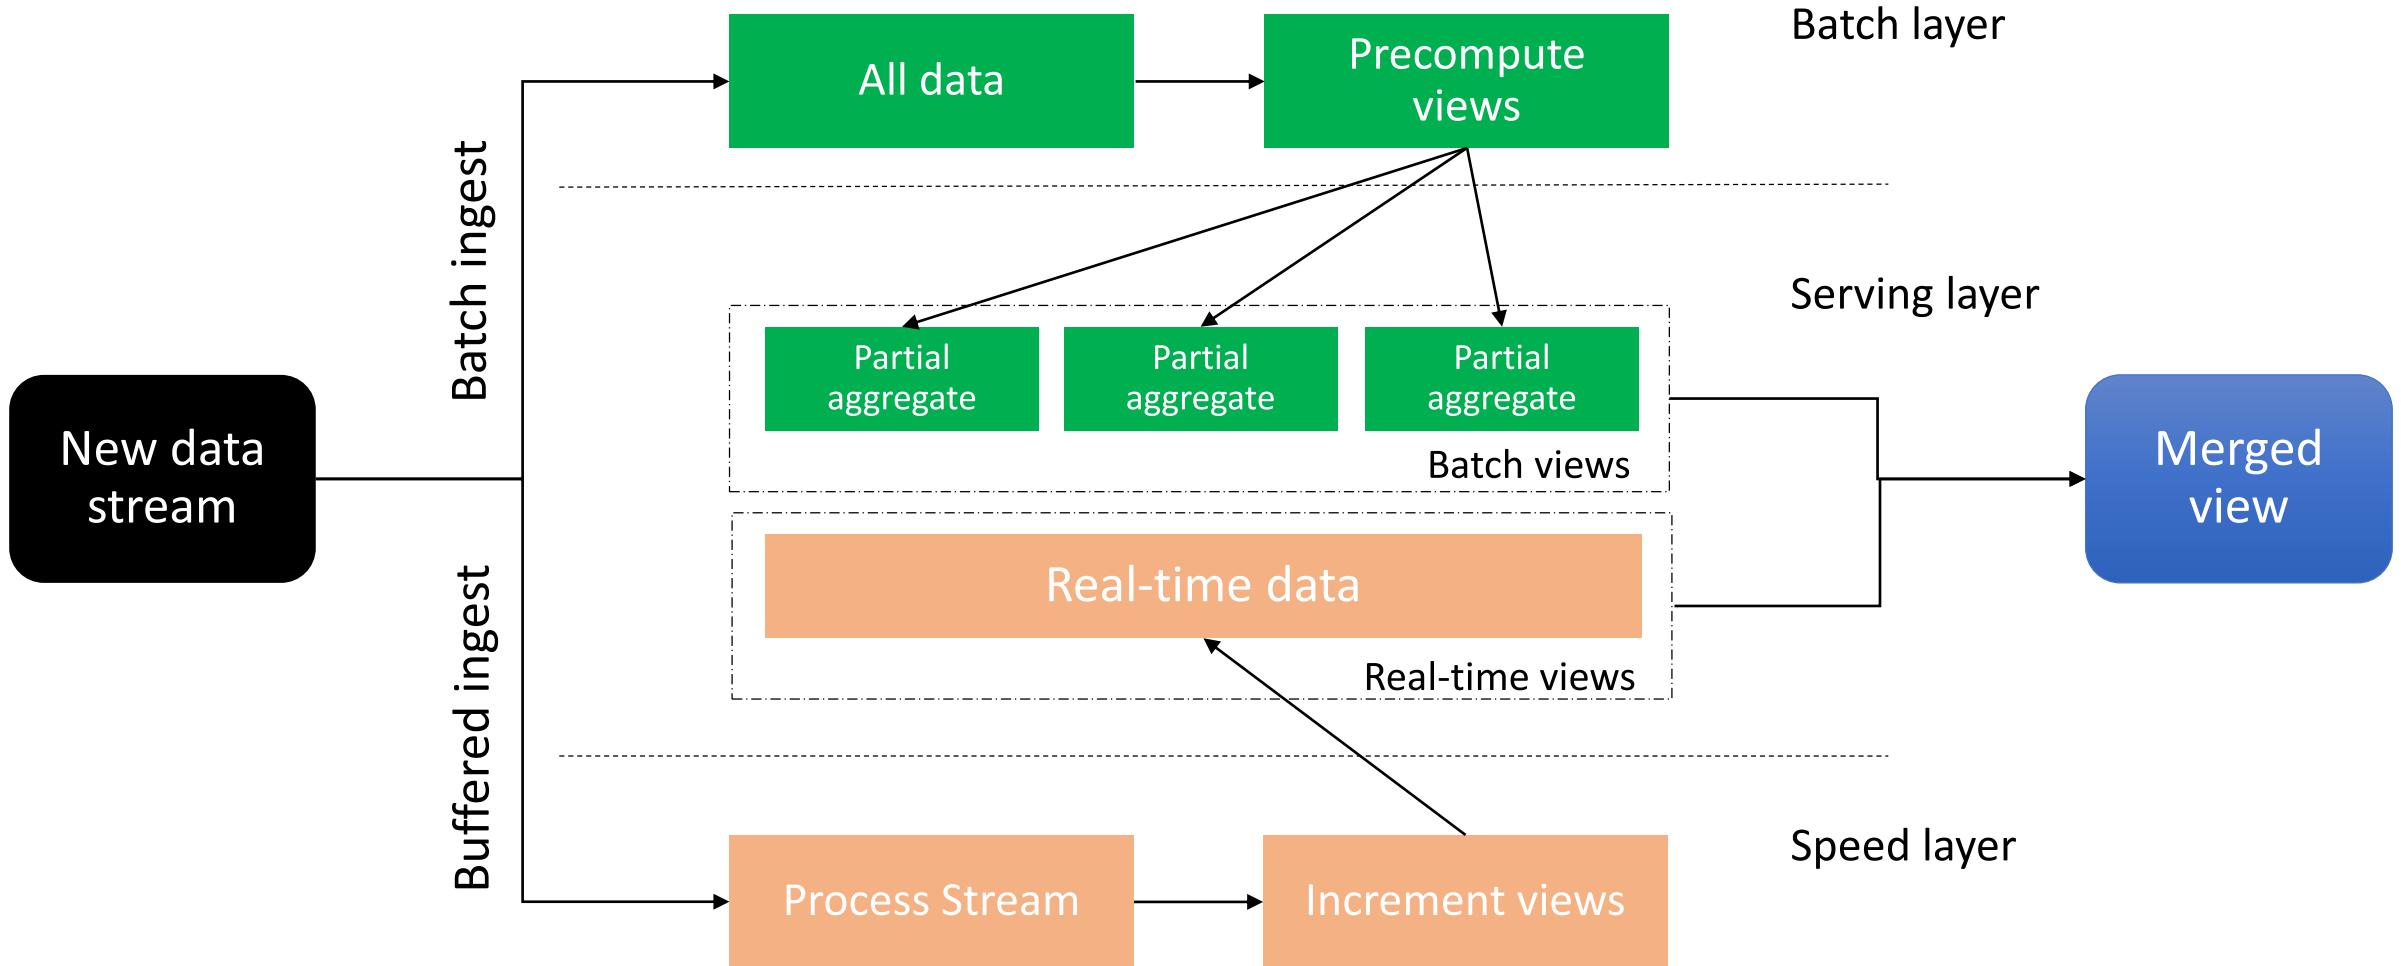
Azure HDInsight

Hadoop in Azure

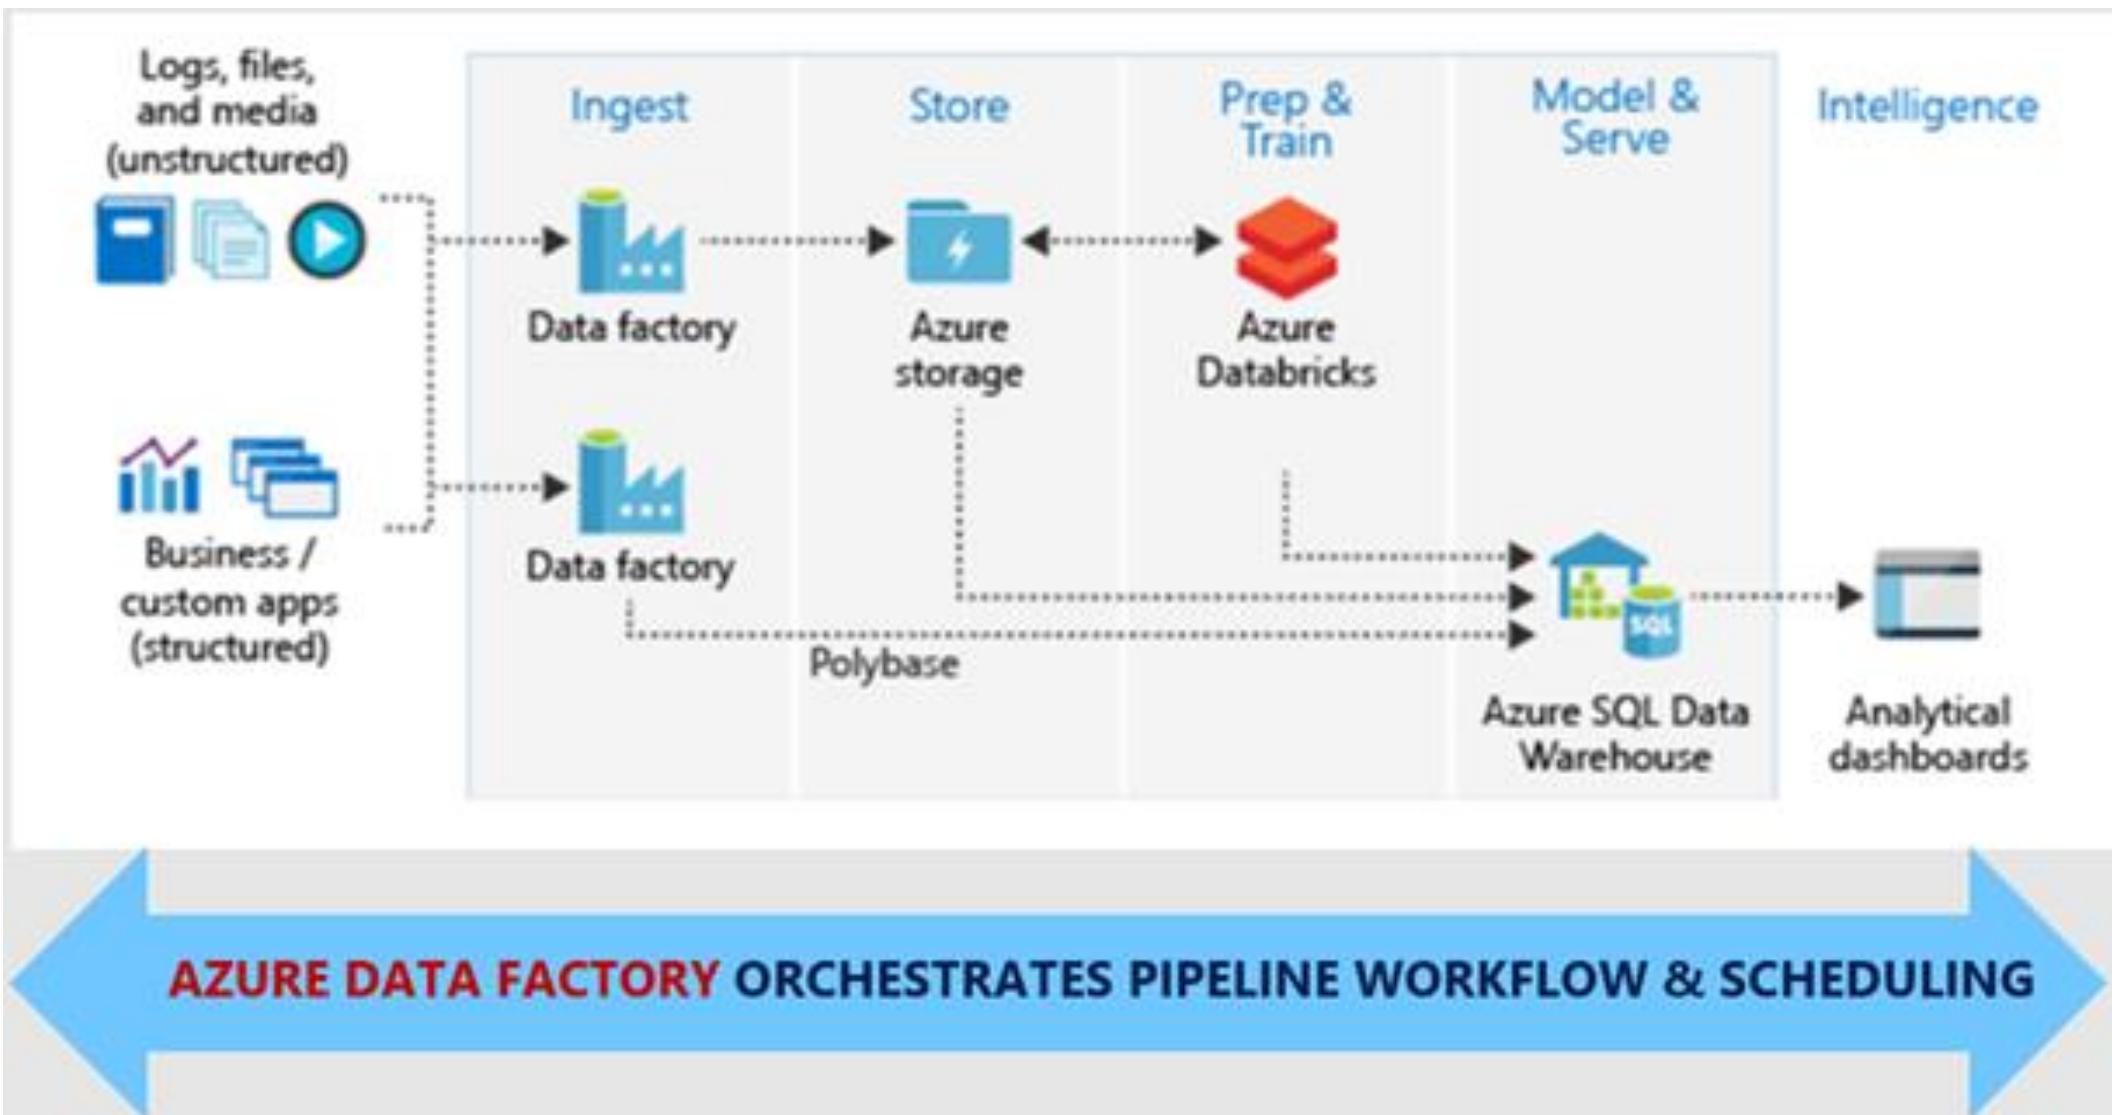
- Massively Parallel Processing (MPP) managed Hadoop cluster
- Spark SQL for real-time queries from in-memory caching
- Storm and Spark Streaming for near real-time processing of IoT
- HBase as a columnar NoSQL transactional database
- Hive SQL queries for batch processing



Lambda Architecture – High Level View



Azure Databricks



Azure Databricks

Azure Databricks Unified Analytics Platform

PEOPLE

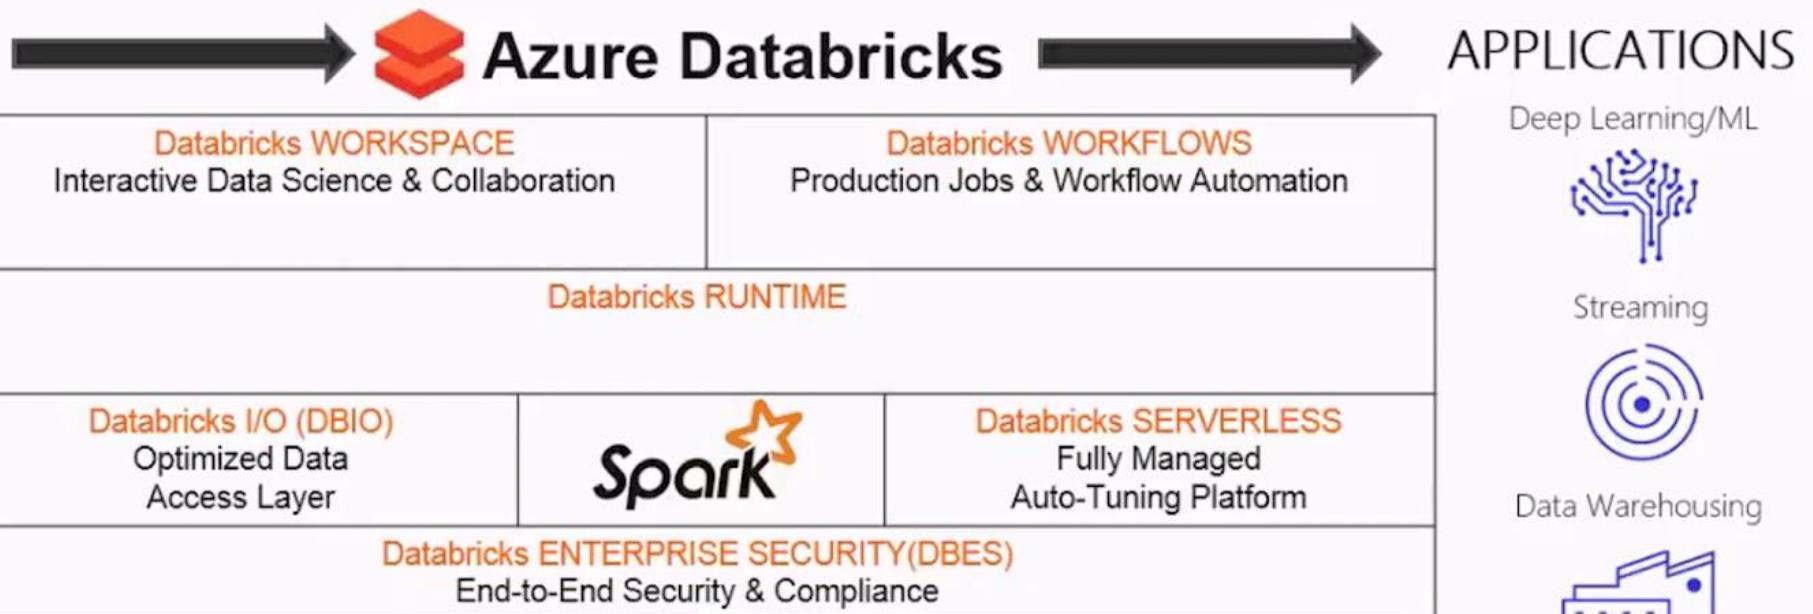
Data Science



Data Engineering



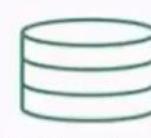
Line of Business



Cloud Storage



Databricks File System



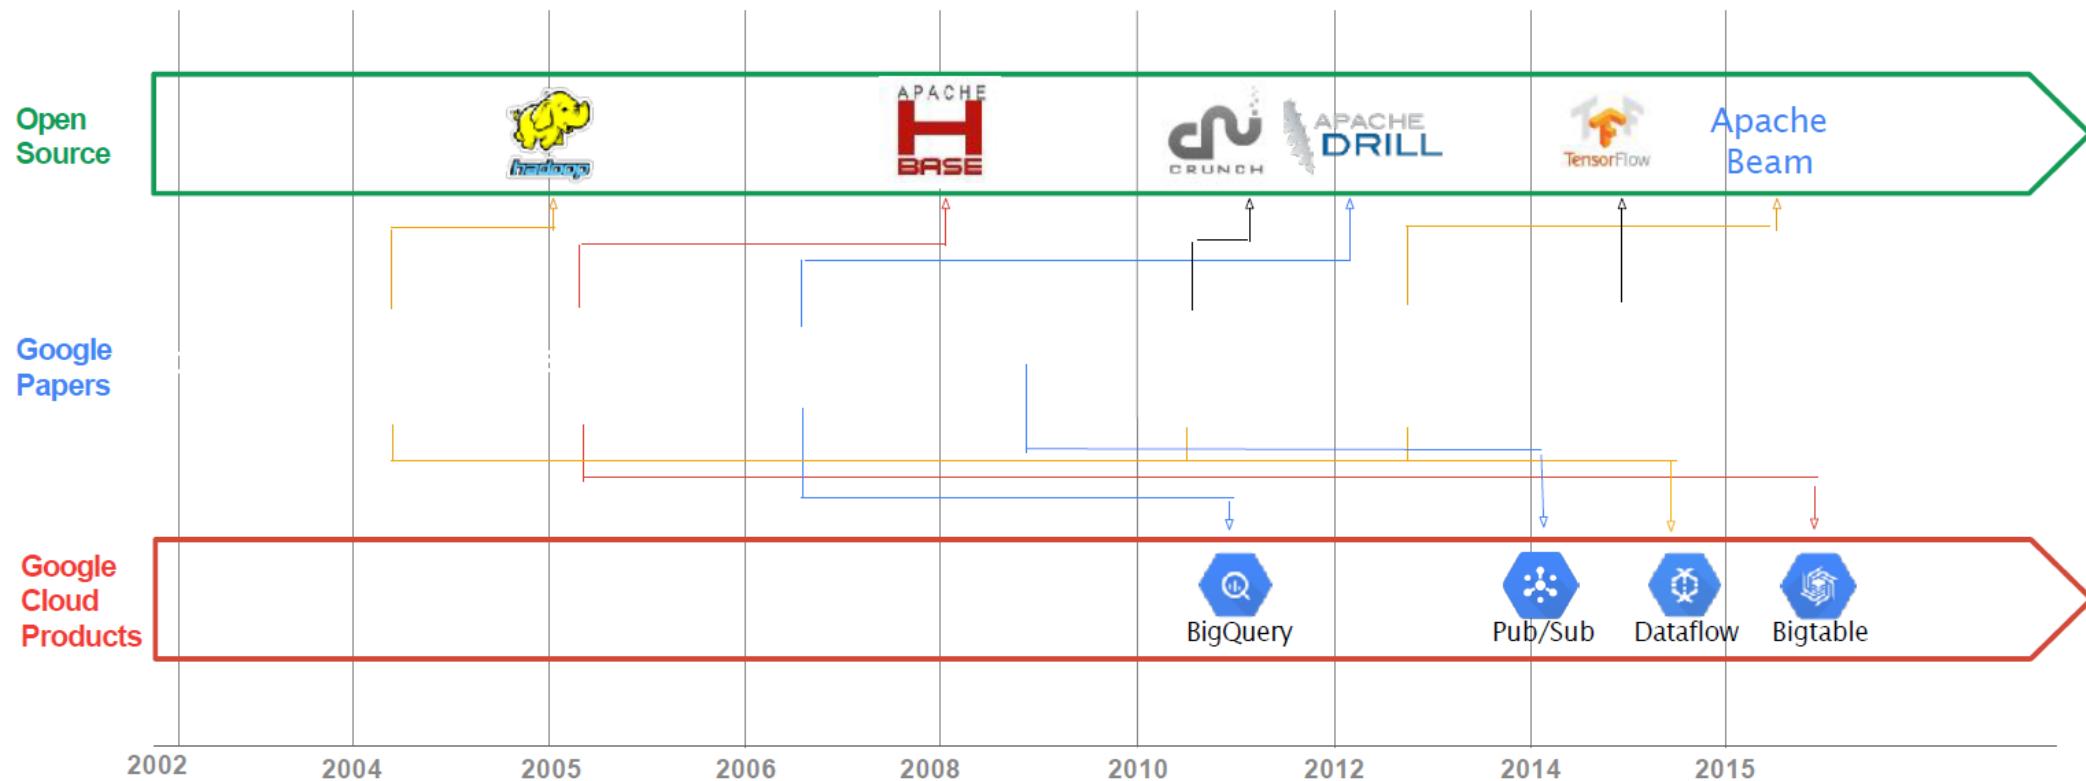
Data Warehouses

And many others...

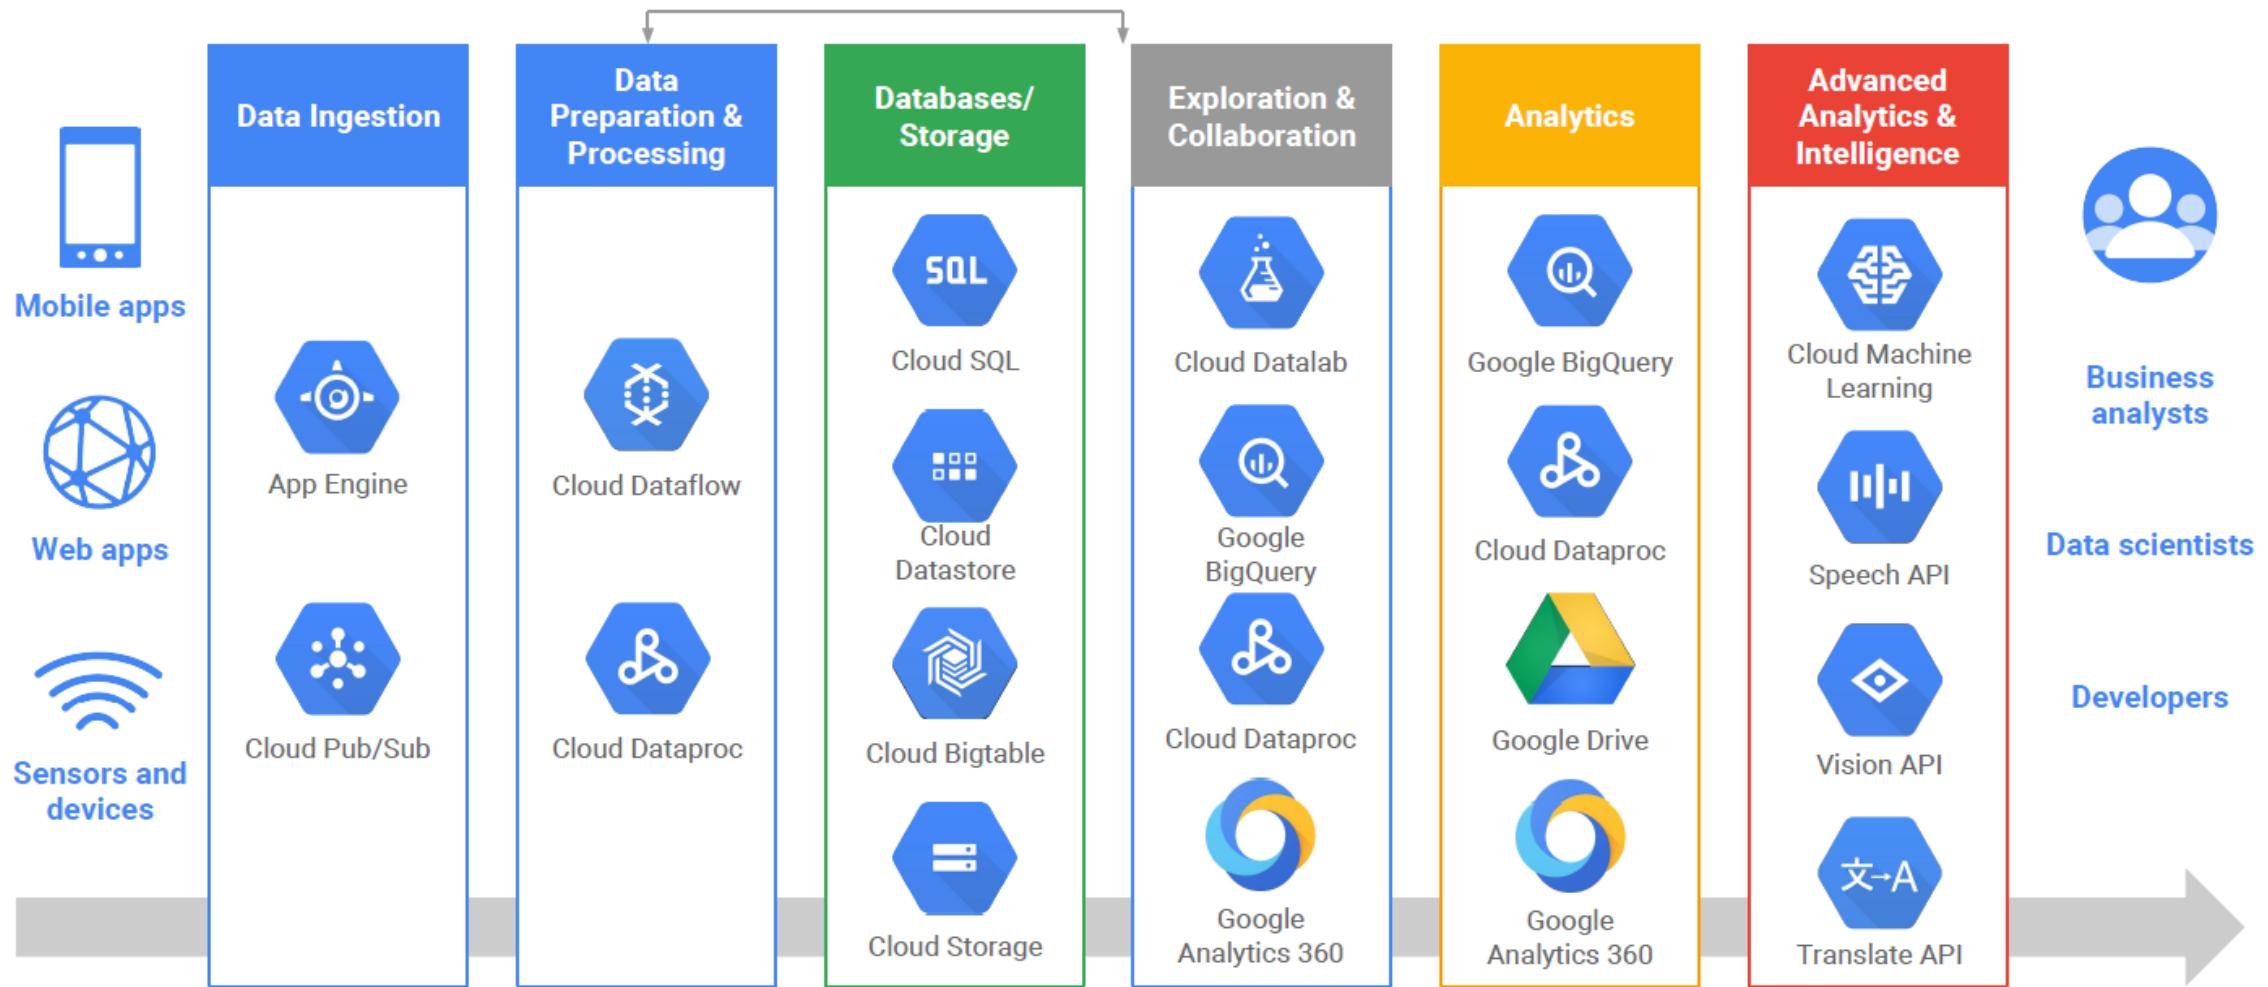
BigData on Google Cloud

<https://cloud.google.com/solutions/big-data>

10 years of Tackling Big Data Problems

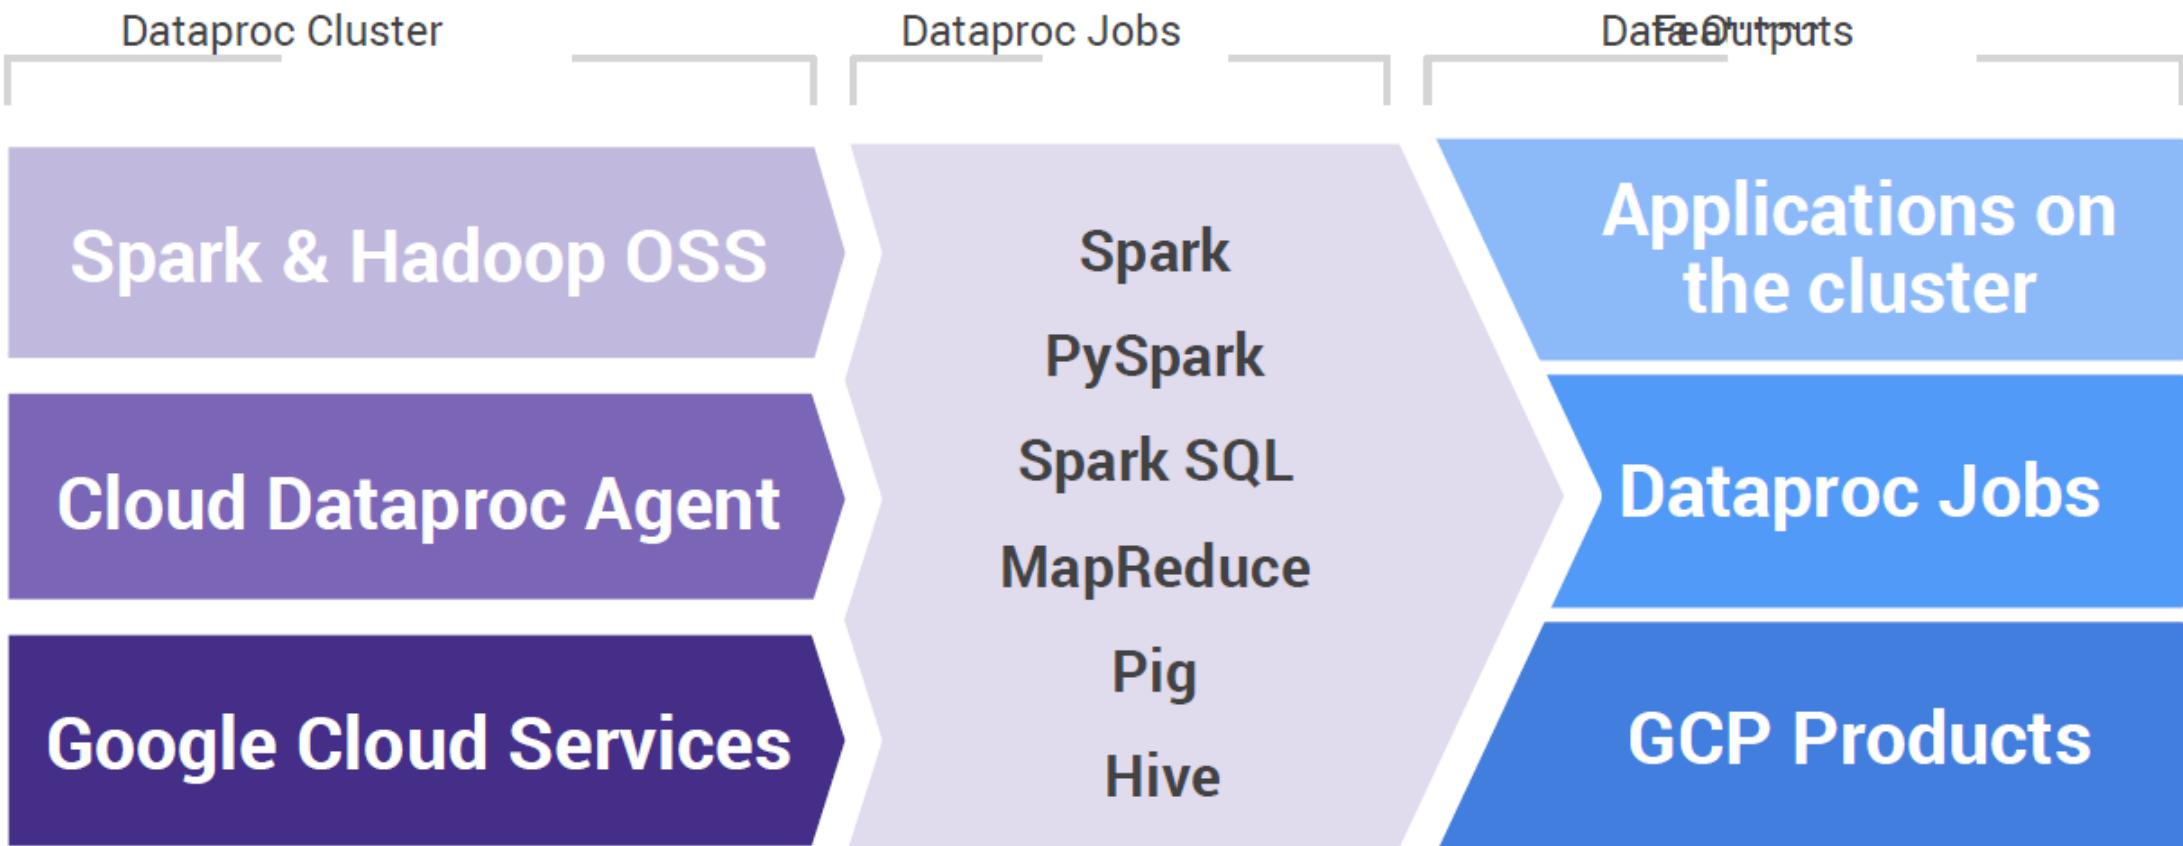


Transform Data into Actions



Google Cloud DataProc

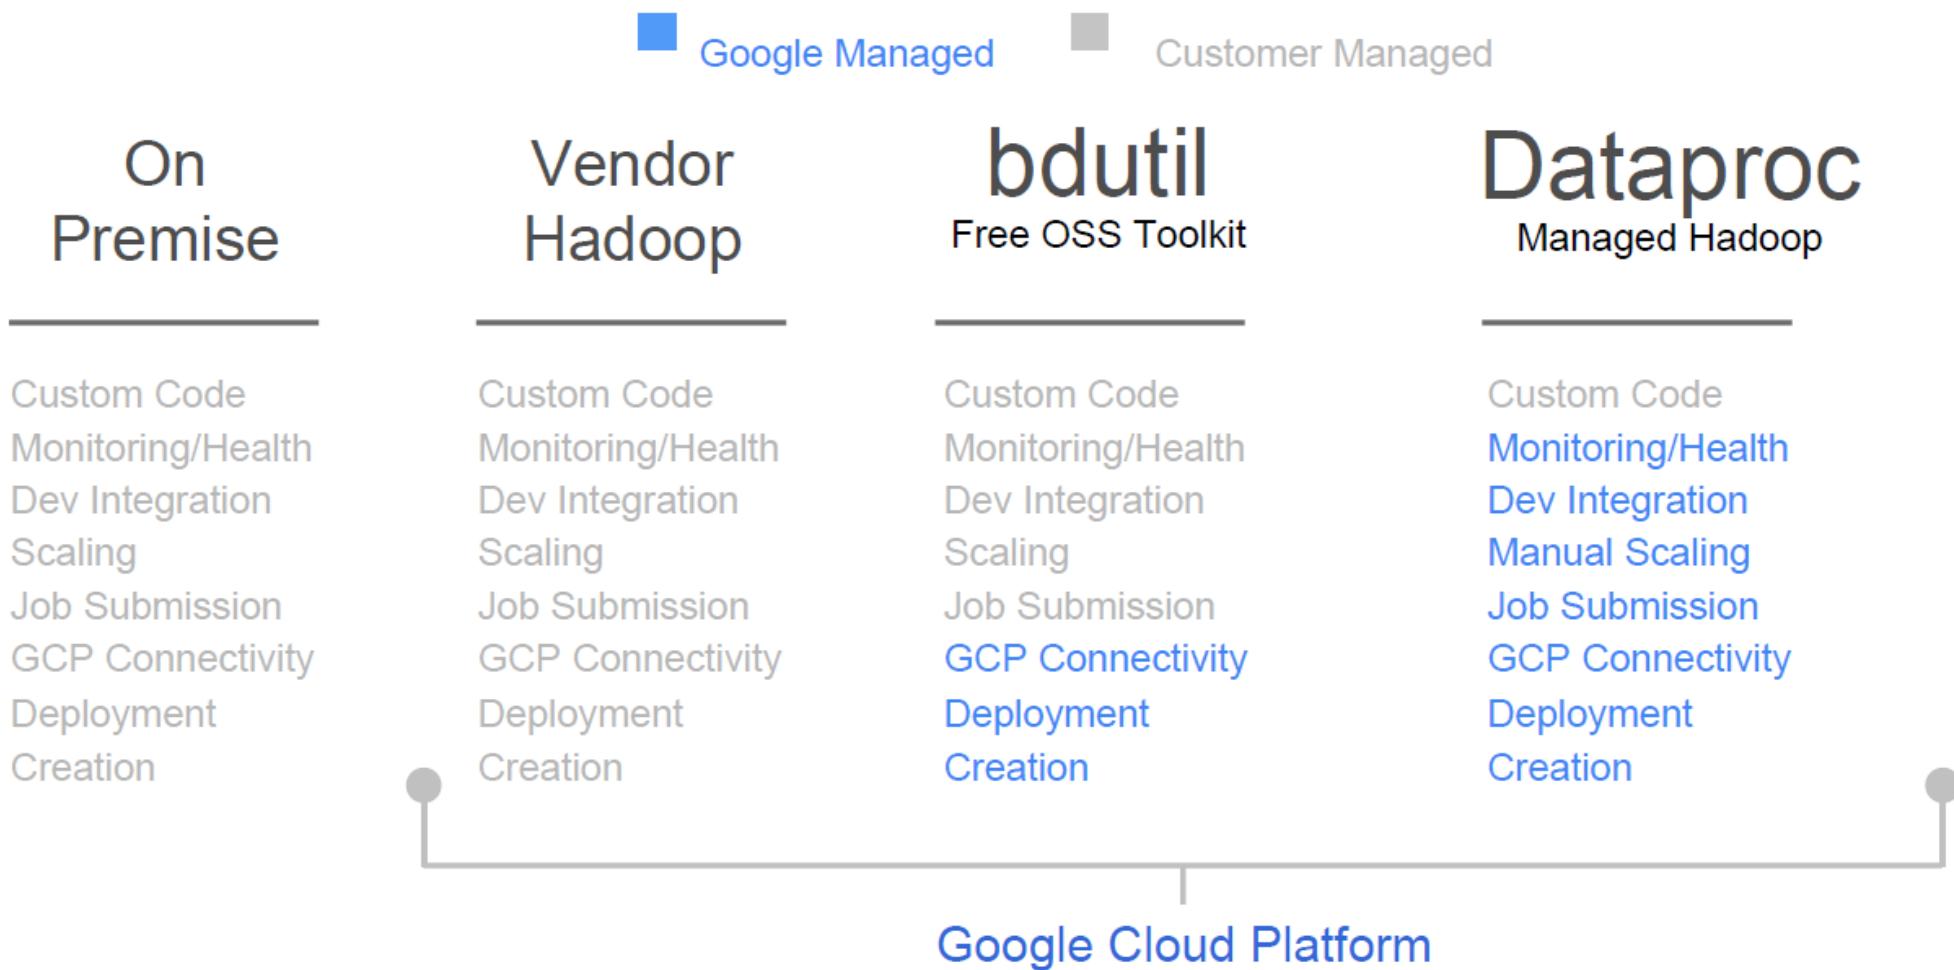
- Apache Spark and Apache Hadoop should be **fast, easy, and cost-effective**



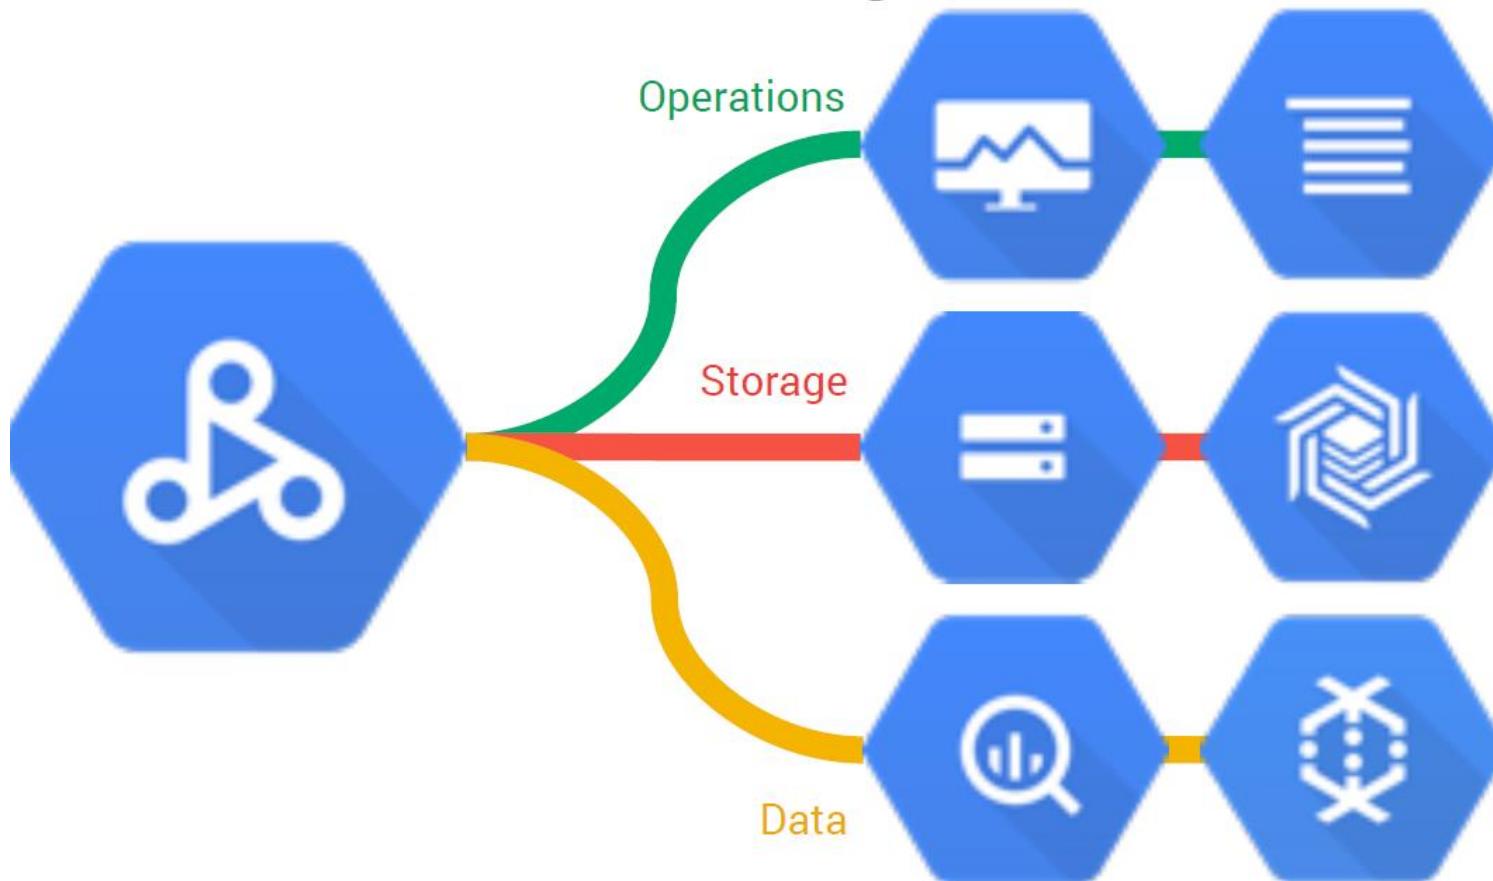
Easy, Fast. Cost-effective

- **Fast**
 - Things take seconds to minutes, not hours or weeks
- **Easy**
 - Be an expert with your data, not your data infrastructure
- **Cost-effective**
 - Pay for exactly what you use

Running Hadoop on Google Cloud

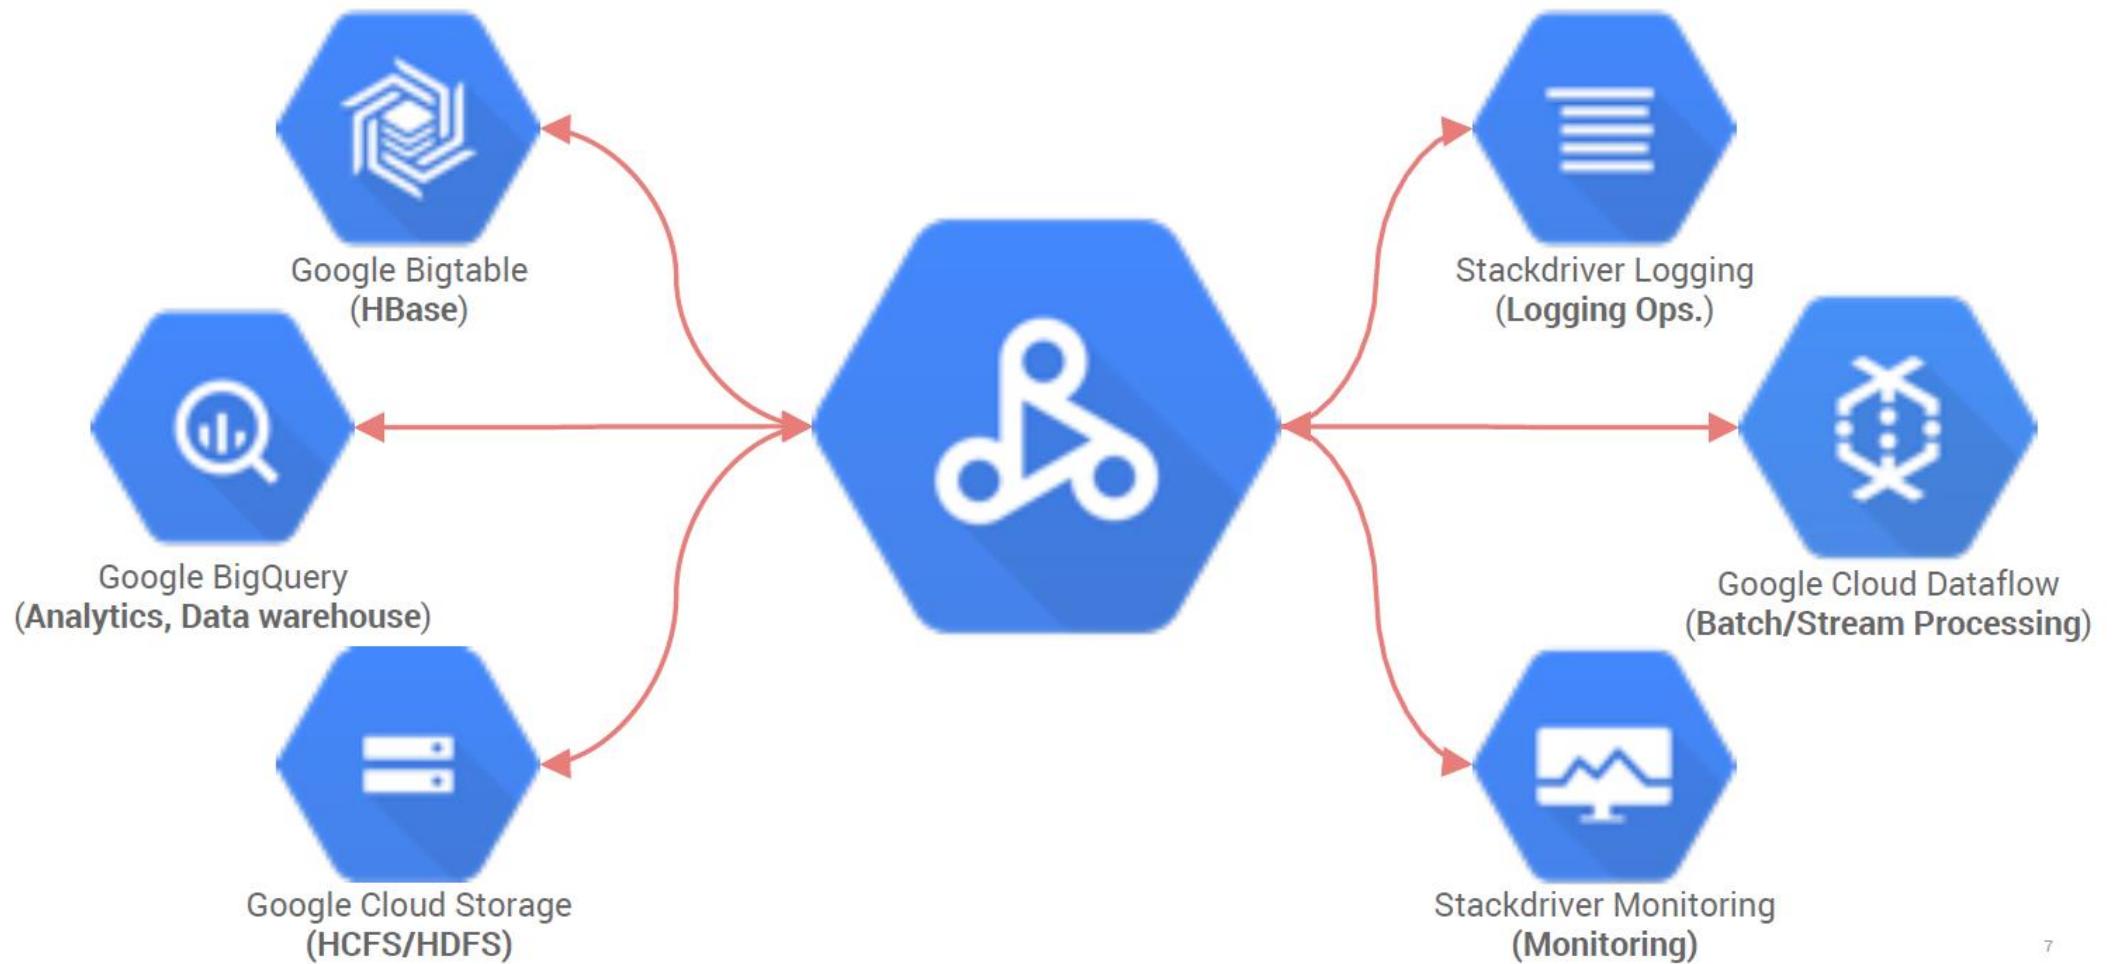


Cloud DataProc - integrated



Cloud Dataproc is natively integrated with several Google Cloud Platform products as part of an **integrated data platform**.

Where Cloud DataProc fits into GCP



Google BigQuery

- makes complex data analysis simple
 - Scales automatically
 - No setup or administration
 - Stream up to 100,000 rows p/sec
 - Easily integrates with third-party software

The screenshot shows the Google BigQuery Query Editor interface. The title bar says "100B Benchmark with 3 wildcards". The main area contains the following SQL query:

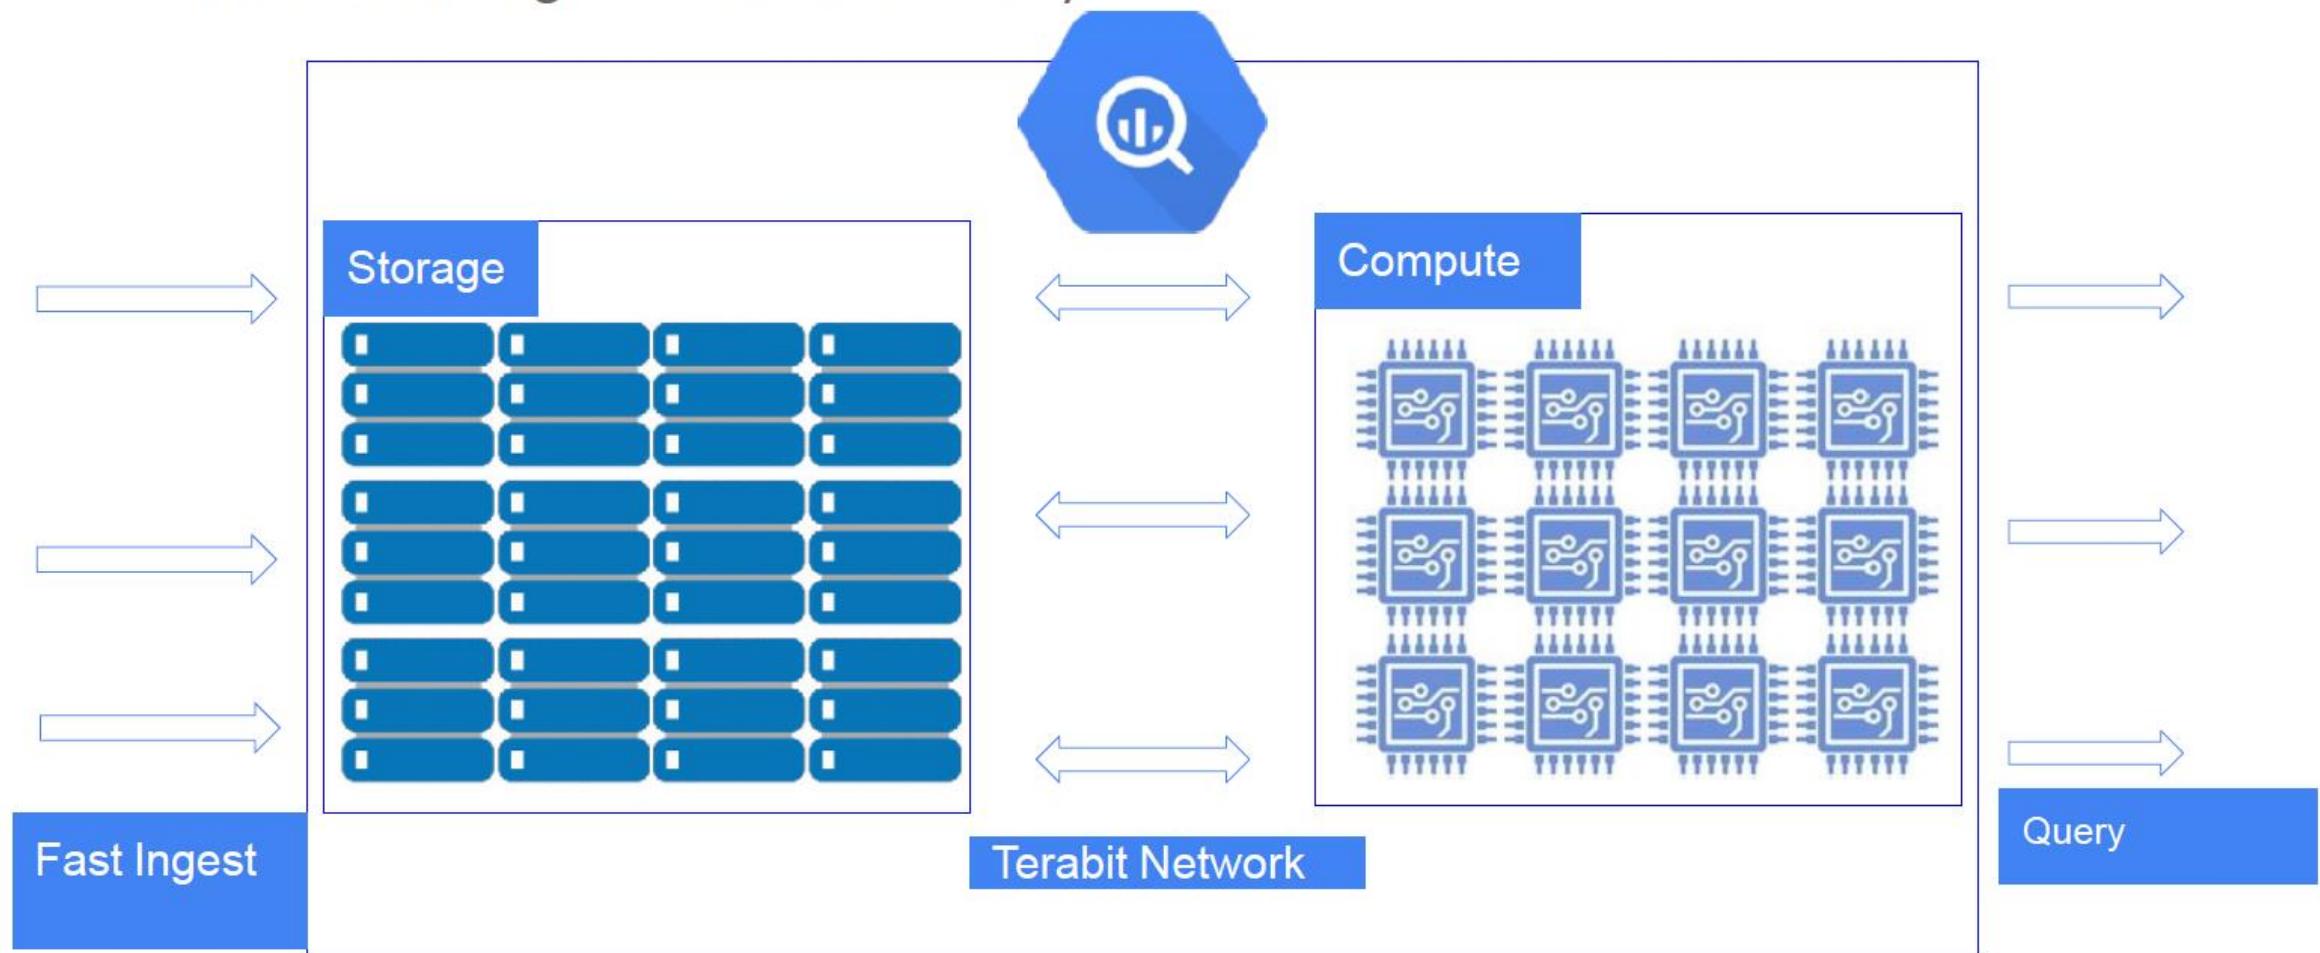
```
1 SELECT language, SUM/views) AS views
2 FROM `bigquery-samples:wikipedia_benchmark.Wiki100B`
3 WHERE REGEXP MATCH(title, "G.*o.*o.*g")
4 GROUP BY language
5 ORDER BY views DESC;
```

Below the query, a button says "No Cached Results". At the bottom, there are several buttons: "RUN QUERY" (highlighted in red), "Save Query", "Save View", "Format Query", and "Show Options". A status message at the bottom says "Query complete (24.7s elapsed, 4.06 TB processed)". A green checkmark icon is in the bottom right corner.

Running an *inefficient* regular expression over 100 billion rows in
less than 60 seconds

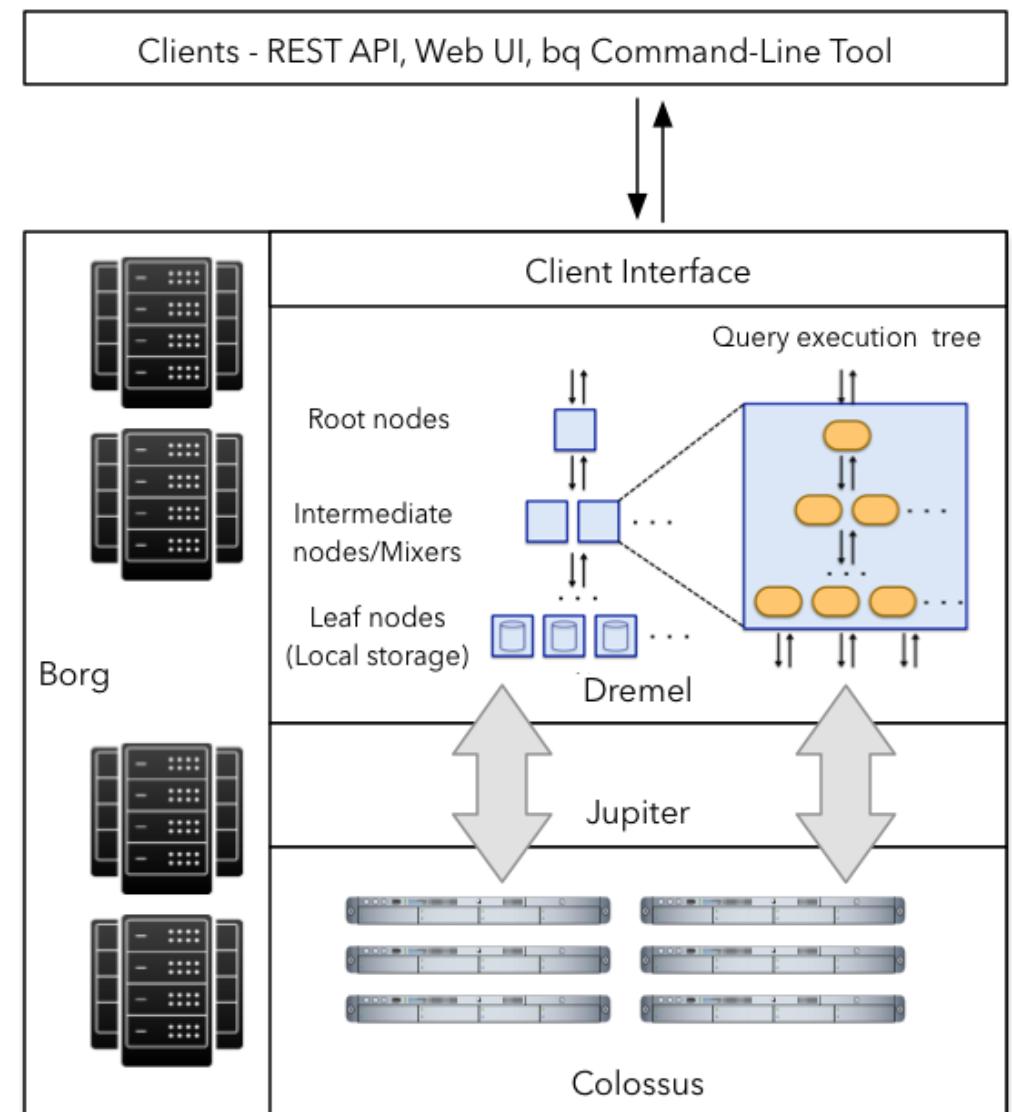
Google BigQuery Architecture

The Power of Google Dremel for everyone

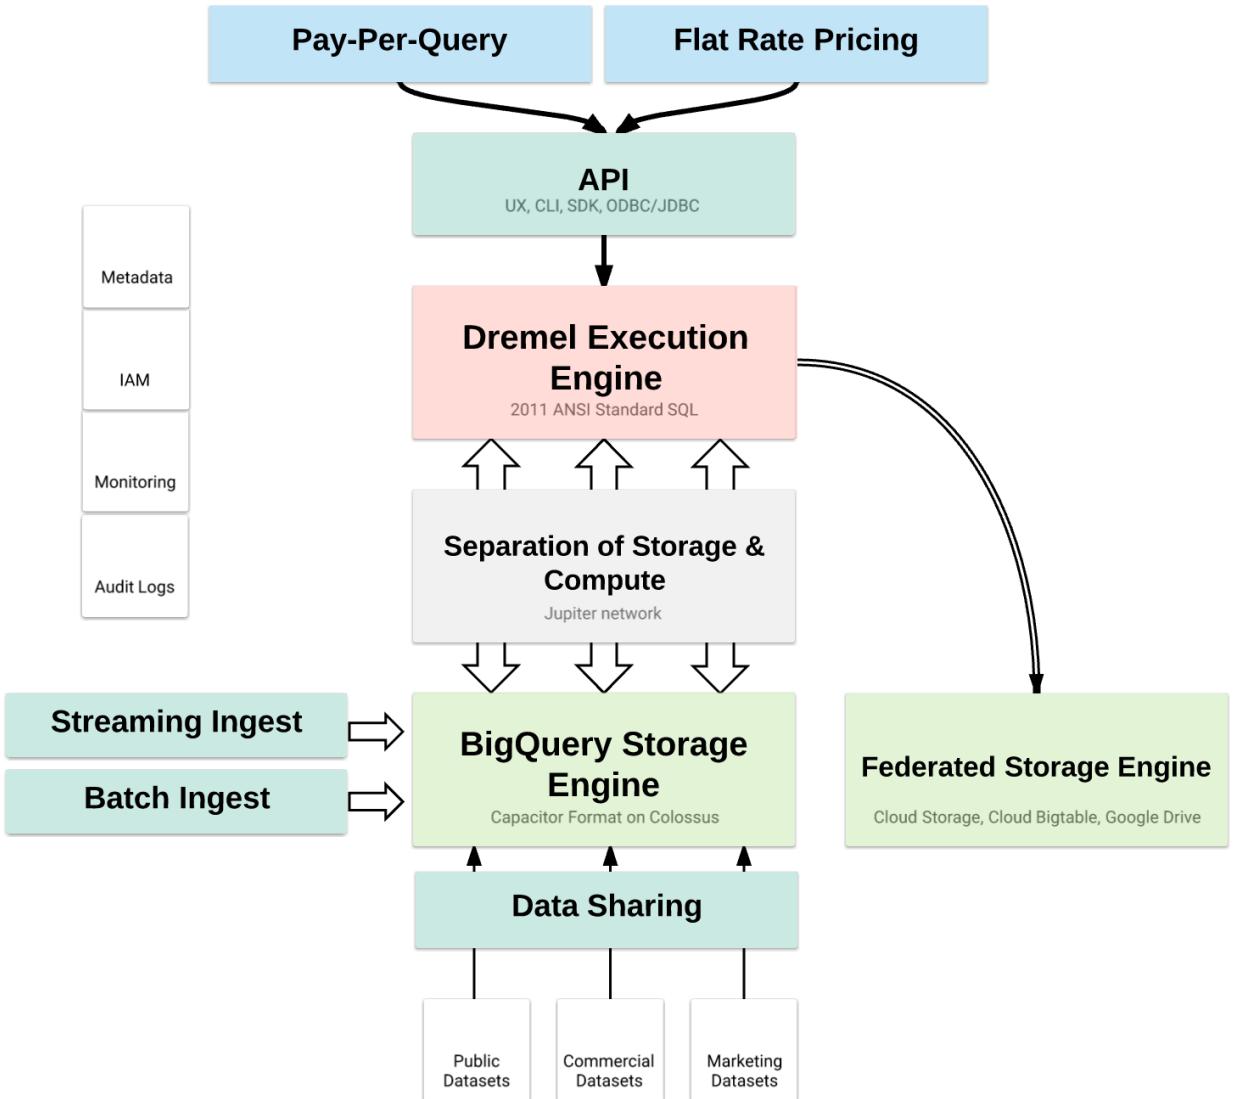


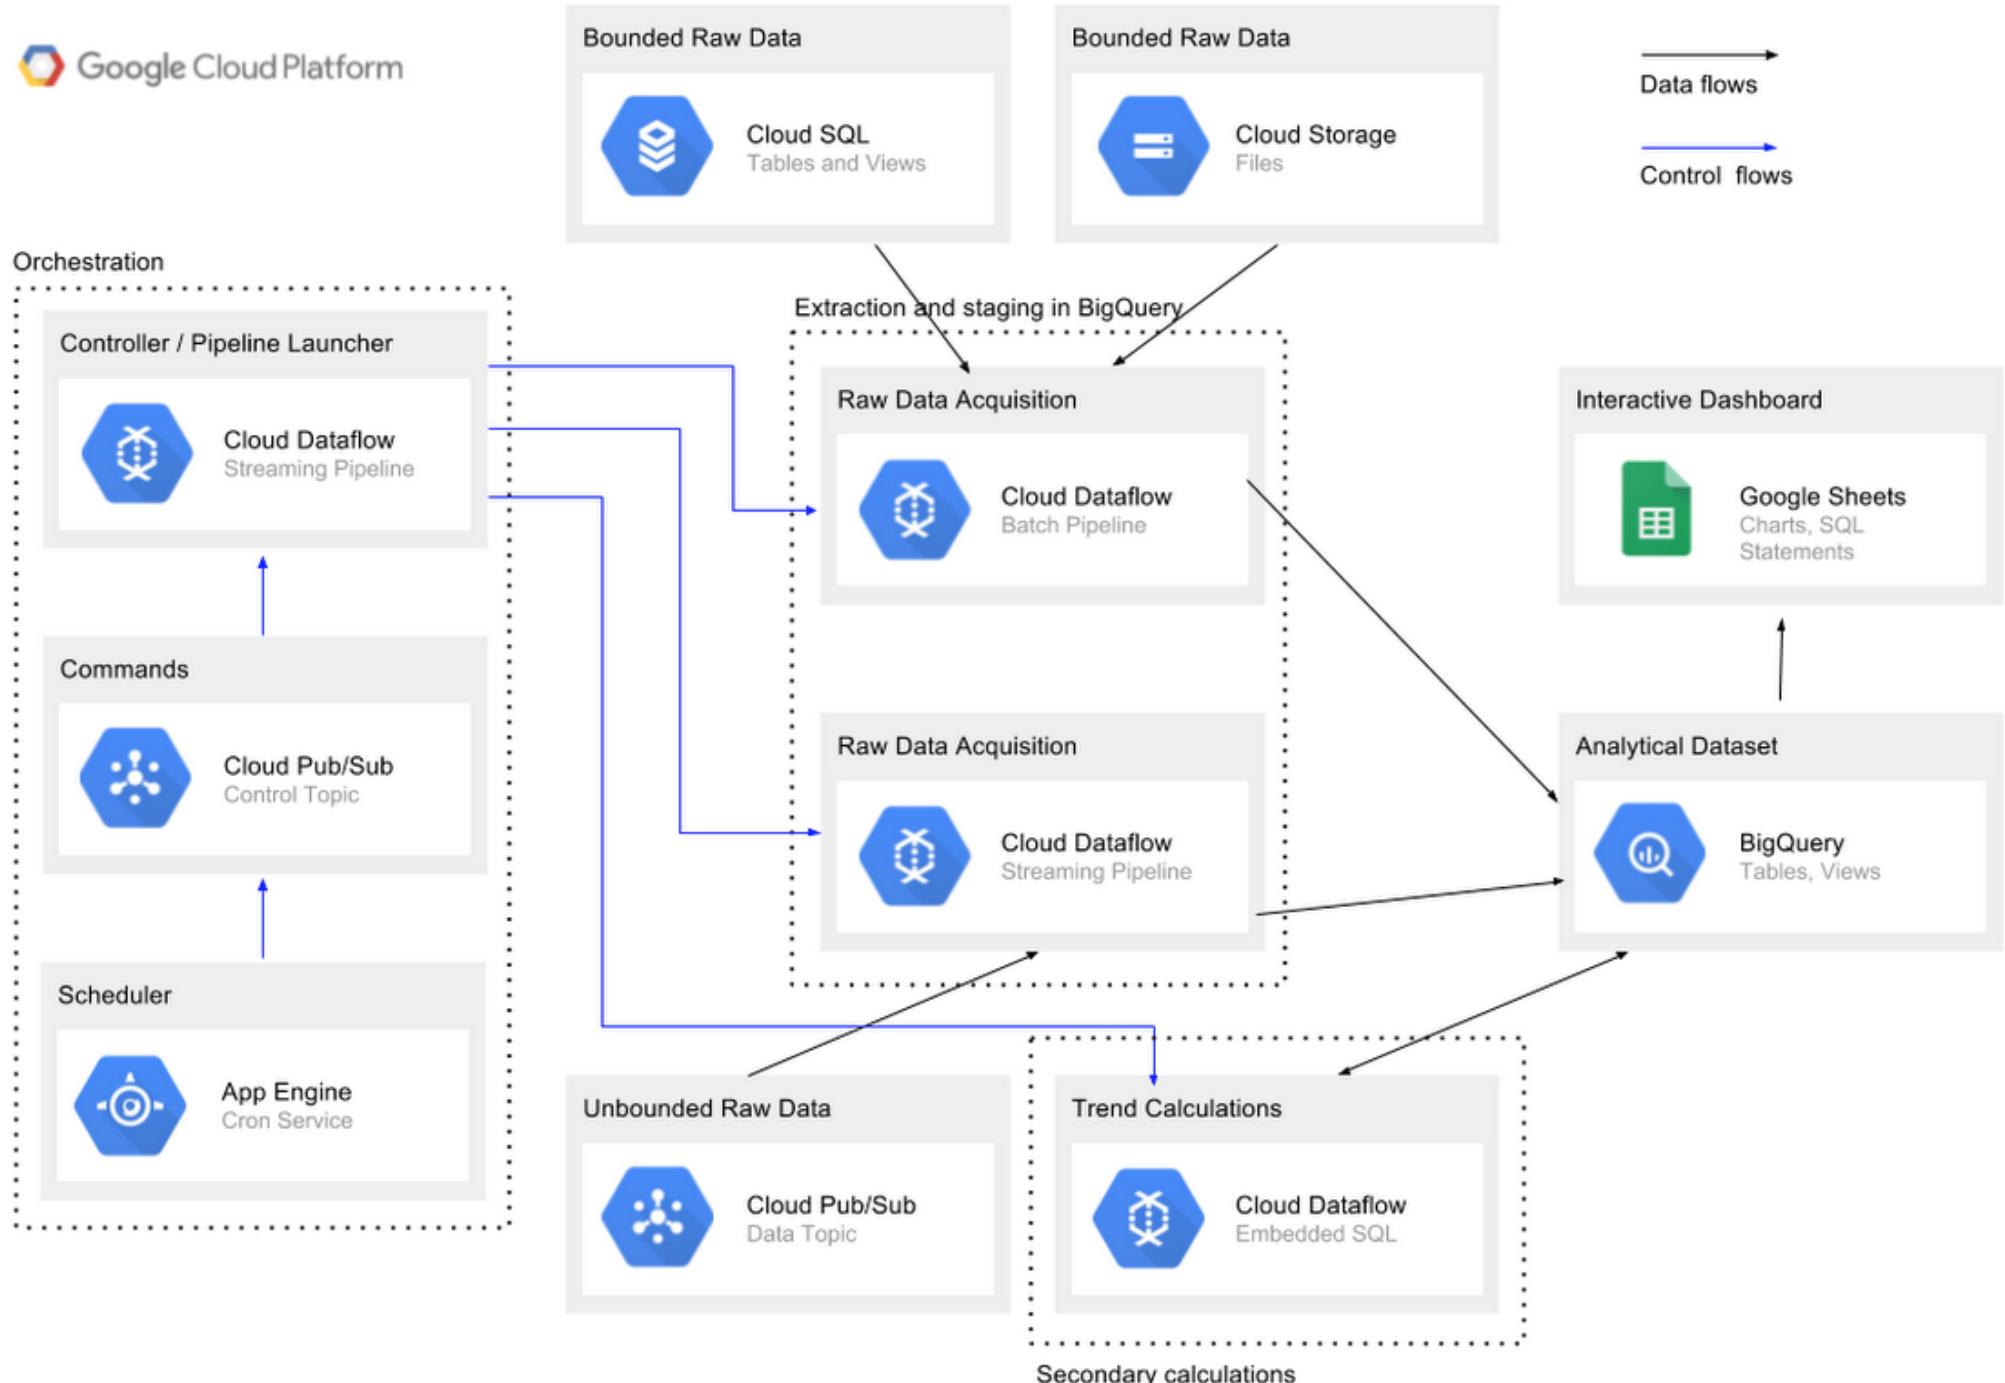
Google BigQuery Architecture

- Google Dremel
- Google Jupiter
- Google Colossus



Google BigQuery Architecture

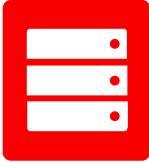




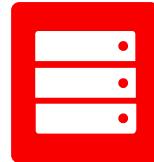
BigData on Oracle Cloud

<https://www.oracle.com/big-data/big-data-service/>

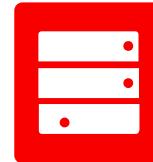
OCI Hardware



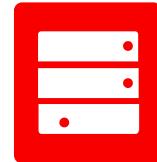
Bare Metal Standard
52 Cores, 768 GB RAM,
up to 1 PB Block Storage



Bare Metal DenseIO
51.2 TB of local NVMe SSD
2x 25Gbe Network Interfaces



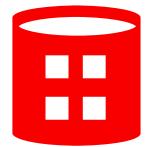
Bare Metal GPU
28 Cores, 192 GB RAM,
2x Tesla P100 GPUs
Pre-Configured Images



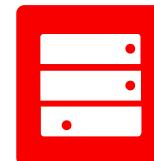
Bare Metal GPU V2
52 Cores, 768 GB RAM,
8x Tesla V100 GPUs
NVLINK Interconnect



File Storage Service
Managed distributed file service
POSIX, NFSv3 mount point



Block Storage
50 GB-2 TB volumes
Up to 400K IOPS per host



Bare-Metal HPC
36 cores, 3.7 GHz
384GB RAM
6.7 TB NVME, 1PB block
RDMA



GPU Visualization
P100 GPU
NVIDIA Quadro Enabled
Teradici & Citrix Support

The Best of Hadoop Ecosystem

There's no way around it. Hadoop is complicated. But Cloudera Enterprise Data Hub makes it easier to take advantage of Hadoop.

With Oracle Big Data Service and Cloudera, you can:

Use all of Cloudera Enterprise Data Hub

Have vertically integrated Spark, Kafka, Impala, Hive,
Solr, and more

Have automatically deployed security and
management features

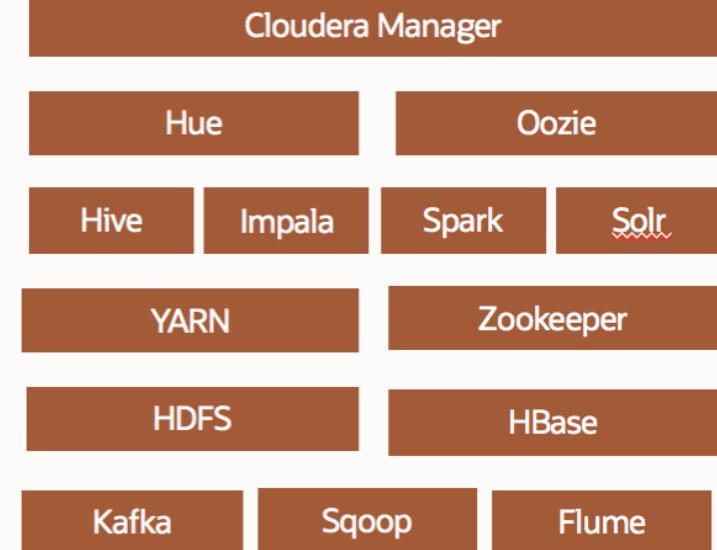
You also can choose your Cloudera version, giving you the ability
to:

Match your current deployment

which is important for test and dev environments

Deploy new versions

allowing you to take advantage of the distribution's latest features



- Authentication: Kerberos
- Authorization: Sentry
- Auditing: Navigator
- Data at-rest encryption: HDFS Transparent Encryption

Oracle Big Data Service offers the compute and storage you want

Development ➡ Test ➡ Data Science ➡ Data Lakes



VM Standard
1-24 OCPUs
15-320 GB RAM
Up to 1PB Block



Dense IO
8-24 OCPUs
120-320 GB RAM
6.4-25.6TB NVMe



Bare Metal Standard
52 OCPUs
762 GB RAM
Up to 1PB Block



Bare Metal HPC Dense IO
36 OCPUs (3.7GHz)
384 GB RAM
6.7 TB NVMe



Bare Metal Dense IO
52 OCPUs
768 GB RAM
51.2 TB NVMe

Performance

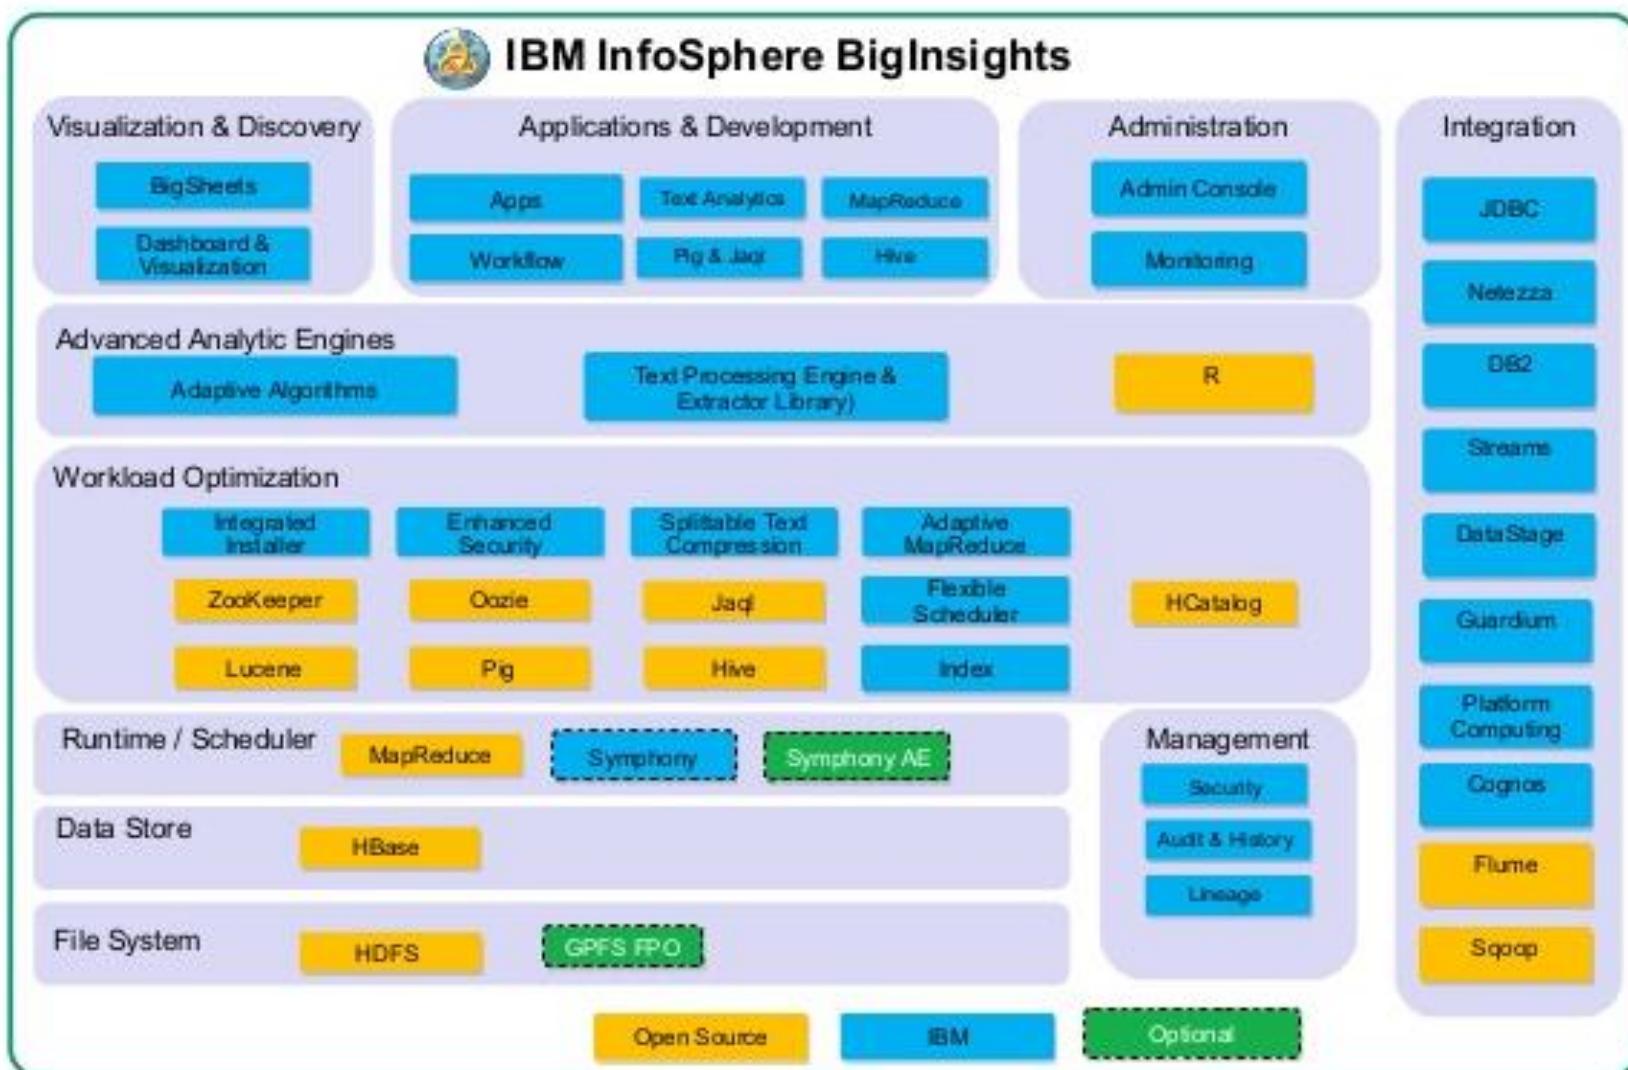
Flexibility

You can use Oracle Big Data Cloud for development, test, data science, or data lakes. In all of these cases, we provide the compute offerings you need for your use case. You can achieve ultimate flexibility by combining bare metal and virtual machines with different storage options, which optimizes your deployment based on your resource needs.

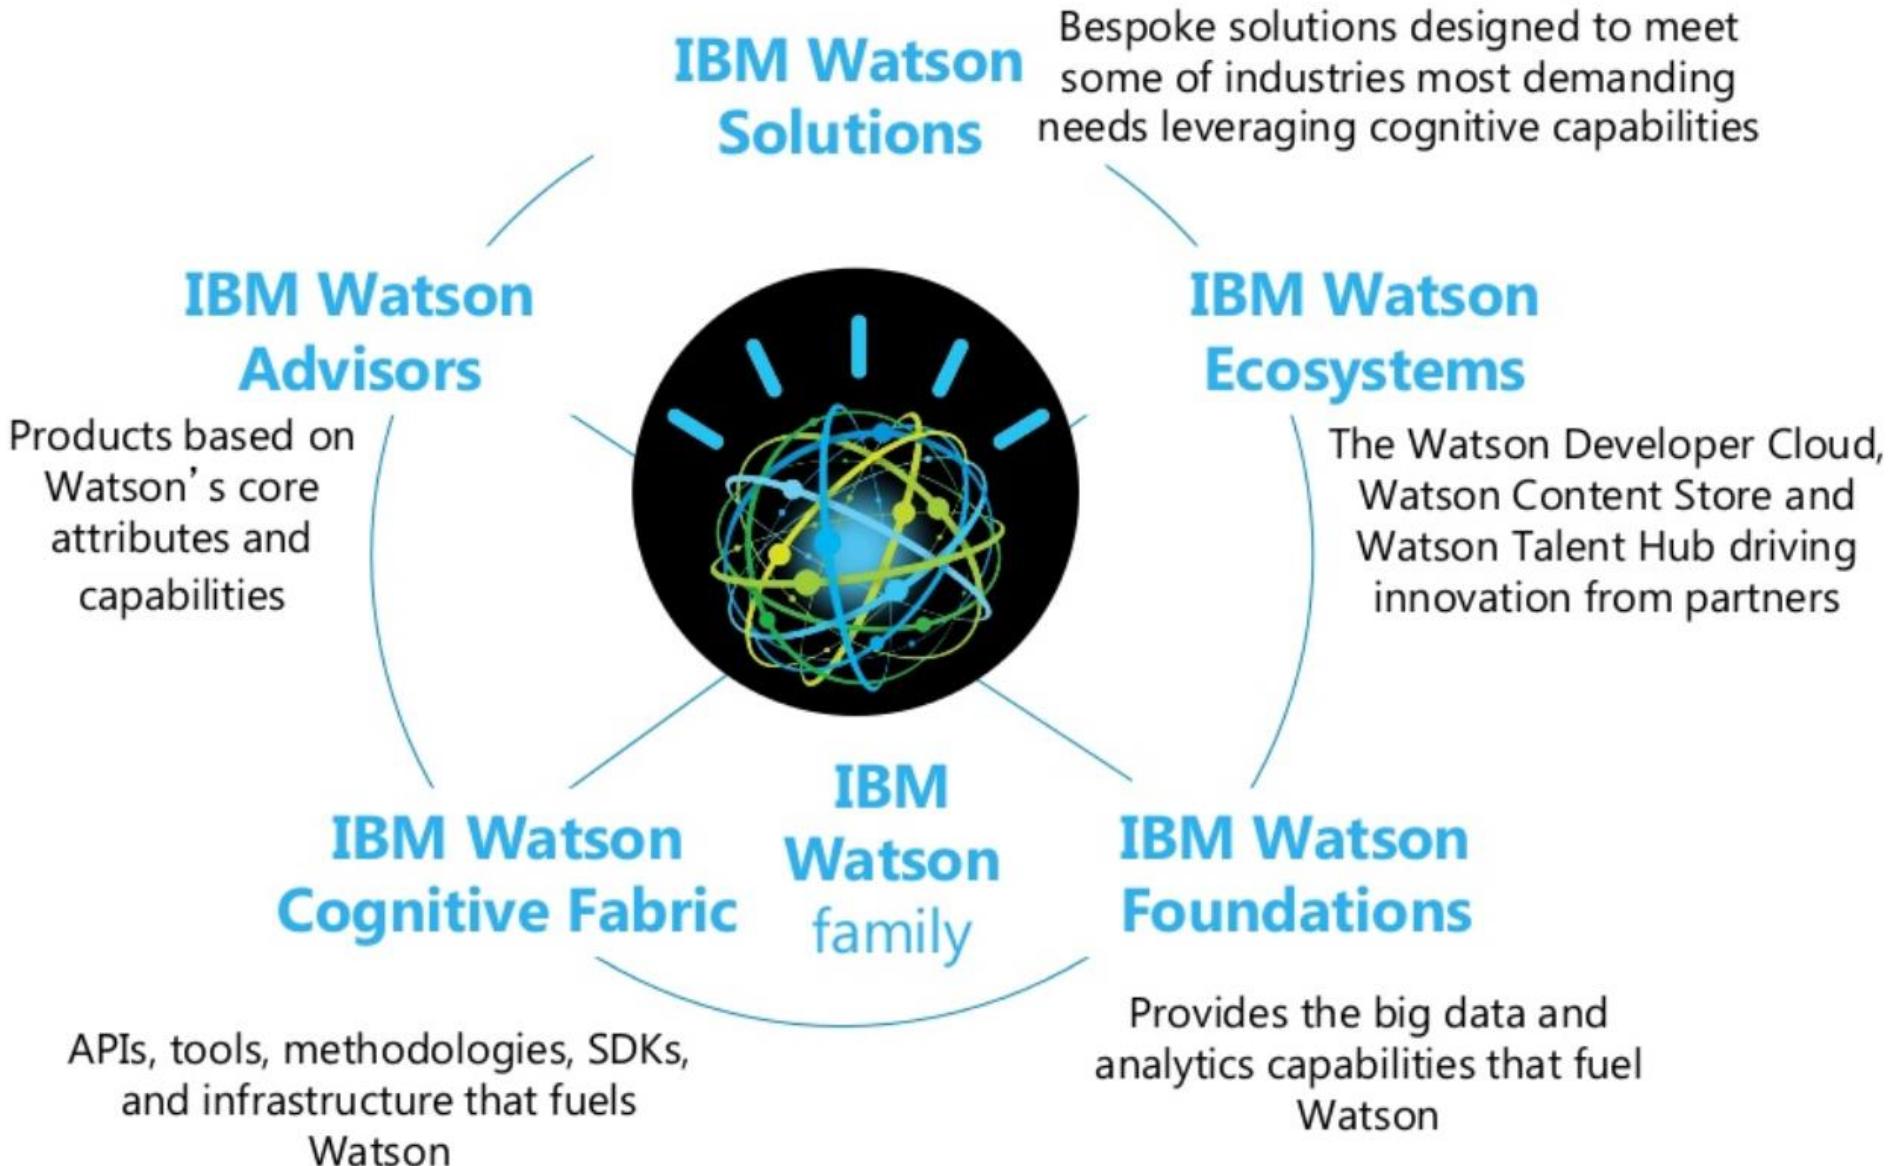
BigData on IBM Cloud

<https://www.ibm.com/uk-en/it-infrastructure/solutions/big-data>

IBM Big Data Platform



IBM Watson Family



Big Data & Analytics Infrastructure

